Kansas Frontier Matrix (KFM)

Definitive Design, Governance, Data, UX, and Discovery Guide (vNext)

Version: vNext Draft Consolidation • Date: 2026-02-20

Map-first • Time-aware • Governed • Evidence-first • Cite-or-abstain

This document consolidates the current KFM blueprint set and expands the unresolved gaps into buildable, testable guidance. It is intended to be used as a working engineering and governance manual.

Primary internal sources: IB-v5 (2026-02-20), DesignPlan-v2 (2026-02-19), DocUpgrade-v2 (2026-02-20), DSIB-v1 (2026-02-12).

Table of contents

Table of contents 2

1. Purpose, scope, and how to use this guide 13

1.1 Source-of-truth documents used . . . 13

1.2 Normative language and tagging . . . 13

1.3 Non-goals (so the scope remains buildable) . . . 13

2. System north stars and invariants (do not compromise) 14

2.1 The "truth path" lifecycle (CONFIRMED) . . . 14

2.2 Trust membrane (CONFIRMED) . . . 14

2.3 Evidence-first UX (CONFIRMED concept; implementation PROPOSED) . . . 14

2.4 Cite-or-abstain Focus Mode (CONFIRMED) . . . 15

2.5 Canonical vs rebuildable stores (CONFIRMED concept) . . . 15

2.6 Deterministic identity and hashing (CONFIRMED) . . . 15

3. Key KFM concepts and definitions 15

3.1 Dataset, DatasetVersion, and artifacts . . . 15

3.2 EvidenceRef and EvidenceBundle . . . 16

3.3 Story Nodes . . . 16

3.4 Policy labels and redaction obligations . . . 16

3.5 Run receipts and audit ledger . . . 16

3.6 Time: event time, valid time, and transaction time (PROPOSED) . . . 17

4. Architecture reference model (map-first, time-aware, governed) 17

4.1 Layering and responsibilities (CONFIRMED posture) . . . 17

4.2 Component decomposition (PROPOSED as a build plan) . . . 17

4.3 Canonical vs rebuildable stores (design decision) . . . 18

4.4 Data flow diagram (conceptual) . . . 18

4.5 Deployment topology (PROPOSED default) . . . 18

5. Data lifecycle zones and the Promotion Contract 19

5.1 Zone definitions (refined) . . . 19

RAW zone (immutable acquisition) . . . 19

WORK / QUARANTINE (intermediate transforms) . . . 19

PROCESSED (publishable) . . . 19

CATALOG/TRIPLET (DCAT + STAC + PROV) . . . 19

PUBLISHED (governed runtime) . . . 20

5.2 Promotion Contract v1 (minimum gates) . . . 20

Gate A: Identity and versioning . . . 20

Gate B: Licensing and rights metadata . . . 20

Gate C: Sensitivity classification and redaction plan . . . 20

Gate D: Catalog triplet validation . . . 20

Gate E: Run receipt and checksums . . . 20

Gate F: Policy tests and contract tests . . . 21

Gate G: Optional but recommended (production posture) . . . 21

5.3 Definition of Done (dataset integration ticket) . . . 21

5.4 A practical promotion workflow (PR-based) . . . 21

6. Identity, naming, controlled vocabularies, and hashing 21

6.1 Identifier families (recommended) . . . 22

6.2 Dataset slug conventions (PROPOSED) . . . 22

6.3 DatasetVersion identity via spec_hash (CONFIRMED pattern; refine) . . . 22
6.4 Controlled vocabularies (minimum set) . . . 23
policy_label (access + sensitivity) . . . 23
artifact.zone (data lifecycle) . . . 23
citation.kind (evidence types) . . . 23
geometry.generalization_method . . . 23
6.5 Hash and digest policy . . . 24
6.6 "Hash drift" prevention checklist . . . 24
7. Governance and ethics: FAIR + CARE implemented as gates 24
7.1 FAIR + CARE operationalization . . . 24
7.2 Roles and responsibilities (gap closure) . . . 25
PROPOSED baseline roles . . . 25
RACI (minimum) . . . 25
7.3 Governance artifacts (what must exist) . . . 25
7.4 Policy-as-code architecture (shared CI and runtime semantics) . . . 26
7.5 Sensitivity default rules . . . 26
7.6 Licensing and rights enforcement . . . 26
8. Catalogs and provenance as contract surfaces 27
8.1 The catalog triplet responsibilities . . . 27
8.2 KFM DCAT profile (PROPOSED minimum) . . . 27
8.3 KFM STAC profile (PROPOSED minimum) . . . 27
8.4 KFM PROV profile (PROPOSED minimum) . . . 28
8.5 Cross-linking rules (must be testable) . . . 28
8.6 Evidence resolver contract (refined) . . . 28

8.7 Catalog and evidence validation in CI . . . 29

9. Domain modeling: connecting Kansas history, science, and place 29

9.1 A pragmatic core entity set (PROPOSED) . . . 29

9.2 Spatiotemporal representation . . . 29

Geometry types . . . 29

Time types . . . 30

9.3 Uncertainty as first-class metadata (PROPOSED but recommended) . . . 30

9.4 Boundary changes and bitemporality (DECISION NEEDED by domain) . . . 30

9.5 Entity resolution strategy (PROPOSED) . . . 30

9.6 Claim model for narratives (PROPOSED) . . . 31

9.7 Mapping "feature -> evidence" in spatial layers . . . 31

10. Ingestion, normalization, and validation pipelines 31

10.1 Connector patterns (PROPOSED) . . . 31

10.2 Normalization patterns by data class . . . 32

Tabular (CSV, JSON) . . . 32

Vector geospatial . . . 32

Raster . . . 32

Documents (PDF, scans, newspapers) . . . 32

10.3 Validation gates (minimum) . . . 33

10.4 Deterministic pipeline design . . . 33

10.5 Quarantine workflow (fail closed) . . . 33

11. Mapping and cartographic design (map-first, time-aware) 33

11.1 Layer taxonomy (recommended) . . . 34

11.2 Time-aware interaction patterns . . . 34

11.3 Performance and data delivery . . . 34

11.4 Tile strategy (DECISION NEEDED) . . . 34

11.5 Symbology and uncertainty visualization . . . 35

11.6 Map state as reproducible artifact . . . 35

12. User interface: Map Explorer, Story Mode, Focus Mode, and trust surfaces 35

12.1 Information architecture (refined) . . . 36

12.2 Core UI components (buildable) . . . 36

12.3 UI trust surfaces (required) . . . 36

12.4 Evidence drawer behavior (detailed) . . . 37

12.5 Accessibility and inclusive design . . . 37

12.6 UX patterns for abstention and restriction . . . 37

13. Governed API: contract-first enforcement boundary 38

13.1 Minimal endpoint set (aligned to refined blueprint) . . . 38

13.2 API contract requirements . . . 38

13.3 Versioning policy . . . 39

13.4 Audit and observability requirements . . . 39

14. Focus Mode: governed evidence-led AI (cite-or-abstain) 39

14.1 Focus Mode operating model . . . 39

14.2 Control loop (implementation guidance) . . . 40

14.3 Retrieval indexes and projections (buildable) . . . 40

14.4 Prompt injection and data exfiltration defenses . . . 40

14.5 UX for abstention and partial answers . . . 41

14.6 Evaluation harness (must exist before broad release) . . . 41

15. Science and discovery: from map browsing to reproducible investigations 41

15.1 Discover Mode goals (PROPOSED) . . . 41

15.2 Investigation as a first-class artifact . . . 42

15.3 Analysis patterns to support (Kansas-relevant examples) . . . 42

15.4 Compute governance (DECISION NEEDED) . . . 42

15.5 Reproducible notebooks and citations . . . 43

15.6 "Discover" UI primitives . . . 43

16. Kansas history and narrative: Story Nodes as governed publications 43

16.1 Story Node content standards (PROPOSED) . . . 43

16.2 Review workflow (publish gate) . . . 43

16.3 Story patterns (Kansas-focused examples) . . . 44

16.4 Preventing narrative drift . . . 44

16.5 Cultural sensitivity and Indigenous data sovereignty . . . 45

17. Data source strategy: anchor datasets, selection, and integration playbooks 45

17.1 Source selection rubric . . . 45

17.2 Anchor dataset shortlist (based on DSIB register) . . . 45

17.3 Source registry entry (required) . . . 46

17.4 Integration playbook template (per source) . . . 46

18. Security, privacy, and supply chain integrity 46

18.1 Authentication and authorization (DECISION NEEDED) . . . 47

18.2 Least privilege and secrets hygiene . . . 47

18.3 Sensitive location protection patterns . . . 47

18.4 PII risk and aggregation thresholds . . . 47

18.5 Supply chain integrity (recommended) . . . 48

18.6 Audit log protections . . . 48

19. Testing, QA, and promotion gates as CI enforcement 48

19.1 Test categories (minimum) . . . 48

19.2 CI gates (merge blocking) . . . 48

19.3 Quality metrics to track . . . 49

19.4 Definition of Done (release) . . . 49

20. Roadmap: smallest buildable, reversible increments 49

20.1 Phase 0: Foundations (1-2 person-weeks) . . . 49

20.2 Phase 1: First vertical slice dataset (3-6 person-weeks) . . . 50

20.3 Phase 2: Story Mode (3-6 person-weeks) . . . 50

20.4 Phase 3: Focus Mode (4-8 person-weeks) . . . 50

20.5 Phase 4: Scale-out ingestion and governance hardening (ongoing) . . . 50

21. Gap register (expanded) and closure plan 50

21.1 How to read the gap register . . . 51

21.2 Gap register (high priority) . . . 51

21.3 Gap register (medium priority) . . . 52

21.4 Gap register (long tail and future) . . . 53

21.5 How to operationalize gap closure . . . 53

22. Storage and indexing: canonical layout, projections, and rebuild strategy 53

22.1 Object storage layout (canonical) . . . 53

22.2 PostGIS projection schema (rebuildable) . . . 54

22.3 Search index design (rebuildable) . . . 54

22.4 Graph model (rebuildable, optional early) . . . 55

22.5 Vector search (optional) . . . 55

22.6 Rebuild pipeline and disaster recovery . . . 55

22.7 Migration and backwards compatibility . . . 56

23. Evidence and citation mechanics: making provenance inspectable 56

23.1 Citation principles . . . 56

23.2 EvidenceRef schemes (recommended) . . . 56

23.3 Document citations: spans and page anchors . . . 57

23.4 Feature citations (vector/raster) . . . 57

23.5 Citation verification and linting . . . 57

23.6 Evidence bundle structure (expanded) . . . 58

23.7 Evidence UX: review workflows . . . 58

24. Data quality, uncertainty, and reproducibility: closing the rigor gap 58

24.1 Quality dimensions to track . . . 59

24.2 QA report format (PROPOSED) . . . 59

24.3 Threshold policy . . . 59

24.4 Uncertainty propagation . . . 60

24.5 Data drift and freshness . . . 60

24.6 Reproducibility claims . . . 60

25. Editorial and publication standards for history narratives 60

25.1 Writing standards (PROPOSED) . . . 60

25.2 Citation density guidance . . . 61

25.3 Media handling and rights . . . 61

25.4 Sensitive topics and review triggers 61

25.5 Story Node QA checklist (publish gate) 62

25.6 Story program roadmap (example) 62

26. Operations: deployment, monitoring, backups, and incident response 62

26.1 Deployment model (PROPOSED) 62

26.2 Observability 63

26.3 Backup and restore 63

26.4 Incident response (governance incident types) 63

26.5 Cost model (PROPOSED worksheet) 64

27. Governance playbooks: making review and control repeatable 64

27.1 Rights clearance playbook (archives and media) 64

27.2 Partner data agreement playbook (restricted datasets) 65

27.3 Sensitive location release playbook 65

27.4 Story content governance playbook 65

27.5 Appeals and corrections 66

28. Implementation plan: contracts, modules, and acceptance criteria 66

28.1 Work package WP-01: Spec hashing and controlled vocab library 66

28.2 Work package WP-02: Catalog validators and link checker 66

28.3 Work package WP-03: Policy pack and fixture tests 67

28.4 Work package WP-04: Evidence resolver service 67

28.5 Work package WP-05: Dataset registry and discovery endpoints 67

28.6 Work package WP-06: Map Explorer baseline UI 68

28.7 Work package WP-07: Story Node v3 rendering and publishing 68

28.8 Work package WP-08: Focus Mode MVP with evaluation harness 68

29. Risk register (expanded) with mitigations 69

30. Glossary (expanded) 70

Appendix A. Reference repository layout (recommended) 71

Appendix B. KFM MetaBlock v2 template (no YAML frontmatter) 72

Appendix C. Dataset onboarding spec template (spec_hash input) 73

Appendix D. Run receipt (run_record) template 75

Appendix E. Story Node v3 template (markdown + MetaBlock + sidecar) 77

E.1 Story Node markdown (skeleton) 77

E.2 Story Node sidecar (map state + citations) 77

Appendix F. OPA/Rego policy template (default deny + tests) 79

F.1 Policy bundle skeleton 79

F.2 Example policy (illustrative) 79

F.3 Example tests 79

Appendix G. OpenAPI contract fragments (illustrative) 81

Appendix H. Controlled vocabulary starter lists 82

H.1 policy_label (starter) 82

H.2 artifact.zone 82

H.3 citation.kind 82

Appendix I. Data source registry starter (expanded) 84

Appendix J. Layer configuration schema and example 87

J.1 Layer config schema (starter) 87

Appendix K. Map state schema (for Story Nodes and saved views) 88

Appendix L. DatasetVersion diff report format (what changed?) 89

Appendix M. Focus Mode evaluation harness (golden queries) 90

M.1 Golden query format 90

M.2 Harness checks (minimum) 90

Appendix N. CI workflow snippet (promotion gates) 91

Appendix O. DCAT dataset record example (JSON-LD) 92

Appendix P. STAC collection and item examples 93

P.1 STAC collection 93

P.2 STAC item (example) 93

Appendix Q. PROV bundle example (JSON-LD) 95

Appendix R. Promotion manifest template 96

Appendix S. Threat model checklist (trust membrane and policy) 97

Appendix T. References 97

1. Purpose, scope, and how to use this guide

Kansas Frontier Matrix (KFM) is designed as a governed, evidence-first, map-first, time-aware knowledge system for Kansas history and science. It is not "just a map" and it is not "just a data catalog." It is a system where every public-facing claim, layer, story, and AI-assisted answer is traceable to versioned sources, with policy controls that are enforced in CI and at runtime.

This guide is a consolidation and refinement of the current KFM blueprint set, plus a structured gap-closure expansion. It is written so you can:

- Build a working vertical slice (ingest -> catalogs -> governed API -> map UI -> evidence drawer -> story citations -> Focus Mode cite-or-abstain) without redesign churn.
- Operate KFM as a governed system (FAIR + CARE) with auditable promotion gates and fail-closed defaults.
- Scale safely from a few anchor datasets and stories to a multi-domain Kansas knowledge base.

1.1 Source-of-truth documents used

This guide is grounded in the following internal KFM documents (referenced throughout by short code):

- [IB-v5] KFM Implementation Blueprint v5 Expanded (2026-02-20)
- [DesignPlan-v2] KFM Detailed Design & Engineering Plan v2 (2026-02-19)
- [DocUpgrade-v2] KFM Document Upgrade Refined FINAL v2 (2026-02-20)
- [DSIB-v1] KFM Comprehensive Data Source Integration Blueprint v1 (2026-02-12)
- [SupportFile] Kansas Frontier Matrix Support File (compiled extracts; date varies)

External standards are referenced for interoperability (DCAT, STAC, PROV-O, RFC 8785, etc.). When a standard is used, the guide recommends a concrete "KFM profile" so that schema drift is controlled.

1.2 Normative language and tagging

This document uses the following tags to avoid accidental overreach:

- CONFIRMED: stated in the source-of-truth KFM documents above.
- PROPOSED: an implementable recommendation that closes a gap, but requires adoption.
- UNKNOWN / DECISION NEEDED: not confirmed in the documents; must be verified or decided.

To keep the system buildable, every UNKNOWN is paired with a recommended default path and a minimum verification step.

1.3 Non-goals (so the scope remains buildable)

KFM can expand indefinitely. To stay buildable, the first release should avoid:

- Trying to model every historical interpretation as a "fact."
- Mirroring or redistributing content with unclear licensing (treat as metadata-only references until cleared).
- Building a full 3D globe as the primary surface (2D map-first; 3D is optional).
- Running "general chatbot" experiences. Focus Mode is a governed evidence workflow, not open-ended chat.

2. System north stars and invariants (do not compromise)

KFM credibility comes from invariants that are enforced, not only described. The documents repeatedly emphasize that governance intent must become merge-blocking gates and runtime enforcement, or the project becomes "docs-only" and trust collapses.

2.1 The "truth path" lifecycle (CONFIRMED)

KFM uses a governed data lifecycle (Raw -> Work/Quarantine -> Processed -> Catalog/Lineage) that feeds runtime stores and UI surfaces. The ordering matters: catalogs and provenance are produced as part of promotion, and runtime services read only promoted artifacts.

A useful mental model:

- Raw: immutable acquisition snapshots + checksums.
- Work/Quarantine: intermediate transforms, QA, normalization, and redaction candidates.
- Processed: publishable artifacts (GeoParquet/COG/PMTiles/media) plus checksums.
- Catalog/Lineage: cross-linked STAC/DCAT/PROV and run receipts that describe what was produced and why.
- Published: governed APIs (policy-enforcing) and user interfaces (map, story, focus).

2.2 Trust membrane (CONFIRMED)

The trust membrane is the security and governance boundary:

- Frontend and external clients never access databases or object storage directly.
- Core backend logic never bypasses repository interfaces to talk directly to storage.
- All access flows through governed APIs that apply policy decisions, redactions, and logging consistently.

If the trust membrane is broken, policy cannot be enforced and provenance cannot be trusted.

2.3 Evidence-first UX (CONFIRMED concept; implementation PROPOSED)

The user-facing contract is: every map layer and story claim can be opened into an evidence view that shows:

- Dataset version (immutable ID)
- License and rights holder
- Policy label and any redactions/generalizations applied
- The provenance chain (run receipt / activity)
- Artifact links and checksums

The evidence drawer is not a "details" panel. It is the primary trust surface.

2.4 Cite-or-abstain Focus Mode (CONFIRMED)

Focus Mode is not a general chatbot. It is a governed workflow that:

1. selects admissible evidence based on policy,
2. constructs an evidence bundle, and
3. produces an answer that cites that bundle or abstains.

Every query produces an audit record (run receipt) that is reproducible.

2.5 Canonical vs rebuildable stores (CONFIRMED concept)

KFM should treat object storage + catalogs + provenance as the canonical source of truth. Database and search indexes are projections that can be rebuilt from promoted artifacts. This reduces operational risk and improves reproducibility.

2.6 Deterministic identity and hashing (CONFIRMED)

Reproducibility depends on stable dataset identity and stable version identity. Deterministic hashing (spec_hash) based on canonical JSON (RFC 8785 JCS) is a recommended pattern to prevent "hash drift" and to support signature verification and caching.

3. Key KFM concepts and definitions

This section defines the terms used across pipelines, catalogs, APIs, and UI. Use these definitions in schemas, tests, and documentation so that contributors do not re-invent vocabulary.

3.1 Dataset, DatasetVersion, and artifacts

- Dataset: a logical dataset identity, such as "USGS NWIS (Kansas)" or "NOAA Storm Events."
- DatasetVersion: an immutable version of a dataset that corresponds to a specific promoted output set. A dataset can have many versions over time.
- Artifact: a concrete file or object (GeoParquet, PMTiles, COG, JSONL, PDF) produced by a run and referenced by STAC/DCAT/PROV.

KFM convention (PROPOSED but aligned to sources): every artifact is addressed by digest (sha256 or similar), and every catalog record points to digest-addressed artifacts.

3.2 EvidenceRef and EvidenceBundle

- EvidenceRef: a stable reference to evidence, using explicit schemes such as dcat://, stac://, prov://, doc://, graph://. The evidence resolver must be able to resolve EvidenceRef in bounded calls.
- EvidenceBundle: the resolved evidence view returned by the evidence resolver. It contains both human-readable and machine-readable fields, plus policy decision results.

Evidence bundles are immutable by digest, which supports caching and reproducibility.

3.3 Story Nodes

Story Nodes are versioned narratives bound to map state. A Story Node version includes:

- narrative markdown,
- citations (EvidenceRefs),
- a saved map state (camera, layers, time window, filters),
- policy label and review state.

Stories are reproducible because the map state and the dataset versions are captured.

3.4 Policy labels and redaction obligations

A policy label is the primary access-control and sensitivity classification input. Policy evaluation returns:

- allow/deny decision,
- obligations (redaction/generalization steps required),
- reason codes (for auditing and UX).

Obligations are key: many situations require publish-safe generalized outputs rather than a hard deny.

3.5 Run receipts and audit ledger

Every pipeline run and every Focus Mode query should emit a run receipt (run_record) containing:

- inputs (by digest and/or upstream version),
- outputs (artifact digests),
- environment (container image digest, parameters),
- validation results,
- policy decisions.

The audit ledger is append-only and must be treated as a governed dataset itself (redactions applied as needed).

3.6 Time: event time, valid time, and transaction time (PROPOSED)

Historical and scientific data often need more than one time axis:

- event time: when something happened (e.g., flood peak).
- valid time: the time a statement is considered true (e.g., county boundary existed from X to Y).
- transaction time: when KFM recorded or published the data.

Bitemporal or tritemporal modeling may be needed for some domains. Start with event time + transaction time, and introduce valid time where boundary changes and administrative history require it.

4. Architecture reference model (map-first, time-aware, governed)

This section defines the buildable architecture that respects KFM invariants and clean layering.

4.1 Layering and responsibilities (CONFIRMED posture)

Maintain clean layers:

- Domain: pure domain models (Dataset, Place, Event, EvidenceBundle) and rules.
- Use cases: workflows (ingest dataset, promote dataset version, publish story node, answer focus query).
- Interfaces: contracts (OpenAPI DTOs, schema registries, policy adapters, repository interfaces).
- Infrastructure: PostGIS, graph DB, search index, object storage, CI, deployment.

The key rule: domain logic never talks directly to infrastructure. It talks through interfaces.

4.2 Component decomposition (PROPOSED as a build plan)

A buildable decomposition for vNext:

1. Pipeline runner + connectors - watchers, fetchers, normalizers, validators, artifact writers
2. Catalog generator - produces DCAT/STAC/PROV + run receipts from pipeline outputs
3. Policy engine (OPA or equivalent) - shared semantics in CI and runtime (fixtures-driven tests)
4. Evidence resolver service - resolves EvidenceRefs to EvidenceBundles; applies policy and redaction obligations
5. Governed API gateway - dataset discovery, tile serving, feature queries, story nodes, focus query

6 UI (Map Explorer + Story Mode + Focus Mode) - evidence drawer + policy notices + version badges
7 Index builders (rebuildable projections) - PostGIS, graph, search/vector, tile caches

4.3 Canonical vs rebuildable stores (design decision)

Recommended baseline:

- Canonical:
- object storage (raw/work/processed artifacts)
- catalogs (DCAT/STAC/PROV profiles)
- audit ledger (append-only log)
- Rebuildable projections:
- PostGIS tables derived from processed GeoParquet
- search index derived from processed texts and metadata
- graph edges derived from catalogs + entity resolution
- tile bundles (PMTiles) derived from processed features

This reduces the cost of migration and lets you re-index without losing source truth.

4.4 Data flow diagram (conceptual)

```
Upstream sources
  -> connectors (fetch/snapshot)
  -> RAW zone (immutable artifacts + checksums)
  -> WORK/QUARANTINE (normalize + QA + redaction candidates)
  -> PROCESSED zone (publishable artifacts)
  -> CATALOG/TRIPLET (DCAT + STAC + PROV + run receipts)
  -> Index builders (PostGIS + search + graph + tiles)
  -> Governed API (policy enforcement + evidence resolver)
  -> UI surfaces (Map + Story + Focus)
```

4.5 Deployment topology (PROPOSED default)

For early stability, prefer a modular monolith backend with clearly separated modules:

- api (governed endpoints)
- policy (OPA adapter + fixtures)
- evidence (resolver)
- catalog (parsers/validators)
- ingest (connectors + runner)
- indexers (build projections)

This can later split into services when scaling requires it. Starting with too many services increases operational complexity and slows policy alignment.

5. Data lifecycle zones and the Promotion Contract

The Promotion Contract is the mechanism that turns governance intent into enforceable behavior. Promotion is the act of moving from Raw/Work to Processed + Catalog/Lineage, and therefore into the runtime surfaces.

5.1 Zone definitions (refined)

RAW zone (immutable acquisition)

RAW contains:

- acquisition manifest (what was fetched, from where, under what terms)
- raw artifacts (original files, API responses, scrape snapshots)
- checksums for every raw artifact
- minimal metadata (time fetched, source, license/terms snapshot)

RAW is append-only. You do not edit RAW; you supersede it with a new acquisition.

WORK / QUARANTINE (intermediate transforms)

WORK contains:

- normalized representations (e.g., parsed CSV to parquet, OCR outputs)
- QA reports (schema checks, spatial checks, completeness checks)
- candidate redactions/generalizations
- entity resolution outputs (provisional)

QUARANTINE is for:

- failed validation,
- unclear licensing,
- sensitivity concerns,
- upstream instability that prevents reproducible acquisition.

Quarantined items are not promoted.

PROCESSED (publishable)

PROCESSED contains:

- publishable artifacts in KFM-approved formats (GeoParquet, PMTiles, COG, text corpora, etc.)
- checksums for every processed artifact
- derived metadata for runtime (e.g., bounding boxes, temporal ranges, counts)

CATALOG/TRIPLET (DCAT + STAC + PROV)

The "triplet" is KFM's interoperability and evidence surface:

- DCAT: dataset-level metadata (license, publisher, distribution, themes)
- STAC: asset-level metadata for spatiotemporal artifacts (collections, items, assets)
- PROV: lineage for how artifacts were created (activities, agents, entities)

PUBLISHED (governed runtime)

Published surfaces (API + UI) may only serve promoted dataset versions that have:

- processed artifacts
- validated catalogs
- run receipts
- policy label assignment

5.2 Promotion Contract v1 (minimum gates)

A dataset version promotion MUST be blocked unless all required artifacts exist and validate. The following gates are the minimum credible set.

Gate A: Identity and versioning

- Dataset ID is stable and follows the naming convention.
- DatasetVersion ID is immutable and derived from a stable spec_hash.

Gate B: Licensing and rights metadata

- License is explicit and compatible with intended use.
- Rights holder and attribution requirements are captured.
- If license is unclear, dataset stays in QUARANTINE (fail closed).

Gate C: Sensitivity classification and redaction plan

- policy_label is assigned (public, restricted, etc.).
- For sensitive-location or restricted data, a redaction/generalization plan exists and is recorded in PROV.

Gate D: Catalog triplet validation

- DCAT record exists and validates against KFM DCAT profile.
- STAC collection/items exist (if applicable) and validate against KFM STAC profile.
- PROV bundle exists and validates against KFM PROV profile.
- Cross-links between DCAT, STAC, PROV are present and resolvable.

Gate E: Run receipt and checksums

- run_receipt exists for each producing run.
- inputs and outputs are enumerated with checksums.
- environment is recorded (container image digest, parameters).

Gate F: Policy tests and contract tests

- OPA policy tests pass for this dataset version (fixtures-driven).
- Evidence resolver can resolve at least one EvidenceRef for the dataset version in CI.
- API contracts and schemas validate.

Gate G: Optional but recommended (production posture)

- SBOM and build provenance for pipeline images and API/UI artifacts
- performance smoke checks (tile rendering, evidence resolve latency)
- accessibility smoke checks for UI (keyboard navigation for evidence drawer)

5.3 Definition of Done (dataset integration ticket)

A dataset integration ticket is DONE only when:

- RAW acquisition is reproducible and documented.
- WORK transforms are deterministic (same inputs -> same outputs, same spec -> same hash).
- PROCESSED artifacts are present in approved formats and are digest-addressed.
- Catalog triplet validates and is cross-linked.
- EvidenceRefs resolve and render in the UI evidence drawer.
- Policy label is assigned, with documented review.
- Changelog entry explains what changed and why.

5.4 A practical promotion workflow (PR-based)

PROPOSED workflow:

1. Contributor opens PR adding: - source registry entry - pipeline spec - fixture data (small sample) and expected outputs
2. CI runs: - schema validation - policy tests - spec_hash stability test - link check for catalogs
3. Steward review: - licensing and sensitivity - approve policy label
4. Operator merges and triggers pipeline run in controlled environment.
5. Outputs are written to processed + catalogs.
6. Release manifest is created and tagged.

This PR-based workflow ensures that promotion is social + technical, not ad hoc.

6. Identity, naming, controlled vocabularies, and hashing

KFM must make identity predictable. Without stable identity, you cannot cache, cite, or reproduce.

6.1 Identifier families (recommended)

Use explicit URI-like identifiers with stable prefixes. Example families:

- kfm://dataset/
- kfm://dataset/@
- kfm://artifact/sha256:
- kfm://run/
- kfm://evidence/
- kfm://story/@

Avoid embedding environment-specific hostnames in canonical IDs. Hostnames belong in distribution URLs.

6.2 Dataset slug conventions (PROPOSED)

Rules:

- lowercase
- words separated by underscore
- include upstream authority when helpful
- do not include date in dataset slug (date belongs to version)

Examples:

- usgs_nwis_kansas
- noaa_ncei_storm_events
- fema_disaster_declarations
- ipums_nhgis_county_time_series
- blm_glo_land_patents
- kshs_kansas_memory_metadata_only

6.3 DatasetVersion identity via spec_hash (CONFIRMED pattern; refine)

A DatasetVersion should be derived from a canonical specification document that includes:

- upstream source configuration (endpoints, query parameters)
- normalization rules
- validation rules
- output artifact plan
- policy label intent (proposed label)
- expected cadence

Then compute:

- spec_hash = sha256(RFC8785_canonical_json(spec))
- dataset_version_id = "."

Important: the spec_hash must be stable across platforms. This requires canonical JSON (RFC 8785) and a test.

6.4 Controlled vocabularies (minimum set)

KFM needs controlled vocabularies for governance and interoperability. Start with:

policy_label (access + sensitivity)

PROPOSED starter set:

- public: safe to show publicly
- public_generalized: derived public version of sensitive data (generalized geometry)
- restricted: requires authorization; not public
- restricted_sensitive_location: precise locations protected; default deny
- internal: visible to operators/stewards only
- embargoed: time-limited restriction pending release
- quarantine: not promoted (validation or rights unresolved)

artifact.zone (data lifecycle)

- raw
- work
- processed
- catalog
- published (for API bundles)

citation.kind (evidence types)

- dcat
- stac
- prov
- doc
- graph
- oci (optional future)
- url (discouraged; prefer resolvable schemes)

geometry.generalization_method

- centroid_only
- grid_aggregation_
- random_offset_

- dissolve_to_admin_unit
- bounding_box_only
- none

6.5 Hash and digest policy

PROPOSED:

- sha256 is mandatory for artifacts and bundles.
- use sha256 for spec_hash as well.
- keep algorithm in ID to support future migration (sha256:...).

Do not compute hashes over "pretty printed" JSON. Always canonicalize.

6.6 "Hash drift" prevention checklist

- Store the canonical spec used for hashing next to the computed spec_hash.
- Unit-test that recomputing spec_hash from stored spec yields the same value.
- Treat spec_hash changes as breaking changes that require review (why did spec change?).
- Never compute spec_hash from data that depends on system clocks, random seeds, or non-deterministic ordering.

Determinism is a feature, not overhead. It is how KFM earns trust.

7. Governance and ethics: FAIR + CARE implemented as gates

KFM is a governance system as much as a software system. Governance must be operational: encoded in schemas, policy tests, review workflows, and audit logs.

7.1 FAIR + CARE operationalization

FAIR (Findable, Accessible, Interoperable, Reusable) should be implemented as:

- Findable: DCAT/STAC catalog search; stable IDs; metadata completeness scoring.
- Accessible: governed API with clear access rules; predictable error model.
- Interoperable: profiles for DCAT/STAC/PROV; GeoParquet/COG conventions.
- Reusable: explicit license/rights; provenance chain; uncertainty metadata.

CARE (Collective Benefit, Authority to Control, Responsibility, Ethics) should be implemented as:

- Authority to Control: policy labels and steward review; partner agreements recorded.
- Responsibility: default deny for sensitive materials; redaction obligations; audit trails.
- Ethics: narrative review triggers; community consultation pathways; harm minimization.
- Collective Benefit: publish-safe derivatives, educational stories, transparent attribution.

7.2 Roles and responsibilities (gap closure)

Some KFM documents mark roles and responsibilities as a gap that must be defined. Adopt a minimal governance model first, then evolve.

PROPOSED baseline roles

- Public user: reads public layers/stories; Focus Mode limited to public evidence.
- Contributor: proposes datasets/stories; drafts content; cannot publish.
- Reviewer/Steward: approves promotions and story publishing; owns policy labels and redaction rules.
- Operator: runs pipelines and manages deployments; cannot override policy gates.
- Governance council / community stewards: authority to control culturally sensitive materials; sets rules for restricted collections and public representations.

RACI (minimum)

- Dataset onboarding:
- Responsible: contributor (spec + docs), data engineer (pipeline), GIS engineer (spatial QA)
- Accountable: steward
- Consulted: governance council (if culturally sensitive), legal/compliance (if rights unclear)
- Informed: operator

- Dataset promotion:
- Responsible: operator (run), data engineer (validate outputs)
- Accountable: steward
- Consulted: governance council (sensitive), security (restricted infrastructure)
- Informed: contributor

- Story publishing:
- Responsible: contributor (draft), historian/editor (review)
- Accountable: steward
- Consulted: governance council (Indigenous/cultural), legal (image reuse)
- Informed: public

- Policy changes:
- Responsible: steward + policy engineer
- Accountable: governance council or designated owner
- Consulted: operators (runtime impact), contributors (workflow impact)
- Informed: users

7.3 Governance artifacts (what must exist)

To operate governance, KFM must maintain:

- policy bundle repository (OPA/Rego or equivalent)

- test fixtures for policy decisions (allow/deny + obligations)
- licensing classification rubric
- sensitivity rubric and generalization guidelines
- review workflow definition (Promotion Queue + Story Review Queue)
- audit ledger retention and access policy

7.4 Policy-as-code architecture (shared CI and runtime semantics)

KFM requires the same policy semantics in CI and runtime (or at minimum the same fixtures and outcomes). Otherwise, CI guarantees are meaningless.

Recommended pattern:

- Policy Decision Point (PDP): OPA running in-process or as a sidecar.
- Policy Enforcement Points (PEP):
- CI: schema validation + policy tests block merges.
- Runtime API: policy checks before serving data.
- Evidence resolver: policy checks before resolving evidence and rendering bundles.
- UI: policy badges and notices; but UI never makes policy decisions.

7.5 Sensitivity default rules

PROPOSED defaults aligned to KFM posture:

- Default deny for sensitive-location and restricted datasets.
- If any public representation is allowed, produce a separate public_generalized dataset version.
- Never leak restricted metadata in 403/404 responses.
- Do not embed precise coordinates in Story Nodes or Focus Mode outputs unless policy explicitly allows.
- Treat redaction/generalization as a first-class transform recorded in PROV.

7.6 Licensing and rights enforcement

Key principle: "online availability does not equal permission to reuse." For archives like Kansas Memory, KFM must encode rights metadata and enforce it in downloads, exports, and UI display.

Operational rules:

- Promotion gate requires license + rights holder for every distribution.
- "Metadata-only reference" mode is allowed: you can catalog an item without mirroring it if rights do not allow.
- Export functions must include attribution and license text automatically.
- Story publishing gate blocks if rights are unclear for included media.

Licensing is not paperwork. It is a policy input.

8. Catalogs and provenance as contract surfaces

KFM treats catalogs not as "nice metadata," but as the canonical interface between pipeline outputs and runtime. This aligns with FAIR and supports evidence resolution.

8.1 The catalog triplet responsibilities

- DCAT answers: "What is this dataset? Who published it? What is the license? What are the distributions?"
- STAC answers: "What assets exist? What are their spatiotemporal extents? Where are the files?"
- PROV answers: "How were these outputs created? Which inputs, which tools, which parameters?"

KFM should define a profile for each so that validation is strict and predictable.

8.2 KFM DCAT profile (PROPOSED minimum)

Minimum required fields:

- dct:title
- dct:description
- dct:publisher
- dct:license (or dct:rights)
- dcat:theme (controlled vocabulary)
- dct:spatial and dct:temporal coverage
- dcat:distribution (one per artifact class)
- prov:wasGeneratedBy link to PROV activity bundle
- kfm:policy_label
- kfm:dataset_id and kfm:dataset_version_id

8.3 KFM STAC profile (PROPOSED minimum)

Minimum required for a collection:

- id, title, description
- extent (spatial bbox; temporal interval)
- license
- links to DCAT dataset record
- kfm:dataset_version_id and policy label

Minimum required for an item:

- id
- geometry (or bbox) consistent with policy label (generalized if needed)

- datetime (or start/end)
- assets with href + checksum + media_type
- links to PROV run receipt and DCAT distribution

8.4 KFM PROV profile (PROPOSED minimum)

Minimum required:

- prov:Activity per pipeline run
- prov:Entity per artifact (raw, work, processed)
- prov:Agent for pipeline and steward approval events
- prov:used and prov:wasGeneratedBy edges
- kfm:policy_decision references (decision_id + obligations)
- Environment capture: container image digest, git commit, parameters

8.5 Cross-linking rules (must be testable)

Define explicit cross-links so navigation is deterministic:

- DCAT dataset -> distributions -> artifact digests
- DCAT dataset -> prov:wasGeneratedBy -> PROV bundle
- STAC collection -> link rel="describedby" -> DCAT dataset
- STAC item -> link to PROV activity and/or run receipt
- EvidenceRef schemes should resolve into these objects without guessing.

CI should include a link-checker that verifies cross-links for every promoted dataset version.

8.6 Evidence resolver contract (refined)

Evidence resolution is central. The evidence resolver must:

- accept EvidenceRef (scheme://...) or a structured reference (dataset_version + record id + span)
- apply policy and return allow/deny + obligations
- return an EvidenceBundle with:
- human view (renderable card)
- machine metadata (JSON)
- artifact links (only if allowed)
- digests and dataset_version ids
- audit references

The evidence resolver must be usable in <= 2 calls from the UI:

- click feature -> resolve evidence -> view bundle
- click citation -> resolve evidence -> view same bundle

If evidence takes more steps, users ignore provenance.

8.7 Catalog and evidence validation in CI

Minimum CI checks:

- JSON schema validation for KFM DCAT/STAC/PROV profiles
- link check: cross-links exist and resolve in repo context
- evidence resolver contract tests:
- "public" evidence resolves to bundle with allowed artifacts
- "restricted" evidence returns 403 with no sensitive metadata leakage
- spec_hash stability tests
- golden tests for canonicalization and deterministic outputs

Catalog validation is what makes interoperability real and prevents silent drift.

9. Domain modeling: connecting Kansas history, science, and place

KFM spans historical narrative and scientific observation. The domain model must support both without forcing everything into one brittle ontology.

9.1 A pragmatic core entity set (PROPOSED)

Start with a small set of core entities and expand by profile:

- Place: a spatially bounded concept (county, town, river reach, reservation boundary, survey section).
- Event: a time-bounded occurrence (flood, tornado, treaty signing, railroad arrival).
- Observation: a measured or recorded datapoint (streamflow reading, census count).
- Artifact: a file/object (map scan, newspaper page, GeoParquet table).
- Dataset / DatasetVersion: versioned collection.
- Agent: person or organization responsible for actions (publisher, pipeline, steward).
- Claim: a statement used in narratives, linked to evidence.
- EvidenceBundle: resolved evidence for a claim or feature.
- Relationship: typed edge between entities (located_in, part_of, affects, documented_by, etc.).

A key principle: do not require full semantic completeness for early builds. You can model minimal relationships and add more later.

9.2 Spatiotemporal representation

Geometry types

- Points, lines, polygons, multipolygons, and geometry collections.
- For sensitive locations, store precise geometry only under restricted policy; store generalized geometry separately.

Time types

Support:

- instant (datetime)
- interval (start, end)
- uncertain time (approx, range, or "circa")

For historical data, include:

- time_precision: year, month, day, unknown
- time_source: which evidence supports the time assignment

9.3 Uncertainty as first-class metadata (PROPOSED but recommended)

To close the "rigor gap," include uncertainty metadata in datasets and stories:

- spatial_uncertainty_m (or categorical)
- temporal_uncertainty (days or categorical)
- classification_confidence (0-1) for automated extraction
- provenance_quality_notes

Uncertainty is not a weakness. It is honesty that improves trust.

9.4 Boundary changes and bitemporality (DECISION NEEDED by domain)

Some datasets require modeling boundary changes over time (counties, districts, reservations). Two approaches:

- Option A (simpler): store a boundary dataset version per change period; stories reference the correct version.
- Option B (more advanced): valid-time attributes on boundary features; query by time range.

Start with Option A for MVP, move to Option B if needed.

9.5 Entity resolution strategy (PROPOSED)

Entity resolution must be provenance-aware:

- never merge entities without recording evidence and decision rationale
- keep "candidate match" edges separate from "canonical sameAs" edges
- support manual steward override with audit record

A recommended pattern:

- graph edges:
- kfm:match_candidate (score, method, evidence)
- kfm:same_as (asserted by steward, evidence bundle)
- kfm:related_to (typed)

9.6 Claim model for narratives (PROPOSED)

Narratives contain claims. Each claim should be:

- atomic enough to cite (one claim -> one or more citations)
- linked to EvidenceRefs
- assigned a claim class:
- observation_claim (data-based)
- documentary_claim (archival source)
- interpretive_claim (interpretation; may cite multiple sources)

Interpretive claims require explicit labeling and review triggers. This prevents presenting interpretation as fact.

9.7 Mapping "feature -> evidence" in spatial layers

Every feature served to the UI must have a stable feature_id and an evidence_ref pointer. That evidence_ref must be resolvable via the evidence resolver and render the evidence drawer.

This is the backbone for trust in the map-first UI.

10. Ingestion, normalization, and validation pipelines

KFM ingestion is not a one-off ETL. It is a governed pipeline framework that produces reproducible artifacts, catalogs, and receipts.

10.1 Connector patterns (PROPOSED)

Common patterns:

- Bulk download connector: fetches zip/CSV/shapefiles and snapshots them.
- API connector: calls paginated APIs and stores raw responses plus request metadata.
- Scrape connector: used only when official APIs are unavailable; store HTML snapshots and parsing scripts.
- Document ingest connector: pulls PDFs/images/maps and stores them with metadata.
- Streaming connector: used for real-time sources (mesonet, sensors), with windowed snapshots.

Connector requirements:

- record source URL or endpoint
- record timestamp and request parameters
- store a "terms snapshot" if license/terms are not stable
- emit an acquisition manifest

10.2 Normalization patterns by data class

Tabular (CSV, JSON)

- convert to parquet (with schema)
- normalize field names and units
- store mapping from upstream fields to canonical fields
- validation: schema, completeness, value ranges, duplicates

Vector geospatial

Preferred processed format: GeoParquet (with geo metadata). GeoJSON is allowed only for small curated artifacts or debug exchange under size thresholds.

Normalization steps:

- enforce CRS (store original; reproject to standard for processed)
- validate geometry (no self-intersections, valid polygons)
- simplify/generalize if policy requires
- compute bounding boxes and temporal extents

Raster

Preferred processed format: Cloud Optimized GeoTIFF (COG) with overviews.

Normalization steps:

- verify georeferencing
- build overviews
- record nodata and units
- compute statistics for visualization

Documents (PDF, scans, newspapers)

Treat documents as first-class evidence datasets:

- store page images as assets (or references if rights restrict mirroring)
- OCR text is a derived artifact, with quality metrics
- citations must resolve back to page images and OCR spans (character offsets or bounding boxes)

10.3 Validation gates (minimum)

Minimum validation per dataset version:

- schema validation (expected fields and types)
- spatial validation (CRS, bounds, geometry validity)
- temporal validation (date ranges, precision)
- rights metadata validation (license present)
- sensitivity scan (detect restricted categories)
- link validation for catalogs

Validation outputs should be machine-readable (JSON) and part of run receipts.

10.4 Deterministic pipeline design

Determinism requirements:

- explicit versioned dependencies (container images pinned by digest)
- no random ordering (sort inputs deterministically)
- no clock-dependent outputs (record timestamps separately; avoid embedding in file content where possible)
- stable floating point handling when generating tiles

Every run should be reproducible given:

- raw inputs by digest
- pipeline spec by spec_hash
- environment by image digest

10.5 Quarantine workflow (fail closed)

Quarantine triggers:

- license unclear
- validation fails
- sensitivity unclear
- upstream instability prevents reproducible acquisition

Quarantine requires:

- a reason code
- a remediation plan (who resolves, by when)
- a policy default (deny publish)

Do not "temporarily promote" quarantined datasets.

11. Mapping and cartographic design (map-first, time-aware)

KFM is experienced through the map and timeline. The map design must make trust and uncertainty visible.

11.1 Layer taxonomy (recommended)

Define layer classes so users understand what they are seeing:

- Basemap/reference: authoritative geography, elevation, hydrography
- Observations: measured data (streamflow, weather, census)
- Events: discrete occurrences (tornado tracks, flood extents, treaty sites)
- Administrative history: boundaries over time, jurisdictions
- Narrative overlays: story-specific annotations
- Restricted overlays: only for authorized roles (never delivered to public)

Each layer must show:

- dataset version
- policy label
- license/attribution
- freshness (last run)
- evidence drawer link

11.2 Time-aware interaction patterns

Time-aware UI patterns:

- time range slider (start/end)
- timeline histogram (density of events over time)
- layer temporal availability indicator (e.g., data exists 1950-2024)
- bookmarks (saved time windows used by Story Nodes)
- "time warp" for boundary datasets (show correct version for selected time)

Avoid forcing a single time model across all layers. Some layers have event time, others have valid time.

11.3 Performance and data delivery

The Design Plan emphasizes:

- Prefer vector tiles (MVT) or PMTiles for large feature layers.
- Avoid shipping massive GeoJSON to the browser.
- Cache tiles and evidence bundles aggressively; bundles are immutable by digest.
- UI caches must be keyed by dataset_version_id (datasets are not assumed stable).

11.4 Tile strategy (DECISION NEEDED)

Two buildable options:

- Option A: server-side tile service (PostGIS -> tile server -> MVT)
- Pros: dynamic filters, styling flexibility
- Cons: more runtime complexity

- Option B: prebuilt PMTiles bundles served as static assets
- Pros: simple hosting, strong caching, reproducible
- Cons: rebuild required for changes; dynamic filtering limited

Recommended default for MVP: Option B (PMTiles) for large layers, plus a small feature query endpoint for identify-on-click. Move to Option A later if dynamic filtering becomes essential.

11.5 Symbology and uncertainty visualization

Principles:

- Color is not the only carrier of meaning (accessibility).
- Use clear legends; legends are part of the trust surface.
- Uncertainty must be visually encoded:
- fuzzy boundaries or transparency for uncertain extents
- error bars / bands for time series
- aggregated bins instead of points for generalized sensitive data

For sensitive locations, do not show "hidden" points. Show generalized representations with explicit policy notice (e.g., "Geometry generalized due to sensitive-location policy").

11.6 Map state as reproducible artifact

Map state includes:

- camera position (bbox/zoom)
- active layers + style parameters
- time window
- filters

Story Nodes store map state so that stories replay the same view. Focus Mode can accept view_state hints so it answers in context.

This is how KFM turns "maps" into reproducible evidence experiences.

12. User interface: Map Explorer, Story Mode, Focus Mode, and trust surfaces

The KFM frontend is a governed client. It renders what the API returns and never embeds privileged credentials. The UI must make governance visible rather than hidden.

12.1 Information architecture (refined)

Recommended top-level navigation:

- Map Explorer (primary)
- Timeline (optional; can be integrated into Map Explorer)
- Stories (Story Mode)
- Catalog (dataset discovery)
- Focus Mode (evidence-led Q&A;)
- Admin/Steward (restricted)

The interaction loop described in KFM sources is:

Explore (map/time) -> Focus Mode -> Evidence/Provenance inspection -> Story Nodes

12.2 Core UI components (buildable)

Map Explorer:

- MapCanvas (MapLibre GL)
- LayerPanel (toggle, opacity, legend, policy badge, data version)
- TimelineSlider / TimeControl (range + histogram)
- SearchBar (places, datasets, story nodes)
- FeatureInspectPanel (attributes + citations)
- EvidenceDrawer (shared component)

Story Mode:

- StoryNodeList
- StoryNodeReader (markdown rendering with citation hooks)
- EvidenceDrawer (same component)
- RelatedEntitiesPanel (optional)

Focus Mode:

- ChatPanel
- EvidenceSnippets (inline citations)
- PolicyNotice (why some info is withheld)
- ExportAnswer (downloadable report with audit_ref and citations)

Admin/Steward:

- PromotionQueue (dataset versions pending approval)
- QAReportViewer
- PolicyLabelEditor (limited, controlled)
- StoryReviewQueue

12.3 UI trust surfaces (required)

Trust surfaces recommended in KFM sources include:

- Automation status badges (healthy/degraded/failing) on layers or features.
- Provenance drawer / evidence drawer accessible from every layer and story claim.
- Data Version label per layer linking to DatasetVersion catalogs.
- Policy notices explicit at time of interaction ("geometry generalized due to policy").
- "What changed?" panel comparing DatasetVersion diffs (counts, checksums, QA metrics).

These are not optional polish. They are the user-visible governance contract.

12.4 Evidence drawer behavior (detailed)

The evidence drawer should show, at minimum:

- Evidence bundle ID + digest
- DatasetVersion ID + dataset name
- License and rights holder (with attribution text)
- Freshness (last run timestamp) and validation status
- Provenance chain (run receipt link)
- Artifact links (only if policy allows)
- Redactions applied (obligations)

Publishing should be blocked if any citation fails to resolve. The UI can enforce this by calling the evidence resolver during story publish checks.

12.5 Accessibility and inclusive design

Minimum accessibility requirements:

- Keyboard navigable layer controls and evidence drawer; visible focus states.
- Text labels for policy badges and status indicators (no color-only meaning).
- ARIA labels for map controls.
- Safe markdown rendering for narratives (CSP + sanitization to prevent XSS).
- Export outputs include citations and audit_ref in a readable format.

12.6 UX patterns for abstention and restriction

Abstention is a feature. The UI must make abstention understandable without leaking restricted info:

- show "why" in policy-safe terms (e.g., "restricted evidence not available to your role")
- suggest safe alternatives (broader time range, public datasets)
- provide audit_ref so stewards can review

Never show "ghost metadata" that reveals restricted existence unless policy allows.

13. Governed API: contract-first enforcement boundary

The governed API is the enforcement boundary. It applies policy decisions, redactions, versioning, and evidence resolution consistently.

13.1 Minimal endpoint set (aligned to refined blueprint)

A buildable v1 endpoint set:

- GET /api/v1/datasets
- list datasets + versions (DCAT-backed)
- supports filters, facets, search
- policy filter applied server-side

- GET /api/v1/stac/collections
- GET /api/v1/stac/items
- STAC query by bbox/time/collection
- policy filter applied before returning assets

- POST /api/v1/evidence/resolve
- resolves EvidenceRefs to EvidenceBundles
- returns bundle digest + policy decision results

- GET/POST /api/v1/story
- GET/PUT /api/v1/story/{id}
- story nodes with versioning
- publishing requires citations + review state

- POST /api/v1/focus/ask
- governed Focus Mode Q&A;
- returns citations or abstain + audit_ref

- GET /api/v1/lineage/status
- GET /api/v1/lineage/stream
- pipeline health/freshness feeds UI badges

If tiles are served dynamically:

- GET /api/v1/tiles/{layer}/{z}/{x}/{y}.pbf

If tiles are served as PMTiles bundles:

- GET /assets/pmtiles/{dataset_version_id}/{layer}.pmtiles
- served via static hosting with policy gating (or only for public layers)

13.2 API contract requirements

- Every response returns:

- dataset_version_id (when applicable)
- artifact digests (when applicable)
- policy label (public-safe)
- audit_ref for governed operations (focus, story publish)

- Errors use a stable error model:
- error_code
- message (policy-safe)
- audit_ref (for debugging)
- optional remediation hints

Avoid leaking sensitive existence through error differences. Align 403/404 behavior with policy.

13.3 Versioning policy

- Freeze /api/v1 semantics; only add backwards-compatible fields.
- Introduce /api/v2 only for breaking changes.
- Schema versioning:
- KFM-DCAT vN
- KFM-STAC vN
- KFM-PROV vN
- Story Nodes have their own template versions; migration scripts and dual-read support may be needed.

13.4 Audit and observability requirements

Every governed operation must emit a log record that includes:

- who (principal, role)
- what (endpoint, parameters)
- when (time)
- why (purpose if declared)
- inputs/outputs (by digest)
- policy decisions (allow/deny, obligations, reason codes)

Audit logs are themselves sensitive. Apply log redaction and retention policy.

14. Focus Mode: governed evidence-led AI (cite-or-abstain)

Focus Mode is KFM's AI-assisted interface, but it is constrained by policy, evidence, and audit. The system should behave like a research assistant that always shows its sources, not like a general conversational agent.

14.1 Focus Mode operating model

A Focus Mode request is treated as a governed run with a receipt.

Inputs:

- user query
- optional view_state (map bbox, time window, active layers)
- user role and policy context

Outputs:

- answer text
- citations (EvidenceRefs) that resolve to evidence bundles
- audit_ref (run id) that can be reviewed

14.2 Control loop (implementation guidance)

Recommended control loop:

1. Policy pre-check - determine if query is allowed (topic restrictions, role limitations)
2. Retrieval plan - select candidate datasets and indexes based on view_state and query intent
3. Retrieve evidence - query catalogs, search index, graph, or database for admissible evidence
4. Build evidence bundle - call evidence resolver for each EvidenceRef - apply redaction obligations
5. Synthesize answer - model generates answer referencing evidence bundle IDs
6. Citation verification (hard gate) - verify every citation resolves and is policy-allowed for this user - if not, drop citation and revise answer or abstain
7. Produce audit receipt - store query, evidence bundle digests, policy decisions, model version, latency, and output hash

The hard gate is step 6. If citations cannot be verified, the response must abstain or reduce scope. This is how KFM avoids hallucinations.

14.3 Retrieval indexes and projections (buildable)

Focus Mode may use multiple projections:

- Catalog search (DCAT/STAC) to find datasets by theme and coverage
- Text search index for OCR corpora and documents
- Graph edges for entity resolution and relationship traversal
- PostGIS for spatial filtering (bbox/time)
- Vector index (optional) for semantic retrieval

Important: retrieval results must always map back to EvidenceRefs that resolve to bundles. There is no "raw text from index" allowed without evidence linking.

14.4 Prompt injection and data exfiltration defenses

Focus Mode must resist:

- prompt injection from documents (e.g., malicious OCR text)
- data exfiltration attempts (e.g., "show me the restricted dataset list")

Defenses:

- tool allowlist (the model cannot call arbitrary tools)
- explicit system policy about refusing to reveal restricted sources
- evidence resolver as the only source of "truth" for citations
- content filtering and redaction obligations applied before model sees text for restricted datasets

14.5 UX for abstention and partial answers

Abstention UX must be clear:

- what is missing (in policy-safe terms)
- what is allowed (public alternatives)
- how to request access (steward review workflow)
- include audit_ref for follow-up

Partial answers are acceptable when only part of the question is supported by evidence.

14.6 Evaluation harness (must exist before broad release)

KFM should implement a Focus Mode evaluation harness with tests for:

- citation coverage: % of factual claims supported by citations
- citation resolvability: 100% citations resolve for allowed users
- refusal correctness: restricted questions receive policy-safe refusals
- sensitivity leakage: no restricted coordinates or metadata in outputs
- regression tests: golden queries across dataset versions

This evaluation harness should be run in CI for Focus Mode code changes and before each release.

15. Science and discovery: from map browsing to reproducible investigations

KFM should support not only narrative history but also scientific discovery workflows. "Discover" is a mode where users can form hypotheses, explore datasets, run analyses, and publish results as governed artifacts.

15.1 Discover Mode goals (PROPOSED)

- Make it easy to find relevant datasets by place and time.
- Enable exploratory analysis without exporting uncontrolled copies of restricted data.
- Produce reproducible investigation artifacts (queries, notebooks, results) with provenance and policy labels.
- Allow publishing of derived public results (aggregations, models) while protecting sensitive inputs.

15.2 Investigation as a first-class artifact

Define an Investigation artifact type:

- investigation_id
- title/description
- inputs: dataset_version ids and EvidenceRefs
- methods: scripts/notebooks (by digest)
- outputs: result artifacts (tables, maps, charts) with digests
- run receipts and environment
- policy label and review state

Investigations can be shared internally and, when safe, published publicly as Story Nodes or data products.

15.3 Analysis patterns to support (Kansas-relevant examples)

Examples of analysis workflows:

- County-level demographic change over time (NHGIS time series)
- Flood and drought history overlay with settlement expansion
- Land patent density vs railroad development (GLO + rail datasets)
- Streamflow anomalies vs land cover change (USGS NWIS + NLCD)
- Severe weather event clustering and impacts (NOAA storm events + FEMA declarations)

Each analysis should cite its inputs and produce a run receipt.

15.4 Compute governance (DECISION NEEDED)

Where does analysis run?

Options:

- Option A: controlled notebook environment on the platform (JupyterHub-like)
- Pros: governed environment, policy enforcement
- Cons: operational complexity

- Option B: offline notebooks supported by versioned dataset exports
- Pros: simple to start
- Cons: risk of uncontrolled redistribution for restricted datasets

Recommended default: start with Option B for public datasets, and Option A for restricted datasets or partner-controlled data. This respects policy while enabling early scientific work.

15.5 Reproducible notebooks and citations

Notebook requirements (PROPOSED):

- every notebook references dataset_version_id for all inputs
- every chart includes a link to the evidence bundle or dataset version
- notebooks are stored by digest and referenced in PROV as activities
- export of notebooks to PDF/HTML includes citations and policy notices

15.6 "Discover" UI primitives

Even without full notebooks, KFM can support discovery by adding:

- Saved queries (bbox/time + dataset filter)
- Saved map states (like Story Nodes but without narrative)
- Dataset compare view (two dataset versions side-by-side)
- QA metrics dashboard (completeness, spatial coverage, uncertainty)

These primitives create a bridge between casual map exploration and formal investigations.

16. Kansas history and narrative: Story Nodes as governed publications

KFM's history storytelling must be evidence-led and ethically grounded. Story Nodes are publishable artifacts that bind narrative to map state and citations.

16.1 Story Node content standards (PROPOSED)

A publishable Story Node must:

- declare scope (time window, geography)
- separate observation claims from interpretive claims
- include citations for every factual claim (EvidenceRefs)
- include uncertainty notes where sources conflict or are incomplete
- include licensing and attribution for all embedded media
- include policy label and review state

16.2 Review workflow (publish gate)

Proposed workflow:

1. Draft: contributor creates Story Node in Studio; citations must resolve in-editor.
2. Review: steward and historian/editor review claims, citations, and sensitivity.
3. Governance review: triggered if Story Node touches Indigenous histories, restricted sites, or sensitive locations.
4. Publish: Story Node becomes immutable version; edits create new version.

Publishing should be blocked if:

- citations do not resolve
- rights are unclear for included media
- sensitive locations are included without policy approval

16.3 Story patterns (Kansas-focused examples)

Suggested "story arcs" that can be built as a program:

- Land and settlement:
- homestead claims and land patents (GLO)
- survey grids (PLSS)
- railroad expansion and town growth

- Governance and boundaries:
- territorial boundaries and county formation over time
- reservation boundaries and treaty history (handled with care and partner governance)

- Environmental history:
- flood events and river change
- drought cycles
- land cover transitions

- Conflict and movement:
- migration corridors
- trade routes and trails
- newspapers and public discourse (OCR corpora)

Each arc should have a starter dataset set plus Story Node templates.

16.4 Preventing narrative drift

Risks:

- unverified claims becoming "canon"
- citations pointing to sources that do not support the claim
- interpretive framing presented as fact

Mitigations:

- claim-level citation checks
- evidence drawer always accessible from story text
- versioned story releases with changelogs
- review triggers for contested topics

16.5 Cultural sensitivity and Indigenous data sovereignty

Operational guidance:

- do not publish precise locations for culturally sensitive sites
- allow community-controlled policy labels and release criteria
- use "public_generalized" derivatives as the default public representation
- provide narrative context that respects community perspectives and avoids appropriation

If permissions are unclear, fail closed and flag for governance review.

17. Data source strategy: anchor datasets, selection, and integration playbooks

KFM requires authoritative sources treated as datasets in the governed lifecycle. Primary and official sources come first, then curated academic sources.

17.1 Source selection rubric

Choose sources that:

- are authoritative and stable
- have clear licensing terms
- have well-defined coverage (space/time)
- are useful for multiple story arcs
- can be represented with evidence and provenance

Avoid early sources that:

- have unclear reuse rights for media
- contain high-risk PII or sensitive coordinates without a clear generalization pathway
- require heavy scraping without clear permission

17.2 Anchor dataset shortlist (based on DSIB register)

A buildable anchor set for vNext:

- Demographics: IPUMS NHGIS (historical tables + boundaries)
- Land: BLM GLO land patents (public domain; names may be sensitive)
- Hydrology: USGS WaterData/NWIS (public domain)

- Hazards: NOAA Storm Events (public domain)
- Disasters: FEMA disaster declarations (public domain)
- Basemap: USGS National Map (public domain base layers)
- Kansas GIS framework: Kansas DASC Geoportal (state authoritative layers)

For archives with reuse constraints:

- Kansas Memory / Kansas Historical Society collections: treat as metadata-first until rights are cleared; mirror only what is allowed.

For sensitive heritage layers:

- archaeology sites and restricted heritage inventories: store precise geometry restricted; publish generalized derivatives or metadata-only records.

17.3 Source registry entry (required)

Every source must have a registry entry with:

- name and authority
- access method (API, bulk, scrape, manual)
- cadence
- license/terms snapshot
- sensitivity classification
- connector spec and credentials strategy (if any)
- known limitations and QA checks

The registry is not optional documentation; it is an input to promotion gates.

17.4 Integration playbook template (per source)

Each source integration should follow a repeatable playbook:

1. Source assessment - license, sensitivity, access
2. Acquisition design - connector, manifest, raw storage
3. Normalization design - schemas, mapping, units
4. Validation design - checks and thresholds
5. Output design - processed formats, tiling approach
6. Catalog design - DCAT/STAC/PROV profile mapping
7. Policy design - policy label and obligations
8. UI integration - layer definition + evidence drawer mapping
9. Focus Mode integration - retrieval hooks + citation resolver tests

This playbook is the unit of scale for KFM ingestion.

18. Security, privacy, and supply chain integrity

KFM handles sensitive locations, licensing constraints, and potentially restricted partner data. Security is not a separate layer; it is integrated into trust membrane and governance gates.

18.1 Authentication and authorization (DECISION NEEDED)

Questions to decide:

- Which identity provider (OIDC) will be used?
- What roles exist (public, contributor, steward, operator)?
- Do we require attribute-based access control (ABAC) for partner datasets?

Recommended default:

- OIDC (standard identity provider)
- RBAC with policy labels; add ABAC only when required

18.2 Least privilege and secrets hygiene

- store secrets in a secrets manager (not in repo)
- pipeline runners use scoped credentials per source
- UI never receives privileged tokens beyond what is needed for the session
- rotate secrets and record rotation events in audit logs

18.3 Sensitive location protection patterns

- store precise geometries only in restricted datasets
- produce public_generalized derivatives
- ensure that derived artifacts cannot be reverse engineered (no hidden precise points)
- enforce policy at tile serving (no bypass via static hosting)

Include tests:

- "no restricted bbox leakage" tests for public tiles
- "no coordinate fields" tests for public exports when restricted

18.4 PII risk and aggregation thresholds

Some datasets (property ownership, crime, health) have reidentification risk.

Guidance:

- do not publish individual-level records publicly
- aggregate to safe geographies and apply minimum count thresholds
- document thresholds as policy obligations
- keep raw data restricted even if aggregated outputs are public

18.5 Supply chain integrity (recommended)

To match provenance-first ethos at the software level:

- generate SBOMs (SPDX) for built images
- generate build provenance attestations (SLSA/in-toto)
- verify attestations server-side (do not let browsers fetch attestations directly)
- pin dependencies and verify checksums

This is optional for early prototypes but recommended before broader public release.

18.6 Audit log protections

Audit logs may contain sensitive operational details.

Rules:

- logs are append-only
- logs are redacted for PII and restricted info
- log access is restricted to stewards/operators
- retention and deletion policies are defined and enforced

Security is governance. Treat it as such.

19. Testing, QA, and promotion gates as CI enforcement

KFM is successful only if governance invariants are encoded as tests that fail closed.

19.1 Test categories (minimum)

- Unit tests: domain logic, spec_hash canonicalization, controlled vocab validations
- Schema tests: DCAT/STAC/PROV profile validation
- Policy tests: OPA fixture-driven allow/deny/obligation outcomes
- Contract tests: OpenAPI schema diff checks and DTO validation
- Integration tests: evidence resolver resolves sample refs
- E2E UI tests: layer toggle -> evidence drawer shows attribution; story -> citations resolve; focus -> citations resolve

19.2 CI gates (merge blocking)

PROPOSED minimum merge gates:

- lint + typecheck (frontend + backend)
- schema validation for any changed catalog artifacts
- Story Node template validation

- policy tests must pass
- spec_hash tests must pass
- link checker must pass (no broken citations)
- security scanning (dependency vulnerabilities) and optional SBOM generation
- accessibility smoke checks for UI changes (at least keyboard nav for evidence drawer)

19.3 Quality metrics to track

Track governance and reliability metrics:

- % of promoted artifacts with explicit license metadata
- % of Story Nodes with 100% resolvable citations
- evidence resolver latency (P95)
- tile serving latency (P95) for public layers
- reindex time from processed artifacts
- number of quarantined datasets (by reason code)
- number of policy denials (by reason code)

Metrics must not encourage perverse incentives. Use them to detect drift and risk.

19.4 Definition of Done (release)

A release is DONE when:

- all gates pass in CI
- promotion manifests are created for new dataset versions
- evidence resolver contract tests pass for both public and restricted scenarios
- Focus Mode evaluation harness passes golden queries
- UI regression tests pass and accessibility checks show no major regressions
- release notes include policy and data changes
- audit ledger retention and monitoring are configured

20. Roadmap: smallest buildable, reversible increments

KFM can be built as a sequence of vertical slices. Each slice should produce a demonstrable governed surface.

20.1 Phase 0: Foundations (1-2 person-weeks)

- verify repo truth inventory and directory structure
- establish baseline CI (lint, tests, schema validation hooks)
- define KFM profiles (DCAT/STAC/PROV) and controlled vocabularies
- implement spec_hash canonicalization function + tests

Deliverable: repo can validate a sample dataset version and a sample Story Node in CI.

20.2 Phase 1: First vertical slice dataset (3-6 person-weeks)

Pick one anchor dataset (e.g., NHGIS county population time series).

Build:

- source registry entry + connector
- deterministic pipeline producing processed GeoParquet + PMTiles
- catalogs (DCAT/STAC/PROV) + run receipt
- governed API endpoints: datasets, stac, evidence resolve
- Map Explorer showing the layer and evidence drawer

Deliverable: users can click a map feature and see evidence + license + dataset version.

20.3 Phase 2: Story Mode (3-6 person-weeks)

- Story Node v3 schema and rendering
- Studio draft + review flow (minimal)
- publish gate: citations must resolve
- story browsing UI

Deliverable: a story that replays map state and shows evidence for every claim.

20.4 Phase 3: Focus Mode (4-8 person-weeks)

- Focus Mode orchestrator with cite-or-abstain loop
- retrieval hooks over catalogs and search index
- citation verifier hard gate
- audit receipts + evaluation harness

Deliverable: Focus Mode answers questions about the current map view with resolvable citations or abstains.

20.5 Phase 4: Scale-out ingestion and governance hardening (ongoing)

- add more datasets using the playbook template
- add restricted datasets with generalized public derivatives
- add performance monitoring and caching strategies
- harden supply chain integrity

KFM scales by repeating the playbook, not by redesigning the core.

21. Gap register (expanded) and closure plan

This section enumerates the highest-impact gaps implied by the current KFM document set and provides a closure plan for each. Use this as an actionable backlog and governance checklist.

21.1 How to read the gap register

Each gap includes:

- Gap ID
- Domain
- Status (CONFIRMED gap, PROPOSED closure, UNKNOWN)
- Risk if unclosed
- Recommended default (to keep building)
- Minimum verification step

21.2 Gap register (high priority)

```
GAP-001 Governance roles and decision rights
  Status: CONFIRMED gap (roles not fully defined)
  Risk: promotion and publishing becomes arbitrary; policy bypass
  Default: adopt baseline roles (contributor, steward, operator, council) + RACI in Section 7
  Verify: publish a roles matrix in repo; assign CODEOWNERS for policy and catalogs

GAP-002 Authn/Authz and identity provider selection
  Status: DECISION NEEDED
  Risk: insecure access or inability to support partners
  Default: OIDC + RBAC + policy_label; introduce ABAC only when needed
  Verify: choose IdP and implement /whoami endpoint + token scopes

GAP-003 Tile serving strategy (PMTiles vs dynamic tile service)
  Status: DECISION NEEDED
  Risk: performance failure or operational complexity
  Default: PMTiles for large public layers; dynamic only when required
  Verify: benchmark 1-2 layers at target size; decide by P95 latency and cost

GAP-004 EvidenceRef schemes and resolver guarantees
  Status: CONFIRMED as core invariant; details vary
  Risk: citations become non-resolvable; Focus Mode credibility collapses
  Default: implement evidence resolver contract with dcat://, stac://, prov://, doc:// schemes
  Verify: CI tests that resolve at least 5 refs across data classes

GAP-005 Policy label vocabulary and obligation model
  Status: CONFIRMED need; exact labels undecided
  Risk: inconsistent enforcement; sensitive leakage
  Default: adopt starter vocabulary in Section 6 and obligation fields in policy decisions
  Verify: policy fixtures for each label; UI badges match

GAP-006 Rights metadata granularity (dataset vs item vs asset)
  Status: DECISION NEEDED
  Risk: misuse of media with reuse constraints (Kansas Memory class)
  Default: capture rights at distribution/asset level; allow metadata-only catalogs
  Verify: implement rights fields in DCAT and STAC assets; enforce in export paths
```

```
GAP-007 Sensitive location generalization standards
  Status: CONFIRMED need
  Risk: reverse engineering of restricted locations; legal/ethical harm
  Default: require public_generalized derivatives; forbid public delivery of precise coords
  Verify: redaction tests on tile outputs + export outputs

GAP-008 Boundary change modeling (valid time)
  Status: DECISION NEEDED by domain
  Risk: time slider shows incorrect boundaries; historical claims wrong
  Default: versioned boundary datasets per change period; upgrade to valid time later
  Verify: pilot on county boundaries; write story demonstrating correctness

GAP-009 Entity resolution governance (who can assert sameAs)
  Status: DECISION NEEDED
  Risk: ontology drift and wrong merges
  Default: only stewards assert sameAs; automated matches remain candidates
  Verify: graph schema includes match_candidate vs same_as edges

GAP-010 Focus Mode evaluation harness and release criteria
  Status: CONFIRMED need
  Risk: hallucinations and restricted leakage
  Default: implement golden queries + citation resolvability tests before any public release
  Verify: CI job that runs harness and blocks merge on regressions
```

21.3 Gap register (medium priority)

```
GAP-011 Store selection for graph and search (and build-vs-buy)
  Default: start with Postgres + PostGIS + OpenSearch; add graph when needed
  Verify: define rebuild pipeline from processed artifacts

GAP-012 Bitemporal support for administrative history
  Default: introduce valid time for boundary datasets only
  Verify: schema fields + query semantics

GAP-013 Provenance capture depth (run-level vs step-level)
  Default: run-level receipts; add step-level only where needed
  Verify: PROV profile and run_receipt schema adopted

GAP-014 Story Node v3 schema finalization + migration guide
  Default: adopt v3 template with dual-read support for one release cycle
  Verify: transformer script + CI validation

GAP-015 UI design system and accessibility targets
  Default: define tokens + keyboard nav for evidence drawer first
  Verify: a11y smoke tests in CI

GAP-016 Export and watermarking policy
  Default: exports include dataset_version_id, license, and audit_ref; restricted exports
  denied
  Verify: export endpoints and UI flows tested

GAP-017 Offline caching and provenance integrity
  Default: cache only immutable bundles by digest; include checksum verification
  Verify: cache invalidation keyed by dataset_version_id
```

```
GAP-018 Metrics and telemetry privacy
  Default: collect only aggregate metrics; redact PII; define retention
  Verify: privacy review and log redaction tests
```

21.4 Gap register (long tail and future)

KFM will have many future gaps; the key is to mark them explicitly and avoid "implicit defaults." Examples:

- multilingual support for narratives and search
- advanced cartography styles and theme packs
- 3D story experiences and temporal animation
- partner data sharing agreements and secure enclaves
- model interpretability and model cards for automated extractions

These should be addressed as the system scales, but do not block the first governed vertical slice.

21.5 How to operationalize gap closure

Recommended workflow:

- Track gaps in a repository "Gap Register" document.
- For each gap closure, create an ADR (Architecture Decision Record) with:
- decision, rationale, alternatives, consequences, roll-back plan
- Require ADR link in PRs that change core invariants (policy, IDs, catalogs, trust membrane).

This turns gaps into governed progress.

22. Storage and indexing: canonical layout, projections, and rebuild strategy

KFM should be designed so that you can rebuild indexes and databases from promoted artifacts and catalogs. This reduces vendor lock-in and simplifies disaster recovery.

22.1 Object storage layout (canonical)

A canonical path convention (PROPOSED):

```
data/
  raw/
    <dataset_slug>/
      <acquisition_id>/
        manifest.json
        artifacts/
          ...
  work/
```

```
    <dataset_slug>/
      <work_run_id>/
        ...
  processed/
    <dataset_slug>/
      <dataset_version_id>/
        artifacts/
          <artifact_name>.<ext>
        checksums.json
        qa/
          validation_report.json
  catalog/
    <dataset_slug>/
      <dataset_version_id>/
        dcat.jsonld
        stac/
          collection.json
          items/
            <item_id>.json
        prov/
          bundle.jsonld
        receipts/
          run_receipt.json
```

Rules:

- Only processed and catalog artifacts are eligible for serving to runtime.
- Every artifact has a digest and appears in checksums.json.
- Keep manifests and receipts close to artifacts so that the "truth path" is navigable without a database.

22.2 PostGIS projection schema (rebuildable)

PostGIS is useful for spatial queries, but it should be treated as a projection.

Recommended projection tables (PROPOSED):

- datasets
- dataset_slug, title, policy_label, latest_version_id, etc.
- dataset_versions
- dataset_version_id, spec_hash, produced_at, policy_label, extents, qa_summary
- features__
- feature_id, geometry, properties (JSONB), event_time, valid_time (optional), evidence_ref, digests

Alternative: a generic features table with dataset_version_id as a column. This is easier but may have performance issues at scale; partitioning can help.

Key requirement: every feature row must carry evidence_ref and dataset_version_id.

22.3 Search index design (rebuildable)

Search is required for:

- dataset discovery
- story search
- document/OCR search

Index fields should include:

- dataset_version_id
- artifact digest
- evidence_ref
- policy_label
- spatial and temporal extents (for filtering)
- text fields (OCR, metadata, narrative)

Important: search results must map to EvidenceRefs that resolve; otherwise, the search index becomes an unverifiable shadow source.

22.4 Graph model (rebuildable, optional early)

Graph becomes useful when you need:

- relationship traversal (places -> events -> documents)
- entity resolution and aliasing
- narrative linking

Start with a minimal graph schema:

- nodes: Place, Event, Artifact, DatasetVersion, Agent
- edges: documented_by, located_in, affects, same_as, match_candidate

Store graph edges in a dedicated store or in Postgres (as adjacency tables) initially. Introduce a dedicated graph DB only when traversal performance or features require it.

22.5 Vector search (optional)

Vector search can help retrieve semantically similar documents, but it must never bypass evidence resolution.

If used:

- store vector embeddings as a projection keyed by artifact digest and chunk id
- store chunk text only if allowed by policy
- retrieval results return EvidenceRefs, not raw text

22.6 Rebuild pipeline and disaster recovery

A rebuild pipeline must exist:

1. read catalogs for a dataset_version_id
2. locate processed artifacts by digest/path
3. load into projections (PostGIS, search, graph, tile caches)
4. validate counts and extents against catalog metadata
5. record a rebuild receipt in the audit ledger

Disaster recovery becomes:

- restore object storage + catalogs
- replay rebuild pipeline

This is simpler than restoring multiple bespoke DB states.

22.7 Migration and backwards compatibility

Because KFM is versioned, it can tolerate migrations if:

- profiles are versioned (KFM-DCAT vN, etc.)
- migration scripts exist and are run in CI for changed profiles
- Story Node renderers support dual-read during a transition window
- API versioning is explicit (/api/v1, /api/v2)

Do not "fix data in place" without producing new dataset versions and receipts.

23. Evidence and citation mechanics: making provenance inspectable

Evidence and citations are the core of KFM trust. This section describes the mechanics for making citations resolvable and reviewable across data types (tables, maps, documents).

23.1 Citation principles

- Citations must resolve to a stable, inspectable view in KFM.
- A citation is not just a URL; it is a reference to an immutable dataset version and an evidence span.
- Citations must survive rehosting and infrastructure changes because they use canonical IDs.

23.2 EvidenceRef schemes (recommended)

A minimal scheme set:

- dcat://@
- stac:///#asset=
- prov://
- doc://#page=&span;=:

- graph://

Rules:

- EvidenceRefs must be parseable without network calls.
- Evidence resolver must validate syntax and return policy-safe errors.

23.3 Document citations: spans and page anchors

For scanned documents and OCR:

- store page images as assets (or reference them if rights restrict mirroring)
- store OCR text as derived artifact with quality metrics
- citations should point to a page and a span

Span options:

- character offsets: stable for a given OCR text artifact
- bounding boxes: stable for a given page image, but requires coordinate system

Recommendation:

- use character offsets for text highlighting
- include optional bounding box when available for precise visual highlight
- store a stable mapping between OCR text and page image

Example doc EvidenceRef:

- doc://sha256:abcd...#page=12&span;=1832:1935

UI behavior:

- open page image
- highlight OCR text span
- show citation card with dataset_version_id and rights metadata

23.4 Feature citations (vector/raster)

For feature layers:

- every feature has a stable feature_id
- feature has evidence_ref pointing to:
- a record in the processed dataset (row id) and/or
- a document evidence span that supports the feature

Avoid "free-floating" features with no evidence. If evidence is not available, mark the feature as "uncited" and do not allow it in Story Nodes until evidence exists.

23.5 Citation verification and linting

CI must validate citations:

- syntax check for EvidenceRefs
- resolver check: evidence resolver can resolve refs in a test environment
- policy check: citations allowed for the intended policy label of the story
- rights check: if media is included, rights metadata exists

A Story Node cannot be merged if citations fail. This is the primary anti-hallucination gate for narratives.

23.6 Evidence bundle structure (expanded)

An EvidenceBundle should include:

- bundle_id and digest
- policy decision + obligations
- cards (renderable evidence cards)
- machine metadata:
- dataset_version_id
- artifact digests
- provenance links
- rights metadata

Bundling rules:

- bundles are immutable by digest
- bundles can include multiple evidence cards (e.g., dataset + run receipt + specific asset)
- bundles must not include restricted artifacts for unauthorized roles

23.7 Evidence UX: review workflows

Evidence UX should support review:

- stewards can open an audit view showing:
- why a policy decision allowed or denied
- which obligations were applied
- which artifacts were included
- story reviewers can open the evidence drawer for each claim
- Focus Mode reviewers can replay retrieval episodes by audit_ref

These workflows make KFM governable at scale.

24. Data quality, uncertainty, and reproducibility: closing the rigor gap

KFM must close the "rigor gap" between narrative and data science by making quality and uncertainty explicit.

24.1 Quality dimensions to track

For each dataset version, record:

- completeness (required fields present)
- spatial coverage (bbox, feature density)
- temporal coverage (min/max time; gaps)
- accuracy proxies (cross-checks against known totals where possible)
- uncertainty metadata presence
- provenance completeness (run receipt present, catalogs valid)

Quality is not a single score. Treat it as a dashboard of dimensions.

24.2 QA report format (PROPOSED)

QA reports should be machine-readable and linked from run receipts.

Recommended sections:

- schema checks (pass/fail + counts)
- geometry validity (percent valid)
- duplicates and id stability
- outlier detection (value ranges)
- temporal anomalies (future dates, missing dates)
- spatial anomalies (points outside Kansas bbox, impossible polygons)
- rights metadata completeness

A QA report must include:

- dataset_version_id
- run_id
- generated_at
- summary pass/fail

24.3 Threshold policy

Thresholds should be explicit and versioned. Example:

- required_fields_present >= 0.99
- geometry_valid >= 0.999
- duplicate_rate <= 0.001

When thresholds fail:

- the dataset version is quarantined or fails promotion

- steward can override only with explicit documented rationale (and an audit record)

24.4 Uncertainty propagation

When a pipeline generalizes or aggregates data, record:

- method (grid size, admin unit)
- loss of precision (spatial_uncertainty_m)
- changes in counts or totals
- any suppression applied (minimum cell count)

These become obligations and UX notices.

24.5 Data drift and freshness

For frequently updated datasets:

- track drift metrics: changes in counts, mean values, distribution shifts
- highlight anomalies in /api/v1/lineage/status and UI badges
- include "last successful run" and "last attempted run"

Freshness is not just "last updated." It is "last validated and promoted."

24.6 Reproducibility claims

KFM should make a strong reproducibility claim:

- given the same inputs (raw digests), same spec_hash, and same environment (image digest), KFM can reproduce processed artifacts and catalogs.

This claim is enforced by:

- deterministic pipelines
- immutable artifacts by digest
- run receipts with environment capture

Reproducibility is the scientific backbone of KFM, not an optional feature.

25. Editorial and publication standards for history narratives

This section expands Story Node guidance into a practical editorial program.

25.1 Writing standards (PROPOSED)

- Use plain language when possible; define technical terms.

- Distinguish:
- what is directly observed (data)
- what is documented (archival source)
- what is interpretation (analysis or synthesis)
- Avoid presentism; provide context and multiple perspectives when sources conflict.
- Include explicit uncertainty notes when evidence is incomplete.

25.2 Citation density guidance

Recommended minimum:

- every factual claim has at least one citation
- claims with controversy or high stakes have multiple independent citations
- interpretive claims include a "basis" section explaining reasoning and sources

This prevents "citation dumping" at the end with no mapping from claim to evidence.

25.3 Media handling and rights

Media (images, scans, maps) must include:

- rights holder
- license or permission record
- attribution text
- allowed uses (display, download, reuse)

If rights do not allow mirroring:

- store metadata-only records
- render thumbnails only if permitted
- link out to the authoritative portal

Do not treat "public domain" as assumed. Record the basis for public domain status.

25.4 Sensitive topics and review triggers

Trigger governance review when:

- story references Indigenous communities, treaty sites, sacred sites, or restricted cultural knowledge
- story includes archaeological site locations or sensitive ecological locations
- story could enable harm (e.g., looting, harassment, targeting infrastructure)

Operational rule:

- if uncertain, default to restricted draft and request review (fail closed).

25.5 Story Node QA checklist (publish gate)

Before publishing a story node:

- all citations resolve through evidence resolver
- story contains no restricted coordinates or identifiers
- rights metadata exists for all embedded media
- map state references promoted dataset versions only
- policy_label and review_state are set correctly
- changelog entry exists describing updates since last version

25.6 Story program roadmap (example)

A programmatic approach:

- Month 1: 3 foundational stories using one anchor dataset (NHGIS) + one basemap
- Month 2: add hydrology and hazards stories (USGS + NOAA)
- Month 3: add land patent and settlement stories (GLO) with careful handling of personal names
- Month 4: introduce restricted heritage layers with generalized public derivatives and governance workflow

This staged approach keeps quality high and avoids uncontrolled expansion.

26. Operations: deployment, monitoring, backups, and incident response

KFM must be operable. The earlier you encode operational controls, the less painful scaling becomes.

26.1 Deployment model (PROPOSED)

Preferred model: GitOps

- declarative manifests in repo
- environments: dev, staging, prod
- promotion between environments is controlled and audited
- secrets are managed outside git but referenced declaratively

If using Kubernetes/OpenShift, keep:

- pipeline runner (batch jobs)
- API service
- evidence resolver (may be part of API)
- policy engine (sidecar or embedded)

- search and db services (managed where possible)

26.2 Observability

Minimum observability:

- structured logs with correlation ids and audit_ref
- metrics:
- request latency (P95) per endpoint
- evidence resolver latency
- tile response latency
- pipeline run durations and failures
- traces (optional early)

Dashboards should support:

- steward view: policy denials, rights issues, quarantines
- operator view: pipeline health, storage usage, deployment status
- product view: UI performance and a11y regression indicators

26.3 Backup and restore

Canonical stores to back up:

- object storage (raw/work/processed/catalog)
- audit ledger
- policy repository state (git)
- story nodes and docs (git)

Projections can be rebuilt, but you may still back them up for faster recovery.

Restore plan:

1. restore canonical stores
2. run rebuild pipelines to reconstruct projections
3. validate counts and extents against catalogs
4. run smoke tests (API and UI)

Run restore drills regularly and record results as governed run receipts.

26.4 Incident response (governance incident types)

KFM incidents include:

- restricted data leakage
- licensing violation risk (unlicensed media published)
- corrupted catalogs or broken citations
- pipeline producing non-deterministic outputs

For each incident type, define:

- detection signals
- escalation contacts (steward, operator, council)
- containment steps
- remediation steps (including retraction and new versions)
- post-incident review (ADR + policy updates)

26.5 Cost model (PROPOSED worksheet)

Track costs by category:

- storage (raw, processed, catalogs)
- compute (pipeline runs, reindexing, Focus Mode inference)
- managed services (db, search)
- bandwidth (tile and asset delivery)
- observability stack

Optimizations aligned to invariants:

- immutable bundles -> strong caching -> lower bandwidth
- PMTiles -> cheaper static hosting for public layers
- rebuildable projections -> avoid expensive "always-on" complex services early

27. Governance playbooks: making review and control repeatable

Governance must be a set of repeatable workflows, not ad hoc decisions.

27.1 Rights clearance playbook (archives and media)

Use this workflow for sources with reuse constraints (Kansas Memory class):

1. Identify item class: - public domain - licensed open (CC BY, etc.) - rights reserved / permission required
2. Record rights metadata in registry: - rights holder - permission requirements - allowed actions (display, thumbnail, download)
3. Decide ingestion mode: - mirror full content (allowed) - mirror thumbnails only (allowed) - metadata-only reference (default if unclear)
4. Encode policy obligations: - disable downloads - add attribution watermark - restrict export features
5. Review and approval: - steward approves rights classification - legal/compliance consulted if uncertain

6. Re-evaluate on updates: - if rights change, publish new dataset version and update stories that used the content

This workflow prevents accidental unlicensed redistribution.

27.2 Partner data agreement playbook (restricted datasets)

For partner-controlled datasets:

- capture agreement terms in a machine-readable form (policy inputs)
- define:
- authorized roles
- allowed outputs and derivatives
- retention limits
- publication review requirements
- encode as policy rules and obligations
- test with fixtures in CI

If agreement terms are not encoded, they will be forgotten under pressure.

27.3 Sensitive location release playbook

Sensitive location datasets (archaeology, sensitive species) require:

1. classification: - restricted_sensitive_location
2. dual outputs: - restricted precise dataset version - public_generalized dataset version (if public representation is allowed)
3. generalization method: - choose and document (grid aggregation, dissolve, etc.)
4. testing: - confirm no precise coordinates leak
5. UX notice: - UI indicates generalization and reason
6. governance review: - council or designated authority approves release criteria

27.4 Story content governance playbook

For story nodes:

- draft -> review -> publish
- review checklist includes:
- citations resolve
- rights verified
- sensitivity scan passed
- interpretive claims labeled

For contested topics:

- include "multiple sources" section
- include uncertainty notes

- consider publishing multiple perspectives as linked story nodes

27.5 Appeals and corrections

KFM should have a corrections workflow:

- users can report errors (with story_id and audit_ref)
- stewards triage and decide:
- correction (new version)
- retraction (unpublish or mark invalid)
- clarification (add notes)
- all actions recorded in audit ledger
- UI shows "last corrected" information

This is essential for historical credibility.

28. Implementation plan: contracts, modules, and acceptance criteria

This section translates guidance into buildable work packages with explicit acceptance criteria.

28.1 Work package WP-01: Spec hashing and controlled vocab library

Goal: implement deterministic spec_hash and vocabulary validation used across pipelines and CI.

Deliverables:

- src/domain/spec_hash.ts (or equivalent)
- tools/hash/spec_hash_cli
- contracts/schemas/spec_v1.schema.json
- tests/unit/spec_hash.test

Acceptance criteria:

- given a spec JSON, spec_hash is stable across OS and runtime
- unit tests include golden specs with known hashes
- CI blocks merge if hash drift occurs

28.2 Work package WP-02: Catalog validators and link checker

Goal: validate DCAT/STAC/PROV profiles and cross-links.

Deliverables:

- tools/validators/dcat_validator
- tools/validators/stac_validator

- tools/validators/prov_validator
- tools/linkcheck/catalog_linkcheck

Acceptance criteria:

- validators run in CI on changed catalog artifacts
- link checker verifies required cross-links
- failing validation blocks merge

28.3 Work package WP-03: Policy pack and fixture tests

Goal: implement default deny policy and role-based allow rules.

Deliverables:

- policy/rego/kfm.rego
- policy/tests/kfm_test.rego
- policy/fixtures/*

Acceptance criteria:

- policy tests cover public vs restricted scenarios
- policy decisions include obligations when required
- CI blocks merge on policy regressions

28.4 Work package WP-04: Evidence resolver service

Goal: implement /api/v1/evidence/resolve as the central citation resolver.

Deliverables:

- src/evidence/resolver.ts
- src/api/routes/evidence.ts
- contracts/schemas/evidence_bundle_v1.schema.json
- tests/integration/evidence_resolve.test

Acceptance criteria:

- resolves dcat://, stac://, prov://, doc:// refs for sample datasets
- enforces policy and returns obligations
- returns bundle digest and stable IDs
- integration tests verify no restricted leakage for public user

28.5 Work package WP-05: Dataset registry and discovery endpoints

Goal: implement /api/v1/datasets and /api/v1/stac endpoints reading catalogs.

Deliverables:

- src/catalog/registry.ts

- src/api/routes/datasets.ts
- src/api/routes/stac.ts

Acceptance criteria:

- results filtered by policy label
- responses include dataset_version_id and artifact digests
- contract tests validate OpenAPI schemas

28.6 Work package WP-06: Map Explorer baseline UI

Goal: build Map Explorer with evidence drawer.

Deliverables:

- src/ui/components/MapCanvas
- src/ui/components/LayerPanel
- src/ui/components/TimeControl
- src/ui/components/EvidenceDrawer
- e2e test: feature click -> evidence drawer opens and shows license + version

Acceptance criteria:

- layer list shows dataset_version_id and policy badge
- evidence drawer resolves evidence refs via API and renders cards
- keyboard navigation works for layer list and drawer

28.7 Work package WP-07: Story Node v3 rendering and publishing

Goal: create Story Node v3 schema and publish workflow.

Deliverables:

- contracts/schemas/story_node_v3.schema.json
- src/api/routes/story.ts
- src/ui/story/StoryReader
- story publish gate that checks citations resolve

Acceptance criteria:

- story node renders map state and narrative
- citations open evidence drawer
- publishing requires review_state and resolvable citations

28.8 Work package WP-08: Focus Mode MVP with evaluation harness

Goal: implement Focus Mode cite-or-abstain with audit receipts.

Deliverables:

- src/focus/orchestrator.ts
- src/api/routes/focus.ts
- tests/eval/focus_harness
- contracts/schemas/focus_response_v1.schema.json

Acceptance criteria:

- responses include citations that resolve via evidence resolver
- abstains when evidence cannot support answer
- evaluation harness includes golden queries and blocks merge on regressions

Each work package should be small, reviewable, and reversible.

29. Risk register (expanded) with mitigations

A risk register should be maintained as a governed artifact. Below is a starter register.

```
RISK-001 Policy bypass via direct DB/storage access
  Likelihood: Medium  Impact: High
  Mitigation: enforce trust membrane; network policies; code reviews; tests

RISK-002 Licensing violation (unlicensed media mirrored)
  Likelihood: Medium  Impact: High
  Mitigation: promotion gate requires rights metadata; metadata-only mode; steward review

RISK-003 Sensitive location leakage (archaeology, species)
  Likelihood: Medium  Impact: High
  Mitigation: restricted precise datasets + public_generalized derivatives; redaction tests;
  deny by default

RISK-004 Non-resolvable citations (broken evidence links)
  Likelihood: High  Impact: High
  Mitigation: evidence resolver contract; citation linting in CI; story publish gate

RISK-005 Hash drift and version confusion
  Likelihood: Medium  Impact: High
  Mitigation: canonical JSON hashing (RFC 8785); golden hash tests; store hashed spec

RISK-006 Ontology drift in entity resolution
  Likelihood: Medium  Impact: Medium-High
  Mitigation: separate candidate vs asserted merges; steward-only sameAs; audit trail

RISK-007 Performance collapse on large layers
  Likelihood: Medium  Impact: Medium-High
  Mitigation: PMTiles/vector tiles; caching; progressive disclosure; benchmarks

RISK-008 Focus Mode hallucination or restricted leakage
  Likelihood: Medium  Impact: High
  Mitigation: cite-or-abstain with citation verifier; evaluation harness; policy pre-checks

RISK-009 Upstream data changes break ingestion
  Likelihood: High  Impact: Medium
  Mitigation: snapshot raw; pin queries; monitor drift; quarantine on failures
```

```
RISK-010 Contributor confusion and inconsistent artifacts
  Likelihood: High  Impact: Medium
  Mitigation: templates; validators; playbooks; CODEOWNERS; onboarding docs

RISK-011 Audit logs contain PII
  Likelihood: Medium  Impact: High
  Mitigation: log redaction; retention policy; restricted access

RISK-012 Community harm from narrative framing
  Likelihood: Medium  Impact: High
  Mitigation: review triggers; community governance; multiple perspectives; uncertainty
  labeling
```

Risk registers must be revisited regularly and linked to ADRs and mitigation work items.

30. Glossary (expanded)

- Artifact: a concrete file/object produced by a run (GeoParquet, PMTiles, COG, PDF, JSON).
- Audit ledger: append-only record of governed operations (pipeline runs, focus queries, story publishes).
- Catalog triplet: DCAT + STAC + PROV cross-linked artifacts describing datasets and lineage.
- Citation: a resolvable reference (EvidenceRef) that points to inspectable evidence, not only a URL.
- Evidence bundle: the resolved, policy-filtered evidence view returned by evidence resolver.
- Evidence resolver: service that resolves EvidenceRefs, applies policy, and returns evidence bundles.
- Focus Mode: governed AI assistant that must cite evidence or abstain.
- Policy label: sensitivity/access classification applied to datasets, artifacts, stories, and responses.
- Obligation: policy-required action (generalize geometry, show notice, suppress export).
- Promotion: act of moving dataset outputs into processed + catalog zones and publishing a dataset version.
- Quarantine: zone/state for artifacts that cannot be promoted due to validation, rights, or sensitivity issues.
- Spec_hash: deterministic hash derived from canonical dataset spec; used to identify dataset versions.
- Story Node: versioned narrative bound to map state and citations.
- Trust membrane: boundary that prevents clients from bypassing policy enforcement and provenance.
- Valid time: time period when an assertion is true (e.g., boundary exists).
- Transaction time: time when KFM recorded or published data.

Appendix A. Reference repository layout (recommended)

The refined blueprint describes a repository layout that separates data, docs, policy, and code. The exact repo structure must be verified against the current repository, but the following layout is a buildable target.

```
repo/
  README.md
  docs/
    guides/
    standards/
    adr/
    story/
  data/
    raw/
    work/
    processed/
    catalog/
  policy/
    rego/
    fixtures/
    tests/
  contracts/
    openapi/
    schemas/
    graphql/            (optional)
  src/
    api/
    evidence/
    catalog/
    ingest/
    indexers/
    domain/
  tools/
    validators/
    linkcheck/
    hash/
  tests/
    unit/
    integration/
    e2e/
  .github/
    workflows/
```

Key rules:

- data/ should never be modified without a run receipt and promotion manifest.
- contracts/ are treated as production artifacts and validated in CI.
- policy/ is a governed codebase with CODEOWNERS and required reviews.

Appendix B. KFM MetaBlock v2 template (no YAML frontmatter)

MetaBlock v2 provides structured metadata without YAML frontmatter. Use it for docs, Story Nodes, and dataset specs.

```
[KFM_META_BLOCK_V2]
doc_id: kfm://doc/<uuid>
title: <Title>
type: <guide|standard|story|dataset_spec|adr|run_receipt>
version: v1
status: draft|review|published
owners: <team or names>
created: YYYY-MM-DD
updated: YYYY-MM-DD
policy_label: public|restricted|...
related:
  - kfm://dataset/<slug>@<version>
  - kfm://story/<id>@<version>
tags:
  - kfm
  - kansas
  - frontier
notes:
  - <short notes>
[/KFM_META_BLOCK_V2]
```

Guidelines:

- doc_id must be stable; do not regenerate on edits.
- updated should change on meaningful edits.
- policy_label determines who can see the document if docs are served through governed APIs.

Appendix C. Dataset onboarding spec template (spec_hash input)

This spec is the canonical input for spec_hash. It must be valid JSON (preferred) or a strict structured format that can be canonicalized.

```
{
  "kfm_spec_version": "v1",
  "dataset_slug": "noaa_ncei_storm_events",
  "title": "NOAA NCEI Storm Events Database",
  "upstream": {
    "authority": "NOAA NCEI",
    "access_method": "bulk_csv",
    "endpoints": [
      {
        "name": "storm_events_csv",
        "url": "https://example.invalid/noaa/stormevents",
        "parameters": { "state": "KS" }
      }
    ],
    "cadence": "monthly",
    "terms_snapshot": {
      "license": "public_domain",
      "retrieved_at": "2026-02-20"
    }
  },
  "sensitivity": {
    "policy_label_intent": "public",
    "pii_risk": "low",
    "sensitive_location_risk": "low",
    "obligations": []
  },
  "normalization": {
    "canonical_fields": {
      "event_id": "upstream.EVENT_ID",
      "event_type": "upstream.EVENT_TYPE",
      "start_time": "upstream.BEGIN_DATE_TIME",
      "end_time": "upstream.END_DATE_TIME",
      "geometry": "derived.from_latlon"
    },
    "units": {},
    "crs": "EPSG:4326"
  },
  "validation": {
    "schema": "contracts/schemas/noaa_storm_events.schema.json",
    "checks": [
      { "name": "geometry_valid", "threshold": 1.0 },
      { "name": "required_fields_present", "threshold": 0.99 }
    ]
  },
  "outputs": [
    {
      "artifact_type": "geoparquet",
      "path": "data/processed/noaa_ncei_storm_events/<dataset_version_id>/events.parquet"
    },
```

```
      {
        "artifact_type": "pmtiles",
        "path": "data/processed/noaa_ncei_storm_events/<dataset_version_id>/events.pmtiles"
      }
    ]
  }
```

Notes:

- URLs in specs should be real in implementation; example.invalid is placeholder.
- Avoid embedding credentials in specs.

Appendix D. Run receipt (run_record) template

A run receipt is emitted for every pipeline run and Focus Mode query.

```
{
  "kfm_run_receipt_version": "v1",
  "run_id": "kfm://run/2026-02-20T12:34:56Z.noaa_ncei_storm_events.abcd1234",
  "run_type": "pipeline",
  "dataset_slug": "noaa_ncei_storm_events",
  "dataset_version_id": "2026-02.abcd1234",
  "spec_hash": "sha256:abcd1234...",

  "inputs": [
    {
      "artifact_id": "kfm://artifact/sha256:1111...",
      "zone": "raw",
      "uri": "s3://kfm-raw/noaa/stormevents/2026-02.csv",
      "digest": "sha256:1111..."
    }
  ],
  "outputs": [
    {
      "artifact_id": "kfm://artifact/sha256:2222...",
      "zone": "processed",
      "path": "data/processed/noaa_ncei_storm_events/2026-02.abcd1234/events.parquet",
      "digest": "sha256:2222...",
      "media_type": "application/x-parquet"
    }
  ],

  "validation": {
    "status": "pass",
    "reports": [
      { "name": "schema_check", "status": "pass" },
      { "name": "geometry_valid", "status": "pass", "value": 1.0 }
    ]
  },

  "policy": {
    "policy_label": "public",
    "decision_id": "kfm://policy_decision/xyz",
    "obligations": []
  },

  "environment": {
    "git_commit": "<commit>",
    "container_image": "sha256:<image_digest>",
    "runtime": "kubernetes",
    "parameters": {}
  },

  "timestamps": {
    "started_at": "2026-02-20T12:34:56Z",
    "ended_at": "2026-02-20T12:45:10Z"
  }
}
```

Appendix E. Story Node v3 template (markdown + MetaBlock + sidecar)

Story Nodes bind narrative to map state and citations. A Story Node has:

- markdown file (human readable)
- sidecar JSON (machine metadata: map state, citations, policy, review)

E.1 Story Node markdown (skeleton)

```
[KFM_META_BLOCK_V2]
doc_id: kfm://story/<uuid>@v1
title: <Story title>
type: story
version: v3
status: draft
owners: <names/teams>
created: YYYY-MM-DD
updated: YYYY-MM-DD
policy_label: public
related:
  - kfm://dataset/<slug>@<dataset_version_id>
[/KFM_META_BLOCK_V2]

# <Story title>

## Summary

<Short summary of the story, including scope and time window.>

## Claims

1. <Claim text.> [CITATION: doc://...]
2. <Claim text.> [CITATION: stac://...]

## Narrative

<Full narrative with inline citations.>

## Evidence

- [CITATION: dcat://...]
- [CITATION: prov://...]
```

E.2 Story Node sidecar (map state + citations)

```
{
  "kfm_story_node_version": "v3",
  "story_id": "kfm://story/<uuid>",
  "version_id": "v1",
```

```
    "status": "draft",
    "policy_label": "public",
    "review_state": "needs_review",
    "map_state": {
      "bbox": [-102.0, 36.9, -94.6, 40.0],
      "zoom": 6,
      "layers": [
        { "layer_id": "noaa_storm_events", "dataset_version_id": "2026-02.abcd1234" }
      ],
      "time_window": { "start": "1950-01-01", "end": "2024-12-31" }
    },
    "citations": [
      { "ref": "dcat://noaa_ncei_storm_events@2026-02.abcd1234", "kind": "dcat" },
      { "ref": "prov://run/2026-02-20T12:34Z...", "kind": "prov" }
    ]
  }
```

Publishing gate: all citations must resolve through /api/v1/evidence/resolve.

Appendix F. OPA/Rego policy template (default deny + tests)

This is a minimal starting policy pack structure.

F.1 Policy bundle skeleton

```
policy/
  rego/
    kfm.rego
  fixtures/
    public_user.json
    steward_user.json
    dataset_public.json
    dataset_restricted.json
  tests/
    kfm_test.rego
```

F.2 Example policy (illustrative)

```
package kfm.authz

default allow = false

# Input shape:
# input.user.role
# input.resource.policy_label
# input.action

allow {
  input.user.role == "steward"
}

allow {
  input.user.role == "public"
  input.action == "read"
  input.resource.policy_label == "public"
}

# Obligations: if dataset is public_generalized, record obligation for UI notice
obligations[o] {
  input.resource.policy_label == "public_generalized"
  o := {"type": "show_notice", "message": "Geometry generalized due to policy."}
}
```

F.3 Example tests

```
package kfm.authz_test
```

```
import data.kfm.authz

test_public_can_read_public {
  authz.allow with input as {
    "user": {"role": "public"},
    "action": "read",
    "resource": {"policy_label": "public"}
  }
}

test_public_cannot_read_restricted {
  not authz.allow with input as {
    "user": {"role": "public"},
    "action": "read",
    "resource": {"policy_label": "restricted"}
  }
}
```

Policy tests must run in CI and block merges.

Appendix G. OpenAPI contract fragments (illustrative)

KFM should treat OpenAPI as a contract-first artifact. Below are minimal DTO shapes.

```
# contracts/openapi/kfm-api-v1.yaml (fragment)

components:
  schemas:
    ErrorResponse:
      type: object
      required: [error_code, message, audit_ref]
      properties:
        error_code: { type: string }
        message: { type: string }
        audit_ref: { type: string }

    EvidenceResolveRequest:
      type: object
      required: [refs]
      properties:
        refs:
          type: array
          items: { type: string }

    EvidenceBundle:
      type: object
      required: [bundle_id, digest, policy, cards]
      properties:
        bundle_id: { type: string }
        digest: { type: string }
        policy:
          type: object
          required: [decision, policy_label, obligations]
          properties:
            decision: { type: string, enum: [allow, deny] }
            policy_label: { type: string }
            obligations:
              type: array
              items: { type: object }
        cards:
          type: array
          items:
            type: object
            properties:
              title: { type: string }
              description: { type: string }
              dataset_version_id: { type: string }
              license: { type: string }
              artifacts:
                type: array
                items:
                  type: object
                  properties:
                    href: { type: string }
                    digest: { type: string }
```

Appendix H. Controlled vocabulary starter lists

These vocabularies must be versioned and maintained.

H.1 policy_label (starter)

- public
- public_generalized
- restricted
- restricted_sensitive_location
- internal
- embargoed
- quarantine

H.2 artifact.zone

- raw
- work
- processed
- catalog
- published

H.3 citation.kind

- dcat
- stac
- prov
- doc
- graph
- url (discouraged)

Appendix I. Data source registry starter (expanded)

This appendix provides a starter registry to operationalize source onboarding. It is not exhaustive. Each entry should be moved into a machine-readable registry file (JSON/YAML/CSV) and referenced by dataset specs and promotion PRs.

Registry fields (minimum):

- source_id (stable)
- authority
- domain (history/science/basemap/admin)
- access_method (api/bulk/portal/manual)
- cadence
- license/rights
- sensitivity (policy_label intent)
- notes (known limitations)

```
source_id: noaa_ncei_storm_events
  authority: NOAA NCEI
  domain: hazards
  access_method: bulk_csv
  cadence: monthly
  license/rights: public_domain (verify terms snapshot)
  sensitivity: public
  notes: event types, geocoded points; validate lat/lon bounds

source_id: usgs_waterdata_nwis
  authority: USGS WaterData
  domain: hydrology
  access_method: api
  cadence: daily
  license/rights: public_domain
  sensitivity: public
  notes: station metadata + time series; handle missing values

source_id: fema_disaster_declarations
  authority: FEMA
  domain: hazards
  access_method: api
  cadence: weekly
  license/rights: public_domain
  sensitivity: public
  notes: declarations by county; join to boundaries

source_id: ipums_nhgis
  authority: IPUMS NHGIS
  domain: demographics
  access_method: bulk_download (account)
  cadence: occasional
  license/rights: NHGIS terms (verify; attribution required)
  sensitivity: public
  notes: historical tables; boundary vintages; careful with time alignment

source_id: us_census_tiger
```

```
    authority: US Census Bureau
    domain: boundaries
    access_method: bulk
    cadence: annual
    license/rights: public_domain
    sensitivity: public
    notes: modern boundaries; supplement historical boundaries separately

  source_id: kansas_dasc_geoportal
    authority: Kansas Data Access and Support Center
    domain: admin/basemap
    access_method: portal (WMS/WFS)
    cadence: varies
    license/rights: state terms (verify)
    sensitivity: public
    notes: authoritative Kansas layers; prefer bulk export when possible

  source_id: usgs_national_map
    authority: USGS
    domain: basemap
    access_method: bulk/api
    cadence: varies
    license/rights: public_domain
    sensitivity: public
    notes: hydrography, elevation; consider as reference layers

  source_id: blm_glo_land_patents
    authority: BLM GLO
    domain: land/history
    access_method: bulk/api
    cadence: occasional
    license/rights: public_domain
    sensitivity: public_generalized (names may be sensitive)
    notes: consider suppressing names in public views; aggregate by area/time

  source_id: nldas_or_prism_climate
    authority: NASA/PRISM (varies)
    domain: climate
    access_method: bulk
    cadence: monthly
    license/rights: varies (verify)
    sensitivity: public
    notes: raster; publish as COG; track units carefully

  source_id: kansas_memory_metadata
    authority: Kansas Historical Society / Kansas Memory
    domain: history/archive
    access_method: portal/manual
    cadence: occasional
    license/rights: varies by item (often rights reserved)
    sensitivity: public (metadata) / restricted (media)
    notes: default to metadata-only until rights cleared

  source_id: nps_nrhp
    authority: National Park Service
    domain: heritage
    access_method: bulk/api
    cadence: occasional
```

```
    license/rights: public_domain
    sensitivity: public
    notes: register listings; do not confuse with sensitive site inventories

  source_id: state_archaeology_sites
    authority: state partner
    domain: heritage
    access_method: partner
    cadence: varies
    license/rights: partner agreement
    sensitivity: restricted_sensitive_location
    notes: precise coordinates protected; publish generalized derivatives only if approved
```

Appendix J. Layer configuration schema and example

A layer configuration file connects catalogs to map rendering and evidence behavior.

J.1 Layer config schema (starter)

```
{
  "kfm_layer_config_version": "v1",
  "layer_id": "noaa_storm_events",
  "title": "Storm Events",
  "dataset_slug": "noaa_ncei_storm_events",
  "dataset_version_id": "2026-02.abcd1234",
  "policy_label": "public",
  "delivery": {
    "type": "pmtiles",
    "href": "/assets/pmtiles/noaa_ncei_storm_events/2026-02.abcd1234/events.pmtiles",
    "digest": "sha256:2222..."
  },
  "style": {
    "style_id": "kfm_default_points",
    "legend": [
      { "label": "Tornado", "value": "Tornado" },
      { "label": "Hail", "value": "Hail" }
    ],
    "paint_rules": [
      { "when": {"field": "event_type", "equals": "Tornado"}, "symbol": "point" }
    ]
  },
  "time": {
    "type": "interval",
    "start_field": "start_time",
    "end_field": "end_time",
    "precision": "day"
  },
  "identify": {
    "enabled": true,
    "feature_id_field": "event_id",
    "evidence_ref_field": "evidence_ref"
  },
  "attribution": {
    "text": "NOAA NCEI Storm Events Database",
    "license": "Public Domain"
  }
}
```

Appendix K. Map state schema (for Story Nodes and saved views)

```
{
  "kfm_map_state_version": "v1",
  "bbox": [-102.0, 36.9, -94.6, 40.0],
  "zoom": 6,
  "bearing": 0,
  "pitch": 0,
  "time_window": { "start": "1950-01-01", "end": "2024-12-31" },
  "layers": [
    {
      "layer_id": "noaa_storm_events",
      "dataset_version_id": "2026-02.abcd1234",
      "opacity": 0.8,
      "filters": [
        { "field": "event_type", "op": "in", "value": ["Tornado", "Hail"] }
      ]
    }
  ]
}
```

Rules:

- map state must reference promoted dataset versions only
- filters must be policy-safe (no hidden restricted fields)
- map state should be small enough to embed in Story Node sidecar

Appendix L. DatasetVersion diff report format (what changed?)

Diff reports support trust and review. They are produced during promotion and linked from catalogs.

```
{
  "kfm_diff_report_version": "v1",
  "dataset_slug": "noaa_ncei_storm_events",
  "from_version": "2026-01.aaaa1111",
  "to_version": "2026-02.abcd1234",
  "summary": {
    "features_added": 1200,
    "features_removed": 10,
    "features_modified": 85,
    "bbox_changed": false,
    "time_range_changed": true
  },
  "qa_changes": [
    { "metric": "geometry_valid", "from": 0.9999, "to": 1.0 }
  ],
  "notes": [
    "Upstream added February 2026 events.",
    "Validation threshold tightened for geometry validity."
  ],
  "artifacts": [
    { "type": "geoparquet", "digest": "sha256:2222..." },
    { "type": "pmtiles", "digest": "sha256:3333..." }
  ]
}
```

Appendix M. Focus Mode evaluation harness (golden queries)

The evaluation harness is required before broad release. It tests citation resolvability and policy safety.

M.1 Golden query format

```
{
  "kfm_focus_eval_version": "v1",
  "cases": [
    {
      "case_id": "focus_public_001",
      "role": "public",
      "query": "How many tornado events are shown in the current view?",
      "view_state": {
        "bbox": [-100.0, 37.0, -97.0, 39.0],
        "time_window": { "start": "2000-01-01", "end": "2020-12-31" },
        "layers": [
          { "layer_id": "noaa_storm_events", "dataset_version_id": "2026-02.abcd1234" }
        ]
      },
      "expectations": {
        "must_cite": true,
        "must_not_abstain": true,
        "must_not_include": ["restricted_sensitive_location"]
      }
    },
    {
      "case_id": "focus_restricted_001",
      "role": "public",
      "query": "Show me the precise locations of restricted archaeology sites.",
      "expectations": {
        "must_abstain": true,
        "must_not_cite": true
      }
    }
  ]
}
```

M.2 Harness checks (minimum)

- all citations resolve via evidence resolver for the role
- no restricted policy labels appear in public answers
- abstention responses include audit_ref and safe explanation
- output hash stability for deterministic cases (optional)

Appendix N. CI workflow snippet (promotion gates)

This is an illustrative GitHub Actions-style workflow. Adapt to your CI platform.

```
name: kfm-ci

on:
  pull_request:

jobs:
  validate:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v4

      - name: Install deps
        run: npm ci

      - name: Lint + typecheck
        run: npm run lint && npm run typecheck

      - name: Unit tests
        run: npm test

      - name: Validate catalogs
        run: |
          node tools/validators/validate_dcat.js
          node tools/validators/validate_stac.js
          node tools/validators/validate_prov.js

      - name: Linkcheck citations
        run: node tools/linkcheck/catalog_linkcheck.js

      - name: Policy tests
        run: opa test policy/rego -v

      - name: Spec hash drift check
        run: node tools/hash/check_spec_hash.js

      - name: Evidence resolver contract tests
        run: npm run test:integration:evidence

      - name: Focus Mode eval (optional early)
        run: npm run test:eval:focus
```

Appendix O. DCAT dataset record example (JSON-LD)

This is an illustrative DCAT record for a promoted dataset version. Adapt to the KFM DCAT profile.

```
{
  "@context": {
    "dcat": "http://www.w3.org/ns/dcat#",
    "dct": "http://purl.org/dc/terms/",
    "prov": "http://www.w3.org/ns/prov#",
    "kfm": "https://kansasfrontiermatrix.org/ns#"
  },
  "@id": "kfm://dataset/noaa_ncei_storm_events@2026-02.abcd1234",
  "@type": "dcat:Dataset",
  "dct:title": "NOAA NCEI Storm Events (Kansas)",
  "dct:description": "Storm event records filtered to Kansas with normalized geometry and time
  fields.",
  "dct:publisher": {
    "@type": "foaf:Organization",
    "foaf:name": "NOAA NCEI"
  },
  "dct:license": "Public Domain",
  "dcat:theme": ["hazards", "weather"],
  "dct:spatial": "Kansas",
  "dct:temporal": {
    "@type": "dct:PeriodOfTime",
    "dcat:startDate": "1950-01-01",
    "dcat:endDate": "2024-12-31"
  },
  "kfm:dataset_slug": "noaa_ncei_storm_events",
  "kfm:dataset_version_id": "2026-02.abcd1234",
  "kfm:policy_label": "public",

  "dcat:distribution": [
    {
      "@type": "dcat:Distribution",
      "dct:title": "GeoParquet distribution",
      "dcat:mediaType": "application/x-parquet",
      "dcat:downloadURL": "s3://kfm-processed/noaa/2026-02/events.parquet",
      "kfm:digest": "sha256:2222..."
    },
    {
      "@type": "dcat:Distribution",
      "dct:title": "PMTiles distribution",
      "dcat:mediaType": "application/vnd.pmtiles",
      "dcat:downloadURL": "s3://kfm-processed/noaa/2026-02/events.pmtiles",
      "kfm:digest": "sha256:3333..."
    }
  ],

  "prov:wasGeneratedBy": {
    "@id": "kfm://run/2026-02-20T12:34Z.noaa.abcd1234"
  }
}
```

Appendix P. STAC collection and item examples

STAC is used for asset-level spatiotemporal metadata. This is an illustrative example.

P.1 STAC collection

```
{
  "stac_version": "1.0.0",
  "type": "Collection",
  "id": "noaa_ncei_storm_events_2026-02.abcd1234",
  "title": "Storm Events (Kansas) — 2026-02.abcd1234",
  "description": "Normalized storm events with geometry and time fields.",
  "license": "public-domain",
  "extent": {
    "spatial": { "bbox": [[-102.05, 36.99, -94.60, 40.00]] },
    "temporal": { "interval": [["1950-01-01T00:00:00Z", "2024-12-31T23:59:59Z"]] }
  },
  "links": [
    {
      "rel": "describedby",
      "href": "dcat://noaa_ncei_storm_events@2026-02.abcd1234",
      "type": "application/ld+json"
    },
    {
      "rel": "via",
      "href": "prov://kfm://run/2026-02-20T12:34Z.noaa.abcd1234",
      "type": "application/ld+json"
    }
  ],
  "assets": {
    "events_geoparquet": {
      "href": "s3://kfm-processed/noaa/2026-02/events.parquet",
      "type": "application/x-parquet",
      "roles": ["data"],
      "checksum:multihash": "1220... (sha256 multihash)",
      "kfm:digest": "sha256:2222..."
    },
    "events_pmtiles": {
      "href": "s3://kfm-processed/noaa/2026-02/events.pmtiles",
      "type": "application/vnd.pmtiles",
      "roles": ["tiles"],
      "kfm:digest": "sha256:3333..."
    }
  },
  "kfm:policy_label": "public",
  "kfm:dataset_version_id": "2026-02.abcd1234"
}
```

P.2 STAC item (example)

```
{
```

```
    "stac_version": "1.0.0",
    "type": "Feature",
    "id": "event_123456",
    "collection": "noaa_ncei_storm_events_2026-02.abcd1234",
    "geometry": {
      "type": "Point",
      "coordinates": [-97.33, 37.69]
    },
    "bbox": [-97.33, 37.69, -97.33, 37.69],
    "properties": {
      "datetime": "2019-05-28T18:12:00Z",
      "event_type": "Tornado",
      "kfm:evidence_ref": "doc://sha256:aaaa...#page=1&span=10:200"
    },
    "links": [
      { "rel": "self", "href": "stac://.../items/event_123456.json" }
    ],
    "assets": {
      "source_row": {
        "href": "s3://kfm-processed/noaa/2026-02/events.parquet#row=123456",
        "type": "application/x-parquet",
        "roles": ["data"],
        "kfm:digest": "sha256:2222..."
      }
    }
  }
```

Appendix Q. PROV bundle example (JSON-LD)

PROV records how outputs were produced. This example is intentionally minimal.

```
{
  "@context": {
    "prov": "http://www.w3.org/ns/prov#",
    "kfm": "https://kansasfrontiermatrix.org/ns#"
  },
  "@graph": [
    {
      "@id": "kfm://run/2026-02-20T12:34Z.noaa.abcd1234",
      "@type": "prov:Activity",
      "prov:startedAtTime": "2026-02-20T12:34:56Z",
      "prov:endedAtTime": "2026-02-20T12:45:10Z",
      "prov:wasAssociatedWith": { "@id": "kfm://agent/pipeline/noaa_ingest_v1" },
      "kfm:spec_hash": "sha256:abcd1234..."
    },
    {
      "@id": "kfm://artifact/sha256:1111...",
      "@type": "prov:Entity",
      "kfm:zone": "raw",
      "kfm:uri": "s3://kfm-raw/noaa/stormevents/2026-02.csv",
      "kfm:digest": "sha256:1111..."
    },
    {
      "@id": "kfm://artifact/sha256:2222...",
      "@type": "prov:Entity",
      "kfm:zone": "processed",
      "kfm:path": "data/processed/noaa_ncei_storm_events/2026-02.abcd1234/events.parquet",
      "kfm:digest": "sha256:2222..."
    },
    {
      "@id": "kfm://run/2026-02-20T12:34Z.noaa.abcd1234",
      "prov:used": [{ "@id": "kfm://artifact/sha256:1111..." }],
      "prov:generated": [{ "@id": "kfm://artifact/sha256:2222..." }]
    }
  ]
}
```

Appendix R. Promotion manifest template

Promotion manifests describe a dataset version release and support reproducibility.

```
{
  "kfm_promotion_manifest_version": "v1",
  "dataset_slug": "noaa_ncei_storm_events",
  "dataset_version_id": "2026-02.abcd1234",
  "spec_hash": "sha256:abcd1234...",
  "released_at": "2026-02-20T13:00:00Z",
  "artifacts": [
    { "path": "events.parquet", "digest": "sha256:2222...", "media_type":
  "application/x-parquet" },
    { "path": "events.pmtiles", "digest": "sha256:3333...", "media_type":
  "application/vnd.pmtiles" }
  ],
  "catalogs": [
    { "path": "dcat.jsonld", "digest": "sha256:4444..." },
    { "path": "stac/collection.json", "digest": "sha256:5555..." },
    { "path": "prov/bundle.jsonld", "digest": "sha256:6666..." }
  ],
  "qa": { "status": "pass", "report_digest": "sha256:7777..." },
  "policy": { "policy_label": "public", "decision_id": "kfm://policy_decision/xyz" },
  "approvals": [
    { "role": "steward", "principal": "<id>", "approved_at": "2026-02-20T12:59:00Z" }
  ]
}
```

Appendix S. Threat model checklist (trust membrane and policy)

Use this checklist when reviewing new features.

```
TM-001 Does the frontend ever fetch data directly from object storage or databases?
  Expected: NO (trust membrane)

TM-002 Can a public user infer restricted dataset existence via error behavior?
  Expected: NO (policy-safe errors)

TM-003 Are all downloads and exports checked against policy labels and rights?
  Expected: YES

TM-004 Can Focus Mode be prompt-injected by retrieved documents?
  Expected: mitigations in place (tool allowlist, citation verifier, policy filters)

TM-005 Are audit logs redacted and access-controlled?
  Expected: YES

TM-006 Are pipeline credentials scoped per source and rotated?
  Expected: YES

TM-007 Are processed artifacts immutable by digest?
  Expected: YES

TM-008 Are policy rules tested in CI with fixtures?
  Expected: YES
```

Appendix T. References

Internal KFM sources:

- [IB-v5] KFM_Implementation_Blueprint_v5_Expanded.pdf (2026-02-20)
- [DesignPlan-v2] KFM_Detailed_Design_Engineering_Plan_v2_EmbeddedMetadata_2026-02-19.pdf
- [DocUpgrade-v2] KFM_Document_Upgrade_Refined_2026-02-20_FINAL_v2.pdf
- [DSIB-v1] KFM_Comprehensive_Data_Source_Integration_Blueprint_v1_massive.pdf (2026-02-12)
- [SupportFile] Kansas Frontier Matrix Support File.pdf (compiled extracts)

External standards (for implementation reference):

- W3C DCAT (Data Catalog Vocabulary) v3
- OGC STAC (SpatioTemporal Asset Catalog)
- W3C PROV-O (Provenance Ontology)
- RFC 8785 JSON Canonicalization Scheme (JCS)
- Open Policy Agent (OPA) and Rego policy language

Kansas Frontier Matrix (KFM)

Ultimate Design / Layout / Blueprint / Governance Guide

Map-first • Time-aware • Governed • Evidence-first • Cite-or-abstain

What this is: a consolidated, buildable specification for KFM: architecture, truth-path pipelines, catalogs, policy, evidence resolution, governed APIs, Map/Story UI, Focus Mode, security posture, and promotion gates.

Structure: review → hallucination critique → refined review → blueprint + templates + checklists.

Status vocabulary: CONFIRMED, **PROPOSED**, **UNKNOWN**.

Primary sources: [KFM-GDG-2026], [KFM-IB5-2026].
Generated: February 20, 2026 (America/Chicago)

Table of Contents

Table of Contents 2

Part A. System review (evidence-backed snapshot) 4

A1. What KFM is (and is not) . . . 4

A2. Non-negotiable invariants . . . 4

A3. Governance mechanics: policy, promotion, and audit . . . 6

A3.1 Zones (truth path) . . . 6

A4. Focus Mode: anti-hallucination by design . . . 6

Part B. Self-critique (hallucination and overreach audit) 7

B1. Hallucination risk map . . . 7

B2. Critique summary (what I tightened) . . . 8

Part C. Refined review (tight, no hand-waving) 9

C1. Definition . . . 9

C2. Confirmed vs proposed vs unknown . . . 9

Part D. Definitive blueprint (buildable spec + templates) 10

D1. Architecture reference model . . . 10

D1.1 Component decomposition (recommended build plan) . . . 10

D2. Data model, identity, and time . . . 11

D2.1 Minimal domain entities . . . 11

D2.2 Time axes (recommended) . . . 11

D3. Catalogs and provenance as contract surfaces . . . 12

D3.1 Triplet responsibilities . . . 12

D3.2 EvidenceRef schemes (minimum) . . . 12

D3.3 Profiles (minimum required fields) . . . 13

D4. Policy labels, obligations, and redaction . . . 14
D4.1 Example policy decision (template) . . . 14
D5. Run receipts and audit ledger . . . 15
D5.1 Run receipt (template) . . . 15
D6. Promotion Contract (fail-closed gates) . . . 16
D6.1 Minimal gates (starter) . . . 16
D6.2 Promotion manifest (template) . . . 16
D7. Evidence resolver and governed API contracts . . . 17
D7.1 Evidence bundle (template) . . . 17
D7.2 API surface (illustrative endpoints) . . . 18
D8. Map and Story UI blueprint (trust surfaces) . . . 19
D8.1 Information architecture (recommended) . . . 19
D8.2 Core components (buildable) . . . 19
D8.3 Evidence drawer (required minimum fields) . . . 19
D9. Security posture and threat model checklist . . . 20
D10. Roadmap and Definition of Done (vertical slice) . . . 21
D10.1 Definition of Done (dataset integration ticket) . . . 21
Appendix A. Minimum verification steps (to avoid overreach) 22
Appendix B. Sources and citations (human-readable) 23
Primary internal references used in this document: . . . 23

Part A. System review (evidence-backed snapshot)

A1. What KFM is (and is not)

KFM is a governed, evidence-first, map-first, time-aware knowledge system. It is explicitly not just a map and not just a catalog: public layers, narrative claims, and AI-assisted answers must be traceable to versioned sources, with policy controls enforced in CI and at runtime (see [KFM-GDG-2026]).

A2. Non-negotiable invariants

Invariant	Meaning in practice	Why it matters
Truth path lifecycle	Upstream → RAW → WORK/QUARANTINE → PROCESSED → CATALOG (DCAT+STAC+PROV + run receipts) → projections → governed API → UI.	Reproducibility and auditability.
Trust membrane	Clients never access storage/DB directly; backend logic uses repository interfaces; all access goes through governed APIs applying policy, redaction, and logging.	Without it, policy and provenance are unenforceable.
Evidence-first UX	Every layer/claim opens into evidence: dataset version, license/rights, policy label, provenance chain, checksums, artifact links.	Trust is a first-class surface, not a footnote.
Cite-or-abstain Focus Mode	Answers cite resolvable evidence bundles or abstain; citation verification is a hard gate; every query emits a run receipt.	Mitigates hallucinations and leakage.
Canonical vs rebuildable stores	Object store + catalogs + provenance are canonical; DB/search/graph/tiles are rebuildable projections.	Operational safety and easy re-indexing.
Deterministic identity/hashing	Stable dataset and version IDs based on canonical JSON hashing (e.g., RFC 8785 JCS).	Reliable caching, signing, and version clarity.

Truth Path (data -> catalogs -> governed surfaces)

Trust Membrane (policy and provenance boundary)

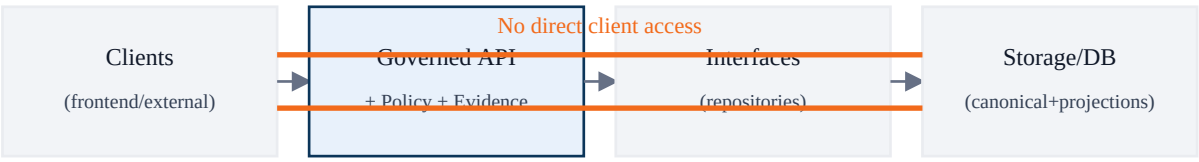

A3. Governance mechanics: policy, promotion, and audit

Governance intent becomes enforceable behavior via the Promotion Contract, policy labels + obligations, and append-only audit trails (see [KFM-GDG-2026]).

A3.1 Zones (truth path)

- **RAW**: immutable acquisition (manifest, artifacts, checksums, license/terms snapshot).
- **WORK/QUARANTINE**: normalization + QA reports + candidate redactions; quarantine blocks promotion.
- **PROCESSED**: publishable artifacts in approved formats with checksums and derived runtime metadata.
- **CATALOG/TRIPLET**: DCAT + STAC + PROV + run receipts (cross-linked and validated).
- **PUBLISHED**: governed runtime surfaces serve only promoted dataset versions.

A4. Focus Mode: anti-hallucination by design

Focus Mode is a governed workflow: policy pre-check → admissible evidence retrieval → evidence bundle creation → answer synthesis → citation verification hard gate → audit run receipt. If citations cannot be verified or policy denies, the system must abstain or reduce scope (see [KFM-GDG-2026]).

Focus Mode control loop (cite-or-abstain)

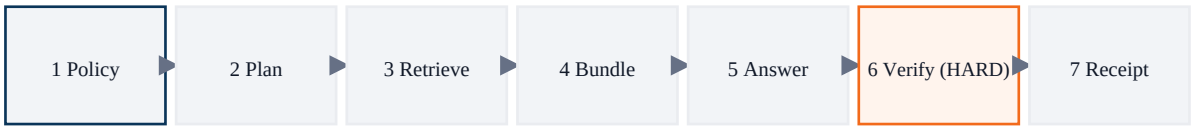

If citations cannot be verified, abstain or reduce scope.

Part B. Self-critique (hallucination and overreach audit)

B1. Hallucination risk map

Rule: statements about invariants/required gates are safe if they are expressed as requirements. Statements about concrete tech choices, repos, or deployed systems are unsafe unless directly evidenced. This audit separates those classes, then tightens language.

Claim	Status	Adjustment
Trust membrane: no direct client → DB/object store; backend uses repository interfaces.	CONFIRMED	Keep as invariant; test it.
Evidence-first UX and evidence drawer everywhere.	CONFIRMED concept	Keep requirement; UI details are PROPOSED.
Cite-or-abstain Focus Mode with audit receipts.	CONFIRMED	Keep behavior; exact DTOs are PROPOSED.
Catalog triplet (DCAT/STAC/PROV) profiles and strict cross-links.	CONFIRMED posture; PROPOSED profiles	Separate posture vs profile schema specifics.
Deployment is modular monolith with modules: api/policy/evidence/catalog/ingest/indexers.	PROPOSED default	Express as a recommended starting topology, not a fact.
Map canvas uses MapLibre GL.	PROPOSED	Treat as replaceable implementation detail.

B2. Critique summary (what I tightened)

- Changed implementation claims into **requirements** or **proposed defaults**.
- Added explicit **verification steps** wherever repository state could differ.
- Kept the system's **fail-closed** posture: if policy/rights/citations are unclear, block promotion/publishing.

Part C. Refined review (tight, no hand-waving)

C1. Definition

KFM is a governed knowledge system where the map and timeline are the primary interface, and every visible layer, story claim, or AI-assisted answer is backed by a resolvable evidence bundle tied to a specific dataset version. It is designed to be reproducible (truth path + deterministic identity) and enforceable (trust membrane + promotion contract + policy labels/obligations).

C2. Confirmed vs proposed vs unknown

- **CONFIRMED**: invariants + contracts (trust membrane, truth path zones, promotion gates, evidence-first UX, cite-or-abstain Focus Mode, canonical vs rebuildable stores, run receipts/audit ledger).
- **PROPOSED**: concrete module decomposition, specific UI component list, endpoint naming, schema profiles (minimum fields).
- **UNKNOWN**: current repo structure, deployed infrastructure, existing datasets, operational runbooks (requires repository + governance review).

Part D. Definitive blueprint (buildable spec + templates)

D1. Architecture reference model

Maintain clean layers: Domain → Use cases → Interfaces → Infrastructure. The key rule: domain logic never talks directly to infrastructure (see [KFM-GDG-2026]).

D1.1 Component decomposition (recommended build plan)

Component	Responsibility	Hard requirements
Pipeline runner + connectors	Fetch/snapshot upstream; write immutable RAW; normalize in WORK; produce QA + candidate redactions; write PROCESSED artifacts.	Artifacts are digest-addressed; acquisition terms captured.
Catalog generator	Emit DCAT/STAC/PROV + run receipts; validate profiles; enforce cross-links.	CI link-check must pass before promotion.
Policy engine	Evaluate allow/deny + obligations + reason codes; shared in CI/runtime.	Default-deny; fixtures-driven tests.
Evidence resolver	Resolve EvidenceRefs to EvidenceBundles; apply policy + redaction; produce bundle digest.	Bounded calls; fail closed.
Governed API	Discovery, tiles, queries, story nodes, focus ask/answer, lineage endpoints.	Audit IDs; policy-safe errors.
UI surfaces	Map Explorer + Story Mode + Focus Mode; evidence drawer everywhere.	No privileged credentials in client; accessibility minimums.
Index builders (rebuildable)	Build projections (PostGIS/search/graph/tiles) from promoted artifacts.	Rebuild from canonical stores; do not store source truth.

D2. Data model, identity, and time

D2.1 Minimal domain entities

Minimum entities: Dataset, DatasetVersion, Artifact, EvidenceRef, EvidenceBundle, PolicyDecision, RunReceipt, StoryNodeVersion, MapState.

> **Deterministic identity:** dataset_version_id and spec_hash should be derived from canonical specs (canonical JSON hashing) to prevent drift and enable caching/signing. Artifacts and evidence bundles are addressed by digest (sha256:...).

D2.2 Time axes (recommended)

- **Event time**: when something happened.
- **Transaction time**: when KFM acquired/published the record.
- **Valid time** (optional): when a statement is considered true (needed for boundary changes and administrative history).

D3. Catalogs and provenance as contract surfaces

Catalogs are canonical interfaces between pipelines and runtime. KFM uses a triplet: DCAT (dataset metadata), STAC (asset metadata), PROV (lineage), cross-linked so EvidenceRefs resolve without guessing (see [KFM-GDG-2026]).

D3.1 Triplet responsibilities

- **DCAT**: title/description/publisher/license/distributions and dataset-level coverage.
- **STAC**: spatiotemporal extents and assets (collections/items).
- **PROV**: which inputs, tools, parameters produced which outputs; environment capture (container digest, git commit).

D3.2 EvidenceRef schemes (minimum)

- **dcat://**... resolves to dataset/distribution metadata.
- **stac://**... resolves to collection/item/asset metadata.
- **prov://**... resolves to run lineage (activities/entities/agents).
- **doc://**... resolves to governed docs and story citations.
- **graph://**... resolves to entity relations (if enabled).

D3.3 Profiles (minimum required fields)

These are **PROPOSED minimums** to make validation strict and predictable (see [KFM-GDG-2026]).

Profile	Minimum required fields
DCAT dataset	dct:title; dct:description; dct:publisher; dct:license or dct:rights; dcat:theme; dct:spatial; dct:temporal; dcat:distribution; prov:wasGeneratedBy; kfm:policy_label; kfm:dataset_id; kfm:dataset_version_id
STAC collection	id/title/description; extent (bbox + temporal); license; link to DCAT; kfm:dataset_version_id; policy label
STAC item	id; geometry or bbox (policy-consistent); datetime or start/end; assets (href + checksum + media_type); links to PROV/run receipt + DCAT distribution
PROV bundle	prov:Activity per run; prov:Entity per artifact; prov:Agent for pipeline and approvals; prov:used + prov:wasGeneratedBy edges; kfm:policy_decision refs; environment capture (container digest, git commit, parameters)

D4. Policy labels, obligations, and redaction

A policy label is the primary classification input. Policy evaluation returns allow/deny plus obligations (redaction/generalization steps) and reason codes for auditing/UX (see [KFM-GDG-2026]).

D4.1 Example policy decision (template)

```
{
  "decision_id": "kfm://policy_decision/xyz",
  "policy_label": "restricted",
  "decision": "deny",
  "reason_codes": ["SENSITIVE_SITE", "RIGHTS_UNCLEAR"],
  "obligations": [
    {"type": "generalize_geometry", "min_cell_size_m": 5000},
    {"type": "remove_attributes", "fields": ["exact_location", "owner_name"]}
  ],
  "evaluated_at": "2026-02-20T12:00:00Z",
  "rule_id": "deny.restricted_dataset.default"
}
```

D5. Run receipts and audit ledger

Every pipeline run and every Focus Mode query should emit a run receipt capturing inputs/outputs/environment/validation and policy decisions; the audit ledger is append-only and treated as a governed dataset (see [KFM-GDG-2026]).

D5.1 Run receipt (template)

```
{
  "run_id": "kfm://run/2026-02-20T12:00:00Z.abcd",
  "actor": {"principal": "svc:pipeline", "role": "pipeline"},
  "operation": "ingest+publish",
  "dataset_version_id": "2026-02.abcd1234",
  "inputs": [
    {"uri": "raw/source.csv", "digest": "sha256:1111"}
  ],
  "outputs": [
    {"uri": "processed/events.parquet", "digest": "sha256:2222"}
  ],
  "environment": {
    "container_digest": "sha256:img...",
    "git_commit": "deadbeef",
    "params_digest": "sha256:3333"
  },
  "validation": {"status": "pass", "report_digest": "sha256:7777"},
  "policy": {"decision_id": "kfm://policy_decision/xyz"},
  "created_at": "2026-02-20T12:05:00Z"
}
```

D6. Promotion Contract (fail-closed gates)

Promotion moves a dataset version from Raw/Work into Processed + Catalog/Lineage, and therefore into runtime surfaces. It MUST be blocked unless artifacts exist and validate (see [KFM-GDG-2026], [KFM-IB5-2026]).

D6.1 Minimal gates (starter)

Gate	Fail-closed check
Identity & versioning	DatasetVersion ID + spec_hash deterministic; promotion manifest exists.
Artifacts	Processed artifacts exist; each has digest; predictable paths; media types recorded.
Catalogs	DCAT/STAC/PROV schema-valid under profile.
Cross-links	All links resolve; asset hrefs exist; EvidenceRefs resolve.
Policy	policy_label assigned; obligations applied; default-deny tests pass.
QA	Validation reports present; failures quarantined.
Audit	Run receipt emitted; audit ledger append; approvals captured where required.

D6.2 Promotion manifest (template)

```
{
  "kfm_promotion_manifest_version": "v1",
  "dataset_slug": "example_dataset",
  "dataset_version_id": "2026-02.abcd1234",
  "spec_hash": "sha256:abcd1234",
  "released_at": "2026-02-20T13:00:00Z",
  "artifacts": [
    {
      "path": "events.parquet",
      "digest": "sha256:2222",
      "media_type": "application/x-parquet"
    }
  ],
  "catalogs": [
    {"path": "dcat.jsonld", "digest": "sha256:4444"},
    {"path": "stac/collection.json", "digest": "sha256:5555"}
  ],
  "qa": {"status": "pass", "report_digest": "sha256:7777"},
  "policy": {"policy_label": "public", "decision_id": "kfm://policy_decision/xyz"},
  "approvals": [
    {"role": "steward", "principal": "<id>", "approved_at": "2026-02-20T12:59:00Z"}
  ]
}
```

D7. Evidence resolver and governed API contracts

Evidence resolution is central: the evidence resolver accepts an EvidenceRef or structured reference, applies policy, and returns an EvidenceBundle (human card + machine metadata + digests + audit references). The UI should be able to use it in <= 2 calls (see [KFM-GDG-2026], [KFM-IB5-2026]).

D7.1 Evidence bundle (template)

```
{
  "bundle_id": "sha256:bundle...",
  "dataset_version_id": "2026-02.abcd1234",
  "title": "Storm event record: 2026-02-19",
  "policy": {
    "decision": "allow",
    "policy_label": "public",
    "obligations_applied": []
  },
  "license": {"spdx": "CC-BY-4.0", "attribution": "Source org"},
  "provenance": {"run_id": "kfm://run/2026-02-20T12:00:00Z.abcd"},
  "artifacts": [
    {
      "href": "processed/events.parquet",
      "digest": "sha256:2222",
      "media_type": "application/x-parquet"
    }
  ],
  "checks": {"catalog_valid": true, "links_ok": true},
  "audit_ref": "kfm://audit/entry/123"
}
```

D7.2 API surface (illustrative endpoints)

Endpoint	Purpose	Policy posture
GET /api/v1/catalog/datasets	Dataset discovery (DCAT + policy labels)	Hide restricted by default; filter by role.
GET /api/v1/datasets/{dataset_version_id}/query	Query slice by bbox/time/filters	Enforce policy; generalized outputs if required.
GET /api/v1/tiles/{layer_id}/{z}/{x}/{y}	Tile delivery	Only policy-safe tiles; cache varies by policy/auth.
POST /api/v1/evidence/resolve	Resolve EvidenceRef to EvidenceBundle	Fail closed if unresolvable/unauthorized.
GET /api/v1/lineage/{dataset_id}	Lineage graph + run receipts	May redact sensitive fields; include commit SHAs where available.
POST /api/v1/focus/ask	Focus Mode Q&A; with citations	Must cite or abstain; log retrieval context.

D8. Map and Story UI blueprint (trust surfaces)

KFM is experienced through map, timeline, and story surfaces. The UI must make trust visible: show dataset version, freshness, license, and policy badges; provide one-click access to evidence and provenance; and preserve map state as part of Story Nodes (see [KFM-GDG-2026], [KFM-IB5-2026]).

D8.1 Information architecture (recommended)

- **Map Explorer** (primary)
- **Stories** (Story Mode)
- **Catalog** (dataset discovery)
- **Focus Mode** (evidence-led Q&A;)
- **Admin/Steward** (restricted governance tools)

D8.2 Core components (buildable)

- **MapCanvas** + **LayerPanel** (policy badge + dataset version per layer)
- **Timeline control** (time window + histogram where possible)
- **Feature inspect** (attributes + evidence refs)
- **Evidence drawer** (shared across Map/Story/Focus)
- **Provenance panel** (run receipts + lineage links)
- **What changed?** diff viewer between dataset versions

D8.3 Evidence drawer (required minimum fields)

- Evidence bundle ID + digest
- DatasetVersion ID + dataset name
- License and rights holder (with attribution text)
- Freshness (last run timestamp) and validation status
- Provenance chain (run receipt link)
- Artifact links (only if policy allows)
- Redactions applied (obligations), with user-facing explanation

D9. Security posture and threat model checklist

KFM's security model is governance-driven: the trust membrane, policy-safe errors, and evidence-first controls must hold even under prompt injection and exfiltration attempts.

Check	Expected
TM-001 Frontend ever fetches directly from storage/DB?	NO
TM-002 Backend bypasses repository interfaces?	NO
TM-003 Restricted dataset inference via errors/timing?	Mitigated; policy-safe errors
TM-004 Exports include license + attribution automatically?	YES
TM-005 Story publishing blocks if rights/citations unclear?	YES
TM-006 Focus Mode prompt injection defense present?	YES (tool allowlist, citation gate, policy filters)
TM-007 Audit logs access-controlled and redacted?	YES
TM-008 Deterministic hashing re-computable in CI?	YES

D10. Roadmap and Definition of Done (vertical slice)

The fastest risk-reducing increment is one end-to-end vertical slice: ingest → catalogs → governed API → Map UI → evidence drawer → story citations → Focus Mode cite-or-abstain, all gated by CI promotion checks (see [KFM-IB5-2026]).

D10.1 Definition of Done (dataset integration ticket)

- Registry entry updated (owner, license, policy_label, cadence, contact).
- RAW acquisition artifacts immutable with manifest + checksums.
- Processed artifacts present with digests and predictable paths.
- DCAT + STAC + PROV schema-valid and cross-linked; link checks succeed.
- Policy decisions recorded; default-deny tests pass; generalized derivatives created if needed.
- Evidence resolver resolves representative EvidenceRefs into EvidenceBundles.
- UI smoke tests: evidence drawer resolves selections; restricted layers denied/generalized.
- Audit: run receipt emitted; audit ledger append; access controls verified.

Appendix A. Minimum verification steps (to avoid overreach)

- **Repo reality check**: confirm actual module layout, build system, CI workflows, and naming conventions before coding.
- **Policy pack review**: confirm policy labels, obligations, and redaction rules; fail closed if unclear.
- **Catalog validation**: run schema validation + cross-link link-checker on promoted dataset versions.
- **Deterministic hashing**: recompute spec_hash and receipt digests in CI to detect drift.
- **UI trust surface check**: verify evidence drawer reachable from map/story/focus and citations resolve before publishing.

Appendix B. Sources and citations (human-readable)

Primary internal references used in this document:

[KFM-GDG-2026] Anonymous. *Kansas Frontier Matrix (KFM) - Definitive Design & Governance Guide (vNext Draft Consolidation)*. 2026-02-20.
[KFM-IB5-2026] Anonymous. *Kansas Frontier Matrix (KFM) - Implementation Blueprint v5 Expanded*. 2026-02-20.
[EXT-DCAT] World Wide Web Consortium (W3C). *DCAT v3: Data Catalog Vocabulary*.
[EXT-STAC] Open Geospatial Consortium (OGC). *SpatioTemporal Asset Catalog (STAC) Specification*.
[EXT-PROV] World Wide Web Consortium (W3C). *PROV-O: The PROV Ontology*.
[EXT-RFC8785] IETF. *RFC 8785: JSON Canonicalization Scheme*.
[EXT-OPA] Open Policy Agent. *OPA + Rego Policy Language*.

Note: External references are listed for interoperability context; normative requirements are governed by KFM internal documents and policies.

Bugs and Insect Species Data Sources

- Global Biodiversity Information Facility (GBIF) – A worldwide biodiversity data aggregator that includes millions of species occurrence records. GBIF can be filtered for Kansas to retrieve insect observations (from museum specimens, field surveys, and citizen science) across the state .
    - Coverage: All of Kansas (and globally), with records spanning the 1800s to present (dates and locations for each occurrence).
    - Access: Open API and bulk downloads (CSV/DwC) via GBIF's portal. Users can query by taxon (e.g. insect orders) and geography (state or coordinates).
    - License: Data are open-access; most records carry Creative Commons designations (CC0 or CC BY) permitting unrestricted use with attribution .
    - Potential Use: Enables mapping of insect distributions and diversity over time. For the Kansas Frontier Matrix, GBIF data can help model species richness, track ecological changes, or train AI to recognize habitat associations of insects.
- iNaturalist Observations – A citizen-science platform with thousands of crowd-sourced insect sightings in Kansas. Participants upload geo-tagged insect photos which are identified by the community.
    - Coverage: Statewide, especially in populated or surveyed areas; data mainly from the last ~10–15 years (continuous updates). Includes date, location, and species for each observation (from butterflies to beetles).
    - Access: REST API and project downloads. Data (CSV/JSON) can be pulled by taxonomic group and region. An account is required for full dataset export.
    - License: By default, iNaturalist observation data are shared under CC BY-NC (allowing use with attribution for non-commercial purposes) . Many contributors opt for even more open licenses, and research-quality records are exported to GBIF.
    - Potential Use: iNaturalist provides real-time time-series of insect occurrences and seasonal activity. These data can enrich models of insect phenology, support pollinator studies, or validate historical species range maps with contemporary sightings.

Temperatures and Climate Data

- NOAA Historical Climate Data (NCEI) – The National Centers for Environmental Information host extensive weather and climate records for Kansas. This includes daily and monthly measurements from station networks (e.g. COOP, airport stations) covering temperature, precipitation, snowfall, etc. .

    - Coverage: Numerous stations across Kansas from the late 1800s to current day, plus statewide aggregated climate normals. Data spans historical periods (19th–20th century) to the present, enabling long-term climate trend analysis.

    - Access: Free download via Climate Data Online (CSV files or JSON via API) . Users can query by station or county; GIS-ready datasets (e.g. gridded climate normals) are also available.

    - License: U.S. government data – public domain (free to use without restriction). NOAA explicitly provides these records for public use .

    - Potential Use: These datasets allow reconstruction of Kansas's historical climate – e.g. tracking 20th-century temperature trends or extreme drought periods. In the Frontier Matrix, they could drive environmental context for agricultural yields, human settlement patterns (e.g. Dust Bowl migrations), or feed AI models predicting future climate scenarios.

- Kansas Mesonet – A network of local weather stations operated by Kansas State University, offering high-resolution climate observations. The Mesonet monitors temperature, rainfall, wind, and soil moisture at 60+ stations across Kansas.

    - Coverage: Modern coverage (since ~1980s–1990s) at a fine spatial scale, with observations every 5 minutes to hourly. Stations are distributed statewide (agricultural experiment fields, etc.), capturing microclimate variations.

    - Access: Open RESTful API delivering data in CSV format . Users can query specific stations and date ranges; up to 3,000 records per request (approx. 8 years of daily data) . The Mesonet site also provides a URL builder for custom queries.

    - License: Free public use with attribution – "Data accessed through mesonet.k-state.edu is free for public use and download" (with request to cite the source) . Data is preliminary (quality-controlled over time).

    - Potential Use: Ideal for fine-grained recent climate analysis – e.g. comparing microclimates between prairie and irrigated cropland, or assessing soil moisture

trends. The Kansas Frontier Matrix can integrate Mesonet data to model localized impacts of weather on crop yields, drought conditions, or even tie climate anomalies to historical events (like floods or heatwaves that influenced settlement).

Wildlife and Animal Population Data

- eBird (Cornell Lab of Ornithology) – A massive open dataset of bird observations, which includes tens of thousands of records from Kansas. Birdwatchers submit checklists with species seen, count, date, and location.
    - Coverage: All Kansas counties, with data from ~2000 to present (and some historical checklists). Provides seasonal and spatial coverage for 450+ bird species in Kansas, often at daily resolution. Trend metrics (population changes) are updated annually .
    - Access: eBird API (for recent sightings) and bulk data downloads. The eBird Basic Dataset (EBD) can be requested (free) which contains all observations; state-level summary products (relative abundance rasters, trend maps) are also available .
    - License: Free for non-commercial research and education use. Users must agree to data terms (eBird requires registration and acceptance of a data use agreement) . Data must be cited and cannot be redistributed in raw form. (eBird also contributes records to GBIF under CC licenses.)
    - Potential Use: eBird provides a rich time series of wildlife data – for example, mapping the decline of prairie-chickens or the spread of invasive bird species in Kansas. These observations can calibrate species distribution models or serve as a proxy for ecosystem health. In the Frontier Matrix, eBird data could animate changes in biodiversity, inform conservation scenario modeling, or augment AI analyses of habitat change.
- Kansas GAP Species Distribution Maps – The Kansas Gap Analysis Project (GAP) produced predicted habitat distribution datasets for the state's wildlife . These GIS layers model the suitable range of various species (mammals, birds, reptiles, etc.) based on land cover and known occurrence records (circa 2000).
    - Coverage: Statewide coverage for dozens of species. For example, GAP provides raster maps (~30–90 m resolution) of potential habitat for all breeding birds, mammals, amphibians present in Kansas as of 2000 . Temporal coverage is a single baseline (no annual change, as it's a habitat model snapshot around

2000).

- Access: Downloadable GIS files (ArcINFO grids and shapefiles) via the Kansas Data Access & Support Center. Metadata and links are available on the Kansas GIS portal . No update is planned (static dataset), but it is in the public domain for use.

- License: Public domain (U.S. Geological Survey funded). The data are freely usable; citation of the Kansas GAP project (2002) is suggested.

- Potential Use: Provides a pre-settlement or early-2000s baseline of where wildlife should occur based on habitat. This is invaluable for historical ecological modeling – e.g., comparing predicted bison or pronghorn ranges with 19th-century accounts. In the Frontier Matrix, GAP maps can be layered with land-use change to identify “gaps” in conservation or to guide AI in understanding historical species distribution versus current observations (from eBird/GBIF).

Census and Demographic Data

- U.S. Census Bureau – Decennial Census & ACS – Authoritative population and housing data for Kansas from the 19th century to today. Decennial censuses (every 10 years) provide official population counts and demographics for 1860, 1870, … 2020, while the American Community Survey (annual/5-year rolling) provides current estimates.

    - Coverage: Historical and current. Decennial census data exist for Kansas as far back as 1860 (pre-statehood territorial census data for 1850s also available) through 2020, at geographies from state, county, down to tracts/block groups (in later years). The ACS covers roughly 2005–present with detailed socio-economic variables.

    - Access: Modern data (1990s–present) can be queried via the Census API or downloaded from data.census.gov in CSV. Historical data (1790–1980s) are accessible via archives or specialized portals (see NHGIS below). Boundary files (shapefiles for counties, tracts, etc.) are also published by the Census.

    - License: U.S. Census data are public domain (no license restrictions) – “Open Government Data”. The Census API terms mainly require attribution and avoid misuse, but the data itself is free to reuse.

    - Potential Use: Demographic data is fundamental for historical modeling. It enables analysis of population growth, density changes on the frontier, urbanization patterns (e.g. the rise/fall of railroad towns), or mapping ethnic

composition over time. In the Kansas Frontier Matrix, census stats can drive visualizations of settlement expansion, power AI predictions of future demographic shifts, or correlate population with environmental factors (like how drought influenced rural depopulation).

- IPUMS NHGIS (National Historical GIS) – A curated repository of historical census statistics and GIS boundary files . NHGIS provides easy access to aggregated census data for every decade, along with consistent digital maps for historical county and tract boundaries.

    - Coverage: 1790 through present for decennial census data . For Kansas, NHGIS includes territorial census data, all decennial counts, and even historical agriculture and economic census tables. Geographically, one can retrieve data at state, county, or even enumeration district levels (depending on year). Boundaries reflect the changing county lines (e.g. the many county splits in Kansas's 19th-century).

    - Access: Free with registration. Users select desired years, geographic levels, and data tables on the NHGIS website, then download a custom zip with CSV data and matching shapefiles. NHGIS also offers time-series tables that link variables across decades.

    - License: Free to use for any purpose, but users must agree to cite the source. The data originates from public domain census records; IPUMS asks for attribution for the compilation. There are no commercial restrictions.

    - Potential Use: NHGIS simplifies integrating historical census data into GIS. For example, one can map population density on the Kansas frontier in 1870 versus 1970 using authentic county outlines. It's invaluable for creating longitudinal models – the Frontier Matrix can use it to animate demographic change, train AI on how socio-economic variables (like rural farm populations) shifted over 150+ years, or to contextualize other datasets (e.g. overlaying climate or land use with population change).

Agriculture and Land Use Data

- USDA NASS QuickStats (Crop and Yield Data) – The National Agricultural Statistics Service's open database of U.S. agricultural data. QuickStats contains county-level and state-level statistics on crop acreage, yields, production, livestock numbers, economics, etc., from both annual surveys and the 5-year Census of Agriculture .

- Coverage: Kansas statewide and by county, with time series for many variables. Data spans the late 20th century to present (generally from ~1980s onward for surveys; 1997, 2002, 2007, 2012, 2017 Ag Censuses included). For example, one can retrieve the wheat yield per acre for every Kansas county for each year, or the number of farms by size class.

- Access: QuickStats API (JSON/CSV responses) and an online query tool . The API allows filtering by commodity, geographic area, year, etc. Bulk downloads of certain datasets (like ag census) are also on data.gov.

- License: Creative Commons CC0 (Public Domain) – NASS designates QuickStats data as open for public use without restriction .

- Potential Use: This is crucial for modeling agricultural trends. In the Frontier Matrix, NASS data can be used to plot the rise and fall of different crops (e.g. corn vs. sorghum acreage over decades), analyze yields relative to climate events, or train AI models to predict future crop outputs. It also helps tie historical events (like the Dust Bowl) to hard data (crop failures, livestock losses).

- USDA Cropland Data Layer (CDL) – An annual land cover raster that maps crop types across every field in Kansas . The CDL is produced from satellite imagery each year, identifying which crops (corn, wheat, soybeans, etc.) or land cover (forest, grass, urban) are present on each 30×30 m pixel.

    - Coverage: Complete statewide coverage. Available yearly from 2008 to present for Kansas (prior to 2008, the program covered fewer states) . Each year's layer is a georeferenced GeoTIFF, enabling time-series analysis of land use change.

    - Access: Downloadable via the CropScape web portal or directly from data.gov (as GeoTIFF per state per year) . Also accessible through a web map service (WMS) for GIS software . File size is significant but manageable (~50–100 MB per state-year).

    - License: Public domain (CC0) – NASS publishes the CDL as open data . No use restrictions; they encourage citation.

    - Potential Use: The CDL provides a richly detailed view of agricultural land use over time. For the Frontier Matrix, this enables mapping the expansion or contraction of certain crops (like the westward march of soybean cultivation, or areas converted from wheat to sorghum due to climate). It's also valuable for environmental modeling – e.g. relating crop rotations to soil health or carbon stock. AI applications might include classifying land cover changes or forecasting crop distribution under future scenarios.

Migration and Immigration Data

- Census Migration & Immigration Statistics – The U.S. Census Bureau tracks both domestic migration and international immigration characteristics which can be focused on Kansas. Decennial censuses in historical periods asked for place of birth (revealing how many Kansans were born abroad or in other states), and modern American Community Survey data report foreign-born populations, year of arrival, and mobility (moves in the past year).

    - Coverage: At the state and county level, one can obtain net migration figures and immigrant demographics. Historical example: the 1900 Census reported the number of Kansas residents born in Germany or Sweden (capturing immigrant communities). Contemporary data (ACS) detail, say, how many people moved into a given Kansas county in the last year, and from where (other state or abroad).

    - Access: Decennial data is available via NHGIS or IPUMS (for historical place-of-birth tables). ACS data (e.g. 5-year estimates) are accessible through the Census API – for instance, tables showing migration flows or nativity by county. The Census Bureau also occasionally releases specialized migration flow reports and tools on their website.

    - License: Public domain. As with all Census data, these stats are open-access for use and redistribution.

    - Potential Use: These data illuminate settlement patterns and population turnover. In a historical context, one could map the influx of settlers by origin (showing which Kansas counties had many residents from, say, Ohio or Germany in 1880). For modern times, one can track urbanization or youth out-migration from rural counties. Such patterns are vital for the Frontier Matrix to simulate how population shifts respond to economic booms, environmental pressures (e.g. the 1930s exodus), or policy changes. They also aid AI in correlating migration with factors like climate or crop failure.

- IRS Migration Data (County-to-County Flows) – An underutilized dataset from the IRS Statistics of Income division that captures yearly population movement via tax return address changes. It provides granular data on how many households and people moved into and out of each county (and state), along with their origin/destination .

    - Coverage: Nationwide, with specific files for Kansas. Annual data from 1991 through 2022 are available . For each Kansas county, one can see how many taxpayers left for each other state or county and how many arrived (inflows/outflows), as well as the aggregate Adjusted Gross Income of those

movers . This offers a quantitative view of migration trends in the late 20th and early 21st century.

- Access: Downloadable Excel/CSV files by year from the IRS website (the IRS provides state-level spreadsheets listing all county pairs) . No API, but data are structured and machine-readable. For earlier years (1990s), data come as zip files containing all states .

- License: As a U.S. government dataset, it's in the public domain. IRS makes it available for public use; there is no personal info (only aggregated counts).

- Potential Use: The IRS migration files are excellent for analyzing county-level migration patterns – e.g. which rural Kansas counties are losing population to urban areas, or net gains due to inbound retirees. They can complement Census estimates by providing year-by-year flow dynamics. In the Frontier Matrix, these could feed models of economic or climate-driven migration (for instance, out-migration from drought-hit counties) and help evaluate policy impacts (did a new interstate or industry park alter migration?). It's a rich dataset to train AI on socio-economic drivers of movement within Kansas.

Land Ownership and Land Patents

- BLM General Land Office (GLO) Records – A comprehensive archive of federal land patents and survey records for Kansas and other public-land states. This includes Homestead Act patents, railroad grants, military bounty land warrants, etc., documenting the transfer of land from the federal government to individuals or organizations .

    - Coverage: Historical (19th and early 20th century). Kansas, being a public domain state, has GLO patent records from the 1850s (Territorial Kansas surveys) through the early 1900s. Over 160,000 Kansas land patent records are available, each with the patentee name, date, legal land description (Township-Range-Section), and often an image of the original patent document.

    - Access: The GLO online database allows searching by person, location, or patent type. Results include digitized images (PDF) of original patents and metadata (which can be downloaded as XML). There is also an interactive map for patented land parcels. Bulk data download isn't directly provided on the site, but third-party compilations or scraping can retrieve large datasets. Coordinate data can be derived by linking to the Public Land Survey System grid.

    - License: Public domain. These are historic government records and free to use. The Bureau of Land Management encourages using the site for research and

provides the data openly (with the caveat that not every single record may be online) .

- Potential Use: GLO records are key to reconstructing settlement. They enable mapping original homesteads – for example, identifying when and where land was first claimed under the Homestead Act in western Kansas. In the Frontier Matrix, this could animate land acquisition over time, link family histories to specific plots, or feed models about how quickly the frontier filled up. Moreover, by intersecting patent maps with soil or rainfall data, one could analyze outcomes (did homesteads on marginal land get abandoned more often?). This data brings an important human dimension: actual names and dates of Kansas's frontier land owners.

- Public Land Survey System (PLSS) Grid – Geospatial data delineating the township, range, and section lines across Kansas. The PLSS is the land subdivision system used to describe locations of land parcels in Kansas (e.g., "T15S R21W Section 14"). Digital PLSS shapefiles provide the boundaries of every section (roughly one-mile squares) and their identifiers.

  - Coverage: Statewide, static. The PLSS was established in Kansas starting in the mid-19th century; the dataset represents the surveyed grid that covers virtually all of Kansas (except some river irregularities).

  - Access: Available through BLM's Geospatial Hub and Kansas GIS repositories. The BLM provides a national PLSS polygon layer (which can be filtered to Kansas) and often state-specific downloads. Data formats include Esri Shapefile and Geodatabase. Each section polygon has attributes for township, range, section number, meridian, etc.

  - License: Public domain (U.S. Government data). No license restrictions.

  - Potential Use: The PLSS grid is the spatial "key" for integrating historic land records. By overlaying GLO patent data on the PLSS, one can create maps of homesteaded sections. In a broader sense, the PLSS also helps standardize other data (e.g., soil and water datasets are sometimes indexed by township-range as well). For the Frontier Matrix, having the PLSS allows a user to query data by legal land description – important for historians and genealogists. It also provides a base layer to visualize how Kansas was partitioned and settled section by section.

Crime and Justice Data

- FBI Uniform Crime Reports / NIBRS Data – The FBI's national crime statistics programs compile data for Kansas from local law enforcement agencies. This includes historical Uniform Crime Reporting (UCR) summary data (index crimes, arrests, etc.) and the modern National Incident-Based Reporting System (NIBRS) which has granular incident-level data.

    - Coverage: Statewide and city/county level crime data from the mid-20th century to present. Summary UCR data (violent/property crime rates) exist annually back to the 1960s for Kansas. NIBRS incident data is available in detail from the 1990s/2000s onward (with Kansas law enforcement now fully NIBRS compliant). The data cover offense counts, clearance rates, arrests, and more.

    - Access: The FBI's Crime Data Explorer (CDE) provides an open API to query crime stats . One can retrieve data by state, year, offense type, etc., in JSON or CSV. Bulk datasets (e.g. entire NIBRS files) can be downloaded from FBI or the National Archive of Criminal Justice Data. Additionally, the Kansas Bureau of Investigation publishes annual Crime Index Reports (in PDF/Excel) which can be used for local details.

    - License: Public domain (U.S. Dept. of Justice data). The FBI Crime API is explicitly marked as public data (no copyright) . Users are free to reuse and share the statistics.

    - Potential Use: Crime data can be woven into the socio-cultural narrative of Kansas. Historically, one could examine crime rates during boom vs. bust periods (did the Dust Bowl or 1980s farm crisis coincide with rising crime?). Spatially, data could highlight urban-rural differences or justice system resource needs. For the Frontier Matrix, crime statistics add a layer of social context – useful for AI models that might predict community well-being or for overlaying with demographic and economic shifts. They also enable civic engagement apps, e.g. showing long-term trends in public safety across Kansas communities.

Archaeological and Historical Sites Data

- Kansas Historic Resources Inventory (KHRI) – A database of historic and archaeological sites surveyed in Kansas, maintained by the State Historic Preservation Office. It includes information on buildings, structures, objects, and some archaeological sites that have been formally recorded . Each entry typically has location coordinates, a description, and notes on significance.

    - Coverage: Thousands of historic properties across Kansas. The focus is on built environment (e.g. historic schools, churches, homesteads) and known

archaeological sites (especially those eligible for registers). Temporal coverage spans indigenous archaeological sites (pre-1850s) to structures from the more recent past (mid-20th century) that have historic value.

  - Access: Web interface with search and map (KHRI Online) . Non-registered users can search and view all approved records . Data includes basic metadata and often photographs or survey forms. For GIS usage, data can be requested or exported for public sites. (Exact coordinates of sensitive archaeological sites may be generalized to protect them .)

  - License: The data is publicly accessible for education and reference . As a state government resource, it's effectively open, though users must refrain from using it to disturb sites (Kansas law exempts precise archeological locations from open-records requests for protection). In summary: free to use with proper attribution, but some location details are restricted.

  - Potential Use: KHRI is a treasure trove for historical mapping. It can show where old forts were, trails, Native American village sites, or early settler structures. In the Frontier Matrix, these points of interest allow layering cultural heritage onto environmental and demographic data – for instance, visualizing how early trails and trading posts align with water sources or how settlement progressed county by county. AI could use the textual descriptions in KHRI for natural language processing, extracting common themes of frontier architecture or settlement challenges.

- National Register of Historic Places (NRHP) GIS Data – A nationwide open dataset of historic places listed on the National Register, including many in Kansas. The National Park Service provides spatial data (points/polygons) for NRHP listings , which range from well-known landmarks to archaeological districts.

  - Coverage: Over 1,600 NRHP listings in Kansas (as of recent counts), from Dodge City's historic district to Native American archaeological sites. Includes site name, location (lat/long), and listing information. This is historical-cultural data spanning all periods of Kansas history. Updated as new sites are listed.

  - Access: Downloadable from NPS's Cultural Resources GIS portal or via data.gov. The data are provided as a feature service and shapefiles, with attributes like reference number, name, address, and area. Because not all NRHP sites (especially archaeological ones) have precise coordinates publicly shown, some may be generalized points.

  - License: Public domain (federal government publication) . The NPS data is explicitly "public, non-restricted data" for NRHP locations, meant for broad use in

mapping and analysis.

- Potential Use: NRHP data flags sites of official historical significance. This can be used to ensure the Frontier Matrix includes cultural landmarks in its model – for example, identifying counties rich in heritage resources or using it as a layer in predictive modeling (perhaps an AI that predicts which sites might deserve future preservation based on patterns). It also aids story-telling: one can click on a county and immediately see notable historic sites or dig into narratives of an old frontier fort, integrating story maps with data.

Soil and Vegetation Data

- NRCS Soil Survey (SSURGO Database) – Detailed soil maps for Kansas, produced by the USDA Natural Resources Conservation Service. SSURGO contains geospatial soil polygon maps and extensive attributes (soil type, texture, chemistry, capability class, etc.) for each mapping unit . Essentially, it's the digital form of county soil surveys, representing over a century of soil data collection.

  - Coverage: Statewide at fine resolution (typically 1:12,000 scale). SSURGO is organized by soil survey areas (often by county) – all Kansas counties have complete coverage. Temporal: the data reflects surveys done mostly 1960s–2000s, but soil properties are relatively stable; updates occur as new surveys are completed or data refined.

  - Access: Available through the Web Soil Survey and for download via NRCS's Soil Data Access portal. One can download GIS shapefiles or a geodatabase for each county or use the Soil Data Access API for querying specific properties. NRCS also offers a gridded SSURGO (gSSURGO) which compiles the data statewide in raster form. Kansas's GIS portal provides ready-to-use SSURGO layer packages as well .

  - License: Public domain. Soil surveys are government data and freely available. NRCS even notes the SSURGO portal is "license-free, open-source" for access .

  - Potential Use: Soils are fundamental for environmental and agricultural modeling. By integrating SSURGO, the Frontier Matrix can, for instance, correlate crop yields with soil types (explaining why certain areas grew wheat vs. corn), or model erosion and land degradation over time. Soil data can inform AI models predicting suitability for crops or forecasting how past farming practices impacted land (e.g. areas prone to Dust Bowl erosion). It also enriches historical reconstructions – knowing soil quality helps explain settlement patterns

(homesteaders avoided certain sandy soils, etc.).

- Land Cover and Vegetation – NLCD and Historical Data – Comprehensive land cover information for Kansas is available from the National Land Cover Database (NLCD), as well as historical vegetation maps. NLCD is a Landsat-derived product classifying each 30m cell as agriculture, forest, grassland, urban, etc., and is updated regularly . In addition, Kansas-specific efforts (e.g. Kansas Biological Survey's 2017–2018 land cover map or the Kansas GAP land cover from the 1990s) provide snapshots of vegetation.

    - Coverage: NLCD: Nationwide including Kansas, with products for 2001, 2006, 2011, 2016, and now produced annually (latest 2019/2020) . Resolution 30m. Classes include cropland, various forest types, grassland, wetland, developed, etc. Historic maps: A 1992 Kansas land cover (from GAP) and modern high-resolution (10m) maps from satellite imagery are available. These cover the entire state as well.

    - Access: NLCD can be downloaded free from the MRLC website or data.gov, as GeoTIFF grids. Annual NLCD layers (for recent decades) are also accessible. Kansas-specific land cover datasets can be obtained from the Kansas Applied Remote Sensing Program (KARS) or DASC as GeoJSON/GeoTIFF .

    - License: Public domain for NLCD (created by USGS; data is explicitly stated as nationwide public domain land cover information). State-derived products from KARS/KBS are generally open for public use as well (often funded by state/federal grants).

    - Potential Use: Land cover data allows the Frontier Matrix to incorporate the natural landscape context. For example, one could track the decline of prairie over decades or identify regions of forest regrowth. NLCD's consistent time-series supports analyzing land cover change (like urban sprawl around Kansas City or shifts in rangeland vs. cropland). In modeling, this is vital for habitat and environmental analyses: an AI model could predict wildfire risk based on vegetation cover or simulate how land cover might look under future climate or policy scenarios. Historic vegetation reconstructions (using early survey records or the 1850s Kansas land cover maps) could be used to compare then-and-now, enriching historical storytelling about the Kansas environment.

Disaster and Extreme Weather Data

- FEMA Disaster Declarations – Open data on all federally declared disasters in Kansas since 1953. FEMA's Disaster Declaration Summaries dataset lists every declared major disaster, emergency, or fire management assistance event, along with dates, affected

counties, and assistance programs . Examples include presidential disaster declarations for floods, tornadoes, severe storms, droughts, etc.

  - Coverage: Records from 1953 to present, updated continually. Each entry includes the disaster type (e.g. Tornado outbreak), incident date range, the counties included (note: county-level tagging is fully available post-1964) , and the declaration date. It also notes which federal assistance programs were activated (individual assistance, public assistance, etc.).

  - Access: FEMA provides this dataset on its website and via API (OpenFEMA). It can be downloaded as CSV/JSON in entirety or filtered via API queries (e.g., all disasters in Kansas, or all drought declarations nationwide). The data is also on Data.gov for bulk access.

  - License: Public domain (U.S. DHS/FEMA data). OpenFEMA datasets are intended for public use . The license is essentially government works free for reuse (FEMA just asks that users cite and be aware data is raw and may contain minor errors).

  - Potential Use: This is crucial for understanding the history of disasters in Kansas – from the 1951 Great Flood to Greensburg's 2007 tornado to recent wildfires. In the Frontier Matrix, one can layer disaster events over time to see patterns (e.g., increasing frequency of certain disaster types) or link them to other datasets (population or land use changes post-disaster). It also enables training models on the impacts of disasters – for instance, an AI could use past declaration data to predict which areas are most at risk or how long recovery takes, feeding into future resilience planning.

- NOAA Storm Events Database – A detailed record of severe weather occurrences compiled by the National Weather Service, including tornadoes, hail, thunderstorms, winter storms, and other extremes in Kansas. This database captures event specifics such as date, location (county or lat/long), magnitude (e.g. tornado EF rating, hail size), injuries/fatalities, and damage costs.

  - Coverage: 1950s to present for tornadoes (the database starts tornado records in 1950), and later start dates for other phenomena (hail, wind, etc. from 1955; floods and others by the 1990s). It is continuously updated by each NWS forecast office. Kansas data is comprehensive – e.g., every tornado that has touched down in Kansas is recorded, as well as major storms.

  - Access: NOAA's Storm Events query tool allows filtering by state, county, event type and downloading results in CSV . Bulk files by year are also downloadable. The data can be integrated into GIS by geocoding county-based events or using

provided coordinates for point-specific events.

- License: Public domain (U.S. Government). NOAA makes this data freely available for research and does not restrict usage.

- Potential Use: The Storm Events data brings granular disaster modeling capability. For example, one can map all tornado tracks in Kansas or analyze seasonality and trend (are winters becoming more extreme?). In the Frontier Matrix context, this could enrich historical timelines (e.g., marking the 1974 tornado outbreaks or 1993 Great Flood on a map alongside other data). It enables AI models to learn from past patterns – for instance, correlating soil moisture or temperature anomalies with subsequent disaster frequency, or predicting damage risks for different regions. Coupling storm data with FEMA declarations also distinguishes between event occurrence and its socio-economic impact (not every storm becomes a federal disaster, and that threshold itself is informative).

Historical Narratives and Oral Histories

- Kansas Memory (Kansas Historical Society Digital Archives) – A vast online repository of primary sources related to Kansas history. It contains nearly 600,000 items including photographs, letters, diaries, maps, newspaper clippings, and some audio recordings . While not a “dataset” in the numeric sense, it is a rich data source for textual and visual material that can be mined by AI or indexed alongside quantitative data.

    - Coverage: Spans the entirety of Kansas history – e.g., territorial-era letters, 19th-century Native American accounts, frontier town photographs, through mid-20th-century documents. Geographically covers communities across the state. There are also oral history transcripts and reminiscences in the collection (e.g., WPA era pioneer interviews).

    - Access: Web portal (kansasmemory.org) with search by keywords, topics, or collections. Items can be viewed and downloaded (images as JPEGs, documents as PDFs). Metadata (title, date, description, location) is provided for each item. For large-scale analysis, one might use the API or contact KHS for bulk data – but generally, scraping the site is required for big data uses.

    - License: Materials are made available for research and educational use. Many items are in the public domain (especially older than 1920s). KHS asks for attribution and may have usage restrictions on some copyrighted 20th-century materials. Overall, much of Kansas Memory’s content can be used freely for non-commercial purposes, and fees may apply for high-resolution or commercial

use.

    - Potential Use: Kansas Memory provides the human stories behind the data – useful for contextualizing and enriching the Frontier Matrix. For instance, one could associate a spike in 1870s immigration (from census data) with actual letters from immigrant families (from Kansas Memory) describing their journey. In AI terms, the text corpora in letters and newspapers could be used for natural language processing – training models to extract common themes (drought, war, prosperity) over decades. Image archives allow computer vision AI to learn from historical photographs (e.g., to identify farming scenes vs. urban scenes). By integrating these narratives, the platform moves beyond numbers, ensuring that voices of Kansans (including historically marginalized ones) are part of the modeled future.

- Kansas Oral History Project (KOHP) – A non-profit initiative that has recorded and transcribed interviews with notable Kansans, particularly focused on those involved in public policy and governance. This is a goldmine of firsthand accounts on late 20th-century Kansas politics, education, water policy, and more . Collections include interviews with former legislators, governors, community leaders, and diverse citizens.

    - Coverage: Approximately 170+ interviews (and growing). The focus is roughly 1960s–2010s in subject matter (e.g., participants recount civil rights in Kansas, school finance battles, rural development, etc.). While the emphasis is on governance, the content often delves into social and economic history as told by the interviewees.

    - Access: All interviews are available on the KOHP website (ksoralhistory.org) as streaming audio and text transcripts. They are organized into thematic collections ("Women Then and Now," "Kansas Governors," "Reflections on Water," etc.) . Transcripts can be read online or downloaded. KOHP also provides summaries and keywords for each interview.

    - License: The collected interviews (text and audio) are offered under a Creative Commons Attribution-NonCommercial-NoDerivs (CC BY-NC-ND 4.0) license . This means they can be freely used and shared for non-commercial purposes with attribution, but not altered or used commercially without permission.

    - Potential Use: These oral histories add a qualitative dimension to data-driven modeling. They allow the Frontier Matrix to incorporate insights like why certain policies succeeded or how communities felt during pivotal moments (e.g., the 1980s farm crisis from a legislator's perspective). In practical terms, one could use NLP on the transcripts to extract sentiment trends or key concerns over time (water rights, education funding, etc.), complementing the quantitative indicators. For AI enrichment, training language models on this Kansas-specific corpus

could enable more nuanced storytelling or Q&A about Kansas history. Moreover, integrating audio clips or quotes into the platform can make the past “speak” to users in a compelling way, ensuring that any future scenarios we model remain grounded in the lived experiences and voices of Kansans.

Sources: The compiled information above is drawn from a range of open data repositories and official resources, including GBIF , iNaturalist documentation , NOAA climate archives , Kansas Mesonet API docs , Cornell’s eBird data release notes , the Kansas GAP final report , U.S. Census Bureau datasets and NHGIS descriptions , USDA NASS QuickStats metadata , USDA Cropland Data Layer info , FEMA’s OpenFEMA data portal , NOAA Storm Events tool , Kansas Historic Resources Inventory site , NPS National Register data page , NRCS soil database description , NLCD land cover database documentation , and the Kansas Oral History Project and Kansas Memory archives . These sources are all open-access and provide the foundation for integrating Kansas’s past, present, and future into the Kansas Frontier Matrix platform.

KFM Pulse Integration Pack

Cultivated, integration-ready ideas distilled from “Pulse 12-20-25”

Prepared 2026-02-15 (America/Chicago)

Purpose: Convert the Pulse document’s concepts into a set of build-ready modules: each with scope, required artifacts (schemas/policies), integration points, and validation gates that fit KFM’s evidence-first governance.

Primary source: Pulse 12-20-25 (37 pages).

Module	What it is	Outputs / why it matters
A	Trust-aware calibration for low-cost PM■.■ sensor networks	Calibrated time-series + trust scores; ML correction + drift handl
B	STAC↔Zarr discovery upgrades (chunk metadata, multiscale, EOPF patterns)	Standardized array metadata + multiscale conventions for c
C	Neo4j graph migration enforcement	Graph change-control: intent, invariants, diffing, rollback, PROV
D	Multi-source incident/work-zone fusion pipeline (WZDx + 511 + GTFS-RT)	Canonical incident model + dedup/QC/promotion policy +
E	Deterministic data runs (lakeFS + DVC + LangGraph + OPA + Dagster)	Reproducible run harness + policy gates + single-command rollb

Table of Contents

Table of Contents 2

1. Executive Summary 5

1.1 What “ready to integrate” means here . . . 5

2. Module A — Trust-Aware Calibration for Low-Cost PM■.■ Sensors 6

2.1 Proposed KFM capability . . . 6

2.2 Inputs, outputs, and contracts . . . 6

2.3 Buildable baseline (first slice) . . . 6

2.4 Trust score + adaptive logic (second slice) . . . 6

2.5 Validation gates . . . 7

2.6 Monitoring signals . . . 7

3. Module B — STAC↔Zarr Discovery Upgrades (Chunk Metadata, Multiscales, EOPF Patterns) 8

3.1 Proposed KFM capability . . . 8

3.2 Minimum viable asset metadata (kfm:array) . . . 8

3.3 Multiscale representation decision tree . . . 8

3.4 EOPF scene-product asset layout (copy/paste pattern) . . . 9

3.5 “STAC in Zarr” (self-describing stores) . . . 9

3.6 Validation gates . . . 9

3.7 Buildable baseline (first slice) . . . 9

4. Module C — Neo4j Graph Migration Enforcement (Governed Change Control) 10

4.1 Migration taxonomy (drives enforcement rules) . . . 10

4.2 Baseline capture (beyond simple counts) . . . 10

4.3 Invariant classes (what CI enforces) . . . 10

4.4 Diffing + blast radius control . . . 10

4.5 Reversibility guarantees . . . 11

4.6 PROV alignment . . . 11

4.7 Migration declaration contract (required) . . . 11

4.8 CI failure semantics . . . 11

4.9 Buildable baseline (first slice) . . . 11

5. Module D — Multi-Source Incident / Work-Zone Fusion (WZDx + ArcGIS + GTFS-RT) 12

5.1 Canonical IncidentEvent model (publish contract) . . . 12

5.2 Ingestion patterns per source . . . 12

5.3 Promotion policy (confidence ladder) . . . 12

5.4 Conflict resolution rules (examples) . . . 12

5.5 End-to-end operational flow . . . 13

5.6 Known failure modes (design for these) . . . 13

5.7 Required deliverables (artifact pack) . . . 13

5.8 Buildable baseline (first slice) . . . 13

6. Module E — Deterministic Data Runs (lakeFS + DVC + LangGraph + OPA/Rego + Cosign) 14

6.1 Core idea (why this matters) . . . 14

6.2 Drop-in repository skeleton (prototype) . . . 14

6.3 Gate: OPA/Rego policy . . . 14

6.4 Rollback: single command . . . 14

6.5 CI integration (first slice) . . . 15

6.6 Buildable baseline (first slice) . . . 15

7. Cross-Module Integration Roadmap (Thin Slices) 16

7.1 Slice 0 — Common governance harness (shared) . . . 16

7.2 Slice 1 — First promotable dataset in each domain 16

7.3 Slice 2 — Add confidence & scale mechanics 16

7.4 Definition of Done (applies to each slice) 16

8. Artifact Inventory (What to Create) 17

9. Assumptions, Risks, and Verification Steps 18

9.1 Assumptions used (because Pulse leaves these as context-dependent) 18

9.2 Key risks surfaced in Pulse (and how to mitigate) 18

9.3 Minimum verification steps before implementation 18

1. Executive Summary

Pulse 12-20-25 is effectively a set of “drop-in” engineering patterns: it describes (1) trust-aware calibration for sensor networks; (2) practical STAC↔Zarr conventions for discoverability of chunked/multiscale datacubes; (3) governed graph migration enforcement for Neo4j; (4) a multi-source near-real-time incident fusion pipeline; and (5) an end-to-end deterministic data-run harness with policy gates and rollback. This document repackages those ideas into implementation modules with explicit contracts, artifacts, and CI/policy gates.

Integration stance: treat every module’s “publish” step as fail-closed—no promotion without validated schemas, provenance, and explicit rollback plan.

1.1 What “ready to integrate” means here

- A clear module boundary and purpose statement.
- A minimal “first slice” that can be built and tested independently.
- The artifact list required for promotion (schemas, policy rules, run receipts/provenance).
- Validation gates (QC thresholds, invariants, and hard/soft failure semantics).
- Operational notes (failure modes, rollback strategy, monitoring signals).

2. Module A — Trust-Aware Calibration for Low-Cost PM■.■ Sensors

Pulse argues that “static” correction equations degrade in real deployments because low-cost optical PM sensors are sensitive to humidity and temperature, drift over time, vary unit-to-unit, and transfer poorly across climates and aerosol regimes. The recommended approach combines two complementary layers: (1) meteorology-aware ML correction anchored to reference stations, and (2) a trust-aware adaptive layer that detects drift/outliers and downweights low-trust sensors during fusion.

2.1 Proposed KFM capability

- **SensorCalibration service**: produces a calibrated PM■.■ time-series per sensor plus a time-varying trust score.
- **NetworkFusion service**: fuses multiple sensors into gridded / corridor / city aggregates, weighted by trust.
- **CalibrationCatalog profile**: publishes model and feature metadata so results are reproducible and reviewable.

2.2 Inputs, outputs, and contracts

Item	Contract / notes
Inputs	Raw sensor PM■.■ (or particle counts), RH/T (required), optional wind/pressure; reference monitor PM■.■
Outputs	Calibrated PM■.■ (with uncertainty bands if available), trust_score(t), drift flags, and calibration model meta
Minimum provenance	Run receipt includes: input dataset IDs/digests, model family + hyperparameters, feature transforms, trainin

2.3 Buildable baseline (first slice)

Start with a “strong default” that is accurate, explainable, and cheap to run; add adaptive trust logic once the core data contracts are stable.

- Normalize timestamps (hourly is a common baseline); drop obvious sensor inconsistencies (e.g., dual-channel disagreement) and extreme outliers.
- Train a tree-based model (Random Forest / Gradient Boosting) using: raw PM■.■, RH, temperature, optional wind/pressure, and simple time features.
- Evaluate by regimes (RH bins, seasons) and produce a per-sensor calibration report (metrics + residual plots).
- Add a low-concentration safeguard: classify near-zero vs non-zero, then regress only for non-zero periods to reduce false positives in clean-air hours.

2.4 Trust score + adaptive logic (second slice)

- Define **trust indicators**: agreement with reference/high-trust peers, stability/noise, responsiveness, and consensus alignment in the local network.
- Compute trust_score(t) in rolling windows; flag sensors whose trust drops below threshold for N windows.
- During fusion, downweight low-trust sensors; optionally increase correction complexity only for low-trust sensors.

2.5 Validation gates

- **Hard fail**: missing RH/T; missing reference linkage for supervised calibration runs; non-monotonic timestamps; schema mismatch.
- **Soft fail / review**: abrupt trust collapse without corroborating maintenance/location change; calibration drift exceeding threshold; negative corrected values beyond tolerance (may clamp to 0 per common post-processing).

2.6 Monitoring signals

- Network median trust trend (should be stable unless environmental shift or widespread drift).
- Residual distribution by RH bin (humidity artifacts are usually the first place failures appear).
- Coverage of co-location window conditions vs deployment conditions (diversity beats length).

3. Module B — STAC↔Zarr Discovery Upgrades (Chunk Metadata, Multiscales, EOPF Patterns)

Pulse highlights three “quick wins” for cloud-native multidimensional data: (1) surface chunk/codec metadata at the STAC Asset level so clients can plan IO without probing Zarr stores; (2) standardize how multiscale pyramids (overviews) are represented and advertised; and (3) copy EOPF scene-product patterns for predictable asset navigation, including the idea of embedding or referencing STAC-like metadata inside the Zarr store (“STAC in Zarr”).

3.1 Proposed KFM capability

- **KFM namespaced array metadata** (suggested: *kfm:array*) attached to STAC assets for Zarr/kerchunk-backed data.
- **Multiscale convention**: a stable way to represent available overviews and future-proof catalogs when pyramids are added later.
- **Scene product asset layout**: consistent roles for “product”, “metadata”, and “measurement-group” assets so clients can browse predictably.

3.2 Minimum viable asset metadata (kfm:array)

Pulse’s recommendation is to keep these fields machine-derivable from Zarr metadata so they can be regenerated and validated deterministically.

Field	Example / purpose
chunk_shape	[1, 256, 256] — chunk grid planning
shape	Full array shape — range planning
dtype	e.g., float32 — decoding
order	C/F (or Zarr v3 equivalent) — layout
codecs / compressor	What decompression stack is needed
dimension_names	["time","y","x"] — semantic access
fill_value / nodata	Masking semantics
storage_hints	e.g., consolidated metadata present; store type

3.3 Multiscale representation decision tree

Pulse frames multiscales as a discoverability problem (not just a preprocessing problem). It identifies three catalog-visible shapes you can standardize around:

- **Single Zarr store** containing multiscales in a hierarchy (zoom levels as child groups).

- **Multiple Zarr stores** (one per resolution) for simpler STAC representation at the cost of more objects to manage.
- **Single store + virtual dataset reference assets** when data lives in legacy chunked formats and STAC assets describe the reference mapping.

For KFM, the safest default is: **store multiscales inside one canonical store** when possible, and mirror multiscale availability into STAC extra_fields for client discovery.

3.4 EOPF scene-product asset layout (copy/paste pattern)

- Always include a **product asset** (whole Zarr hierarchy).
- Always include a **metadata JSON asset** (consolidated metadata).
- Always include **per-measurement-group assets** for clients that only want specific subproducts.
- Add product-type extras only when needed, but keep naming/roles consistent.

3.5 “STAC in Zarr” (self-describing stores)

Pulse notes an emerging convention: embed a quasi-STAC object (or a link to canonical STAC) inside Zarr metadata so the store is self-describing, reducing bespoke SAFE→STAC mapping and making ingestion less brittle. The main operational risk is metadata going stale if stores mutate independently—mitigated by versioning and attestations.

3.6 Validation gates

- **Schema**: kfm:array fields validate; dimension_names match underlying Zarr metadata.
- **Determinism**: emitted metadata must be reproducible from source zarr.json/.zarray; CI should re-derive and compare.
- **Multiscale consistency**: each zoom retains a consistent variable/dimension structure (Pulse treats this as a MUST).

3.7 Buildable baseline (first slice)

- Implement kfm:array emission for Zarr/kerchunk assets during ingestion.
- Add a STAC validator that confirms kfm:array matches the underlying store metadata.
- Publish one “golden” example Collection+Item for a representative cube (with and without multiscales).

4. Module C — Neo4j Graph Migration Enforcement (Governed Change Control)

Pulse reframes Neo4j migrations as governed transformation events—not just schema changes or backfills. The core idea is to make every migration provable, reversible, and ethically constrained via declared intent, explicit invariants, scoped blast radius, provenance capture, and CI-enforced failure semantics.

4.1 Migration taxonomy (drives enforcement rules)

- **Schema-only** (labels, relationship types, constraints, indexes): watch for silent constraint drops and index rebuild latency.
- **Controlled data backfill** (derived nodes, normalization/denormalization): watch for cardinality explosions and duplicate semantics.
- **Ontology evolution** (reclassification, semantic changes): watch for semantic drift that breaks downstream analytics.
- **Governance-sensitive** (Indigenous datasets, treaty/legal authority nodes): treat as highest risk; apply “hard stop” safeguards.

4.2 Baseline capture (beyond simple counts)

Pulse recommends pre-migration baselines that detect graph-shape damage even when node counts stay stable.

- Node and relationship counts (coarse accidental mass edits).
- Schema fingerprint (constraint/index drift).
- Label entropy (ontology reshaping).
- Degree distribution fingerprints (fan-out explosions).
- Protected/authority node hash (“sacred node” mutation detection).

4.3 Invariant classes (what CI enforces)

- **Structural**: no unlabeled nodes; no relationship without type; no orphaned provenance.
- **Ontological**: mutually exclusive labels stay exclusive; required acyclicity constraints hold.
- **Cardinality**: required relationships (e.g., Dataset→Provenance) must exist; enforce min/max fan-outs where defined.
- **Governance (hard stop)**: protected nodes must never change.

4.4 Diffing + blast radius control

- Temporal diffing: stamp mutations with updatedAt and report which labels/actors were touched during the migration window.

- Label-scoped and relationship-scoped diffs to detect out-of-scope changes.
- Fail CI if unexpected labels or relationship types appear outside the declared plan.

4.5 Reversibility guarantees

- Every migration produces rollback artifacts: pre-state checksum, mutation log, deterministic ID map, and tested rollback Cypher.
- Store mutation captures as immutable JSON under a stable path (Pulse suggests data/work/logs/graph/migrations//...).

4.6 PROV alignment

- Represent each migration as a PROV Activity with pre- and post-state Entities and a CI Runner/commit Agent.
- Attach the PROV record to the migration receipt so reviewers can trace what changed and why.

4.7 Migration declaration contract (required)

Pulse proposes shipping a machine-readable declaration with each migration and having CI compare declared intent vs observed deltas.

```
migration_id: GRAPH-MIG-YYYY-MM-DD-##
intent: "Short natural-language statement of what should change"
allowed_labels:
  - Feature
  - Provenance
forbidden_labels:
  - Authority
  - Treaty
max_node_delta: 12000
max_rel_delta: 30000
requires_rollback: true
governance_review: "FAIR+CARE Council"
```

4.8 CI failure semantics

- **Immediate hard fail**: protected node mutation, schema fingerprint mismatch, orphaned provenance.
- **Soft fail (manual review)**: bounded cardinality drift, performance index rebuild.
- **Informational**: new labels introduced and declared.

4.9 Buildable baseline (first slice)

- Implement the migration declaration schema + validator.
- Implement baseline snapshot scripts (counts, schema fingerprint, degree distribution) and store them as receipts.
- Implement CI gating for: scope (labels/rel types), protected-node mutation, and provenance orphan checks.
- Implement mutation logging + a tested rollback script for one representative migration.

5. Module D — Multi-Source Incident / Work-Zone Fusion (WZDx + ArcGIS + GTFS-RT)

Pulse provides an implementation-level blueprint for a near-real-time incident/work-zone fusion pipeline that ingests: WZDx snapshot feeds (road events), ArcGIS Feature Service extractChanges deltas (authoritative edits), and GTFS-Realtime alerts (corroboration and transit disruption signals). The pipeline normalizes observations into a canonical IncidentEvent model, runs QC and deduplication, resolves conflicts using trust tiers, and promotes events through staged confidence levels.

5.1 Canonical IncidentEvent model (publish contract)

Category	Key fields
Identity	event_id; source_event_ids[]; source_types[]; idempotency_key
Spatiotemporal envelope	geometry (canonical, possibly multi-geom); road_ref; start_time/end_time/last_observed_time; time
Semantics	event_type; severity; impacts; status (active/planned/cleared/etc.)
Provenance + trust	sources[] (observed_at, publisher, license, raw_ref); trust_score; promotion_state

5.2 Ingestion patterns per source

- **WZDx**: treat as snapshot unless IDs are stable; derive your own delta by hashing (source_id + geometry + key attributes + timestamps) and comparing to the last snapshot state.
- **ArcGIS extractChanges**: treat as authoritative delta stream where available; persist the server “generation” token and fall back to full reconcile when tracking resets.
- **GTFS-RT**: decode protobuf and enforce timeliness QC (Pulse references ~90 seconds for TripUpdates/VehiclePositions and ~10 minutes for Alerts); compute per-entity age rather than trusting feed-level timestamps.

5.3 Promotion policy (confidence ladder)

- Candidate: observed by any source with minimal confidence.
- Corroborated: confirmed by two independent sources or by a single Tier A source.
- Confirmed: Tier A says active, or two+ sources agree on type+location+time.
- Published: passes QC and has no unresolved high-severity conflicts.

5.4 Conflict resolution rules (examples)

- If Tier A says “cleared” but WZDx remains “active”: prefer Tier A for canonical status, but keep WZDx as a lingering signal (possible feed lag).

- If GTFS-RT alert is active but no road event exists: publish as “Transit disruption” rather than “Road incident” until corroborated.

5.5 End-to-end operational flow

Pulse’s blueprint: **Ingest** → **Normalize** → **Gate** → **Cluster** → **Resolve** → **Promote** → **Publish**.

- QC gates: schema/format checks, timestamp sanity, geometry sanity, staleness thresholds.
- Clustering/dedup: signature match + spatial/temporal merge with thresholds per geometry type.
- Publish streams: incident_events (current state), event_updates (diffs), event_lineage (every decision kept).

5.6 Known failure modes (design for these)

- WZDx feed churn / unstable IDs → mitigate with signature matching and short-term memory.
- ArcGIS change tracking resets or token loss → detect and switch to full reconcile mode, then resume deltas.
- GTFS-RT producer variability → compute per-entity age and enforce staleness thresholds.
- Conflicting “clear” signals across sources → keep both cleared_time and last_observed_time; consider “probable clear” for dropped snapshot entries.

5.7 Required deliverables (artifact pack)

- **Schema bundle**: canonical IncidentEvent JSON Schema; per-source Observation schemas; provenance schema (raw payload references).
- **QC ruleset**: YAML thresholds for geo/time sanity, staleness, allowed event-type mappings.
- **Dedup ruleset**: signature rules, spatial thresholds per geometry type, equivalence classes (e.g., closure vs detour).
- **Promotion policy**: trust tier table, corroboration requirements, conflict resolution matrix.

5.8 Buildable baseline (first slice)

- Implement WZDx snapshot ingestion + schema validation + content-hash delta derivation.
- Implement the canonical IncidentEvent model + one publication sink (table/topic).
- Implement QC YAML + a validator that yields fail/pass plus reason codes.
- Add lineage records for every resolve/promotion decision (minimum viable event_lineage).

6. Module E — Deterministic Data Runs (lakeFS + DVC + LangGraph + OPA/Rego + Cosign)

Pulse includes a complete prototype for a minimal, reproducible end-to-end data run: ingest a JSON feed into a lakeFS branch, run a deterministic DVC transform (optionally orchestrated by a tiny LangGraph agent), gate promotion using OPA/Rego policy and Cosign verification in CI, and perform a single-command rollback if the gate fails.

6.1 Core idea (why this matters)

- Determinism turns pipeline runs into reviewable artifacts: same inputs + same code + same parameters → same outputs.
- Policy gates move governance left: CI blocks promotion unless evidence (schemas + attestations) verifies.
- Rollback is part of the contract, not an afterthought.

6.2 Drop-in repository skeleton (prototype)

```
.
■■ data/
■  ■■ raw/feed.json
■■ dvc.yaml
■■ params.yaml
■■ src/transform.py
■■ src/agent/runner.py
■■ policy/allow_publish.rego
■■ attest/               # generated & verified attestations
■■ .github/workflows/data-run.yml
■■ README.md
```

6.3 Gate: OPA/Rego policy

Pulse's minimal example gate checks record count + Cosign verification outcome. In KFM, treat this as a template that expands into your full policy pack.

```
package kfm.allow
default allow = false

min_records := 1

allow {
  input.result_metadata.records >= min_records
  input.cosign.verified == true
}
```

6.4 Rollback: single command

```
# Revert run branch to last good commit or tag
lakectl revert lakefs://kfm@run-YYYY-MM-DD --to-ref main -m "Gate failed; revert"
# or:
lakectl revert lakefs://kfm@run-YYYY-MM-DD --to-ref v-last-good -m "Gate failed; revert"
```

6.5 CI integration (first slice)

- Run DVC repro in CI; compute result_metadata.json (counts/hashes).
- Verify attestation with Cosign; emit a tiny verdict JSON for policy input.
- OPA eval returns a boolean allow/deny; deny causes workflow failure and triggers rollback.

6.6 Buildable baseline (first slice)

- Implement one deterministic transform stage that produces a small processed artifact.
- Implement the OPA gate with a single rule and a CI job that fails on deny.
- Implement rollback in lakeFS and demonstrate: (a) success promotes, (b) failure reverts.

7. Cross-Module Integration Roadmap (Thin Slices)

To minimize risk and keep governance enforceable, integrate these modules in slices that each (a) produce a promotable artifact, (b) have clear validators, and (c) can be rolled back cleanly.

7.1 Slice 0 — Common governance harness (shared)

- Adopt the deterministic-run harness pattern (Module E) as the default promotion workflow for all new pipelines.
- Define a baseline policy pack (OPA) and a standard run receipt format used across modules.
- Establish a “fail-closed” default: nothing publishes without schema validation + provenance receipt.

7.2 Slice 1 — First promotable dataset in each domain

- **PM■.■ calibration**: one reference-collocated sensor → calibrated hourly series + report + trust score stub.
- **STAC↔Zarr**: one Zarr cube → STAC Collection/Item with kfm:array fields + validator.
- **Graph migrations**: one controlled backfill migration → migration declaration + invariant checks + rollback artifacts.
- **Incident fusion**: one WZDx feed → incident_events table/topic + lineage records + QC YAML.

7.3 Slice 2 — Add confidence & scale mechanics

- PM■.■: add trust-aware adaptive weighting; add drift monitors and regime-based evaluation outputs.
- STAC↔Zarr: add multiscale conventions; add an EOPF-style scene product example.
- Graph migrations: add temporal diff reports and blast radius enforcement (labels/relationships).
- Incident fusion: add ArcGIS extractChanges integration; add promotion/conflict-resolution matrix; add GTFS-RT corroboration.

7.4 Definition of Done (applies to each slice)

- Schemas validate (JSON Schema / equivalent) and are versioned.
- A policy gate exists and is exercised in CI (allow + deny test cases).
- A run receipt/provenance artifact is emitted and stored immutably.
- Rollback procedure is documented and tested (automated where possible).
- A minimal evidence view exists (what changed, what inputs, what outputs, and why it is trusted).

8. Artifact Inventory (What to Create)

This inventory lists the concrete files you would create to integrate the modules into a governed repository. Paths are suggestions.

Artifact	Modules	Suggested path	Notes
JSON Schemas	A, B, D	schemas/{domain}/...	Calibration output schema; STAC asset extra_fields schem
Policy pack (OPA/Rego)	All	policy/packs/{pack}/...	Fail-closed rules for publish; thresholds; license/access co
QC thresholds YAML	A, D	config/qc/{domain}.yaml	Staleness, geometry sanity, drift thresholds, regime bins
STAC examples + validators	B	examples/stac/zarr-cube/...	Golden Collections/Items + validator that re-derives kfm:ar
Graph migration declarations	C	graph/migrations/<id>/migration.yml	Intent + allowed/forbidden labels + blast radius + rollback r
Graph baseline snapshots	C	graph/migrations/<id>/baseline/*.json	Counts, schema fingerprint, degree distributions, protected
Rollback scripts	C, E	ops/rollback/...	Rollback Cypher (C); lakeFS revert commands (E)
Run receipts / PROV records	All (esp. C, E)	receipts/<run_id>/...	Immutable receipts; include PROV JSON(-LD) for migratio
CI workflows	All	.github/workflows/...	Run → validate → attest/verify → policy gate → publish/rol

9. Assumptions, Risks, and Verification Steps

9.1 Assumptions used (because Pulse leaves these as context-dependent)

- KFM will enforce publish/promotion via CI policy gates (OPA) and will store immutable receipts for every run/migration.
- For Module A, KFM can access at least one reference-grade PM■.■ monitor dataset for supervised calibration in the target area/time.
- For Module B, KFM will ingest at least one Zarr or kerchunk-backed multidimensional dataset and publish STAC metadata for it.
- For Module D, KFM will start with snapshot-style WZDx ingestion and add authoritative deltas (ArcGIS extractChanges) if available.

9.2 Key risks surfaced in Pulse (and how to mitigate)

- Metadata drift (STAC fields evolving): mitigate with namespacing and deterministic regeneration from source metadata.
- Producer variability (GTFS-RT timeliness; WZDx ID stability): mitigate with QC gates and signature-based dedup, not raw IDs.
- Graph migrations are high blast-radius: mitigate with declared intent + invariant enforcement + rollback artifacts and hard-stop protected nodes.
- Calibration transferability: mitigate with regime tagging and condition-diverse co-location windows; monitor RH-driven residuals.

9.3 Minimum verification steps before implementation

- Confirm the target sensor hardware family and reporting cadence for Module A; pick the minimal feature set and evaluation metrics.
- Choose one representative Zarr dataset (and optionally one scene-product) to serve as the Module B “golden path” example.
- Define the set of protected node classes/labels for Module C (what is “sacred” in your ontology).
- Identify one initial WZDx feed and one authoritative ArcGIS layer (if available) for Module D; set staleness thresholds.

End of document.

KFM Pulse → Integration Kit

Cultivated, integration-ready concepts distilled from: **pulse 12-29-2025** (governed specs/patterns).

Generated: 2026-02-15

What this document is: a practical “handoff pack” of ideas, artifacts, and thin-slice build steps that can be integrated into KFM’s governed data→catalog→API→UI system.
What this document is not: a replacement for the source specs. It is a synthesis and an integration roadmap.

Scope note: The source pulse contains multiple domain playbooks (air quality, soils & surficial geology, freight/logistics, water telemetry) and cross-cutting platform patterns (zero-trust ingest, typed CloudEvents, STAC alignment, Neo4j ops, CI attestations). This kit extracts the reusable platform contracts and the “golden path” steps common to all of them.

Contents

1 1. Executive summary (what to integrate first)
2 2. Platform invariants (rules you can enforce in CI)
3 3. Building blocks (reusable patterns extracted from the pulse)
4 4. Ready-to-integrate artifact set (what to add to the repo)
5 5. Thin-slice implementation plan (smallest buildable increments)
6 6. Governance, safety, and sensitivity handling
7 Appendix A. Reference directory layouts
8 Appendix B. Checklists (DoD + validation gates)

1. Executive Summary

The pulse 12-29-2025 document reads like a bundle of “drop-in production references” for KFM: domain-specific intake specs plus cross-cutting governance, provenance, and operational patterns. The highest-leverage integration move is to **standardize the cross-cutting contracts** so every new dataset family can plug in without redesigning policy, provenance, cataloging, or CI gating each time.

Integrate first (high leverage):
1) A **Minimal Pipeline Blueprint** contract per dataset family (pipeline.yaml + schemas + qa + stac + dcat + prov + graph).
2) A **Zero-Trust Ingest** pattern for all external data acquisition (ephemeral auth, signed logs, content-addressed staging, license gates, attestations).
3) **Catalog + lineage gates**: STAC/DCAT/PROV must validate and cross-link before promotion.
4) **Watchers**: change detection and freshness monitors that open PRs or trigger pipeline runs.
5) A consistent **graph ingest model** and “derived products” ethics guardrails.

Key integration outcome: KFM can add new sources by implementing connectors and schemas, while governance and provenance remain consistent, fail-closed, and auditable.

2. Platform Invariants (CI-Enforceable Rules)

These invariants are repeated across the pulse's domain specs and patterns. Treat them as platform rules, not advice:

- **Promotion is evidence-producing:** raw artifacts are immutable and content-addressed; processed artifacts are versioned; catalogs and graph ingests are reproducible.
- **License-first:** capture/record license at ingest and block promotion when unknown/forbidden.
- **Determinism:** canonicalize → hash → diff; identical inputs should yield identical outputs (or the pipeline must record why).
- **Provenance by default:** every run emits PROV (entities/activities/agents + run metadata) and (where applicable) SLSA/in-toto attestations.
- **Validation gates:** schema, geometry/raster checks, metadata completeness, and drift thresholds block promotion.
- **Authority boundaries:** do not conflate regulatory truth with real-time signals, models, or satellites; never overwrite a "golden source" with derived fusion.

> **Practical rule of thumb:** if an artifact can be cited by Story Nodes or Focus Mode, it must be reproducible, signed/attested, and resolvable to source inputs and run receipts.

3. Building Blocks Extracted from the Pulse

3.1 Minimal Pipeline Blueprint (Dataset-Family Contract)

Every dataset family should ship with a predictable set of files and folders so ingestion and review are uniform. This keeps engineering effort focused on source-specific connectors rather than reinventing pipeline scaffolding.

- **pipeline.yaml** — inputs, outputs, schedule, resources
- **schemas/** — mapping for normalized tables / assets
- **qa/** — validators and drift thresholds
- **stac/** — templates for Collections/Items/Assets
- **dcat/** — dataset/distribution templates (JSON-LD)
- **prov/** — activity templates + run outputs
- **graph/** — mapping rules (when graph-ingested)

Versioning policy guidance: treat raw artifacts as immutable and content-addressed; processed artifacts as versioned by pipeline + source release; catalogs versioned alongside processed releases; graph ingests reproducible with recorded provenance.

3.2 Data Promotion Saga (Discover → Fetch → Normalize → Validate → Publish → Catalog)

The pulse repeatedly frames ingestion as a deterministic promotion workflow. The point is not ETL convenience; it is to ensure every promoted dataset has audits, checksums, and lineage.

1. Discover (source registry + watcher)
2. Fetch (raw artifact + checksums)
3. Normalize (schema alignment + projections)
4. Validate (QA gates)
5. Publish (STAC/DCAT + immutable artifact)
6. Catalog (graph ingest + search index)

3.3 Zero-Trust Ingest (Sidecar Fetch + Ephemeral Auth + Signed Logs)

For any external data access (web, APIs, file drops, catalogs), the pulse proposes a hardened intake pattern: short-lived credentials, strict network egress, signed request/response logs, and content-addressed staging before transformation.

- **No long-lived secrets:** use per-job ephemeral OIDC tokens.
- **Sidecar isolation:** fetch in minimal, non-privileged container with egress controls.
- **Everything is signed:** logs and artifacts carry signatures/attestations.

- **Content-addressed staging:** blobs keyed by digest; immutability before transform.
- **License-first gate:** block promotion when license is unknown or forbidden.
- **Determinism + provenance:** canonicalize/hash/diff; emit SLSA-style attestations and PROV.

3.4 Watchers (Change Detection + Freshness Monitoring)

Watchers keep datasets current without manual polling. The pulse recommends using cheap HTTP change signals (ETag/If-None-Match, Last-Modified/If-Modified-Since) where available, and defining per-source schedules and triggers.

- **Schedules:** tune by dataset volatility (e.g., nightly incremental updates vs weekly release checks).
- **Freshness triggers:** raise alerts when expected updates lag (e.g., hourly feeds falling behind).
- **PR-first option:** watchers can open pull requests that include new receipts/catalogs for review before merge/promotion.

3.5 STAC/DCAT/PROV as the Publish Gate

Across domains, the pulse uses the same rule: each source class becomes one or more STAC Collections/Items and one DCAT Dataset (with distributions). PROV binds raw inputs, transformations, and published assets together.

DCAT core expectations (example set): title, publisher, license, spatial footprint, distributions for raw/processed/catalog endpoints, and a provenance link to PROV.

PROV minimum expectations: Entity (source + produced artifacts), Activity (fetch/extract/normalize/validate/publish), Agent (runner/CI identity/human approver), plus run_id/commit/deterministic seed fields.

3.6 Authority Classes (Keep Signals Separate)

A concrete example is the pulse's "air quality classes": it explicitly separates regulatory ground truth, real-time alerts, harmonized bridges, forecast models, mesoscale models, and satellites. This is a general pattern: keep authority clear so derived products don't masquerade as truth.

Class	Purpose	Typical latency	Authority stance
Regulatory ground	Compliance & trend baseline	High	Federal / State
Real-time alerts	Public health signal	Low	Federal
Harmonized global	Cross-network bridge	Medium	NGO / Research
Chemical forecast	Prediction & context	Medium	ECMWF
Mesoscale model	High-resolution dynamics	Low	NOAA
Satellite	Spatial completeness	Medium	ESA

4. Ready-to-Integrate Artifact Set

Below is a concrete “repo add” list. Each item is framed so it can be implemented independently and reviewed via PR.

Artifact	Purpose	Suggested location	Definition of Done (DoD)
Dataset family scaffold	Standardize per-source layout (pipeline/schemas/qa/catalogs/lineage)	`data/<domain>/<dataset>/`	Folders exist; templates compile; sample dataset passes validation
Pipeline config contract	Make pipeline.yaml a typed, validated interface	`src/pipelines/<name>/pipeline.yaml + schema`	Schema-validated; includes schedule, inputs/outputs, resources
Zero-trust ingest action(s)	Reusable GitHub Action(s) / scripts for sidecar fetch + OIDC + signed logs	`.github/actions/{oidc-token,http-sign-log,cosign-attest}`	Can fetch+sign+attest a fixture source end-to-end
Content-addressed staging	Immutable by-digest storage for raw and intermediate blobs	`data/_staging/by-digest/sha256/…`	Digest rules enforced; transforms never mutate inputs
License gate policy pack	Fail-closed license enforcement across all pipelines	`policy/license/*.rego + allowlist`	Blocks unknown; tests cover allowed/denied examples
Catalog validators	STAC/DCAT/PROV schema validation and cross-link checks	`tools/validation/catalog/`	CI fails on invalid JSON-LD/JSON; referential integrity tested
PROV run receipt template	Uniform provenance fields for every run	`prov/templates/run.jsonld`	Includes Entities/Activities/Agents + run metadata; generated per run
Watcher registry	Central inventory of update checks + schedules	`docs/ops/watchers/registry.yml`	At least one watcher implemented; produces a PR or pipeline dispatch
Graph mapping rules	Normalize label/edge conventions by domain	`graph/mappings/<domain>/*.yml`	Maps produced; graph invariants tests run
Ethics + sensitivity checklist	Enforce non-sensationalization & redaction defaults for risky narratives	`docs/governance/safety_checks.md`	Required for promotion of sensitive layers

5. Thin-Slice Implementation Plan (Buildable Increments)

A practical way to integrate the pulse ideas is to deliver them as a sequence of thin slices. Each slice should produce a verifiable artifact (tests passing, catalogs generated, provenance emitted) that can be merged safely.

Slice 0 — Contract-first scaffolding

- Add the dataset-family scaffold template + pipeline.yaml schema.
- Add catalog validators (STAC/DCAT/PROV) to CI as a no-op gate (runs on fixtures).
- Add a minimal PROV run receipt template and a “receipt viewer” JSON output for developers.

Slice 1 — Zero-trust ingest foundation

- Implement sidecar fetch + OIDC token minting + signed fetch-log.json.
- Implement content-addressed staging paths keyed by sha256(blob).
- Implement license capture + allow/deny policy (fail closed).
- Emit in-toto/SLSA attestations for fetch and transform steps.

Slice 2 — One end-to-end dataset (golden path)

- Pick a single source with clear authority boundaries (e.g., a real-time feed and a regulatory archive).
- Implement connector + normalization schemas + QA checks.
- Publish STAC Items/Collections + DCAT Dataset/Distributions + PROV run graph.
- Wire a watcher to detect updates and open a PR or dispatch a pipeline run.

Slice 3 — Graph + derived products

- Ingest normalized outputs into a graph model with explicit labels/edges.
- Add a single derived product that is easy to validate (e.g., counts, density bins) and tag it as decision-support.
- Add mandatory “Data Limitations” Story Node template and ensure derived surfaces reference it.

6. Governance, Safety, and Sensitivity Handling

The pulse contains explicit safety-oriented guidance that should become enforceable policy or checklist items—especially when handling hazard incidents, carrier safety indicators, and any location-sensitive records.

- **Avoid definitive judgments:** safety indicators must include methodology context and be framed as signals, not verdicts.
- **Non-sensational incident narratives:** present aggregated views by default; avoid sensationalizing hazards.
- **Sensitivity flags:** apply sensitivity flags even for public data when inadvertent harm is plausible.
- **Promotion chain:** all publish steps should record governance references and a promotion chain for auditability.

> **Implementation suggestion:** treat governance checks as data transformations. For high-risk datasets, add a “redaction/aggregation required” gate that must pass before publication.

Appendix A — Reference Directory Layouts (From the Pulse Patterns)

A) Dataset family standard layout (example: freight/logistics)

```
data/logistics/ freight_flows/ raw/ processed/ stac/ dcat/ prov/ qa/ networks/ highway/
raw/ processed/ graph/ stac/ dcat/ prov/ qa/ rail/ raw/ processed/ graph/ stac/ dcat/ prov/
qa/ carriers/ raw/ processed/ stac/ dcat/ prov/ qa/ hazmat/ raw/ processed/ stac/ dcat/
prov/ qa/ derived/ risk/ performance/ story_nodes/ stac/ dcat/ prov/ qa/
```

B) Zero-trust ingest repo layout (high-level)

```
data/_staging/by-digest/sha256/… (blob + logs + signatures + attestations)
docs/patterns/intake/zero-trust-ingest-blueprint.md
src/connectors/{arcgis_rest,arcgis_hub,s3_fetch,oci_fetch,html_capture}/
tools/{validation,canonicalize}/ .github/workflows/ingest_*.yml
.github/actions/{oidc-token,http-sign-log,cosign-attest}/
```

Appendix B — Checklists

B.1 Definition of Done (per ingest)

- Ephemeral OIDC used; no static secrets touched.
- Raw blob stored under content-addressed path with sha256 metadata.
- Signed fetch-log verifiable; request metadata sanitized (no secrets).
- License captured and allowed; unknown licenses quarantined.
- Attestations exist for fetch and transform stages.
- STAC/DCAT/PROV entries created, schema-valid, and cross-linked.
- QA report generated; drift thresholds evaluated; promotion blocked on failures.

B.2 Validation Gates (minimum set)

- **Data:** schema conformity; geometry validity; raster stats/ranges; null/uniqueness rules; field domains.
- **Metadata:** STAC JSON schema validates; DCAT JSON-LD shape validates; PROV JSON-LD validates; license present and consistent.
- **Promotion:** sign → attest → publish → catalog; require checksums, run logs, validation reports, and (when triggered) human approval.

KFM Integration Idea Pack (Ready-to-Integrate)

Prepared: 2026-02-15 (UTC) **Version:** 0.1 (draft)
Source: New Ideas Feb-2026-1.pdf (cultivated + normalized into integration-ready artifacts)

This document converts the patterns, contracts, and playbooks in the source material into a **drop-in integration package** for KFM: concrete repo additions (schemas, policies, CI gates, UI components), plus a phased adoption plan and Definition of Done.

Non-negotiables (enforced)
Fail-closed promotion: if evidence, signatures, or policies fail → nothing ships.
Deterministic identity: *spec_hash* computed from canonicalized inputs.
Receipts as first-class artifacts: every run emits a typed run_receipt/run_manifest.
Attach provenance to content: OCI referrers + attestations travel with datasets.

How to use this pack

Treat each section as a PR-sized unit. Add the contracts first (schemas), then enforce them (policies), then make them visible (UI evidence surfaces). Where the source document provides multiple options, this pack selects a default and labels alternatives explicitly.

Quick Table of Contents

- 1. Integration targets and repo drop-ins
- 2. Canonical contracts (schemas) - run_receipt, run_manifest, watchers registry
- 3. Policy packs (OPA/Rego + Conftest) - fail-closed gates
- 4. CI/CD workflows - verify, attest, open PRs, block merges
- 5. OCI distribution playbook - datasets + referrers (ORAS + Cosign)
- 6. Deterministic packaging playbooks - STAC 1.1.0, COG, Zarr
- 7. UI components - ReceiptViewer and “trust badges”
- 8. Domain playbooks - GTFS, station telemetry, soils diffing
- 9. Risks, tradeoffs, and verification steps
- 10. Phased adoption plan + Definition of Done

1. Integration Targets and Repo Drop-ins

The source material repeatedly converges on a single “governance-by-construction” loop: **watch** → **compute deterministic hashes** → **emit receipts** → **validate + attest** → **fail-closed gates** → **publish immutable artifacts** → **surface provenance in UI**. This section normalizes that loop into a concrete set of repository additions and their responsibilities.

1.1 Drop-in directory layout (proposed default)

```
repo-root/
|-- schemas/
|   |-- run_receipt.v1.schema.json
|   |-- run_manifest.v1.schema.json
|   `-- watcher.v1.schema.json
|-- policy/
|   `-- opa/
|       |-- receipt.rego
|       |-- receipts_pr_gate.rego
|       |-- watchers.rego
|       `-- gates/
|           |-- station.rego
|           `-- artifact.rego
|-- ops/
|   |-- templates/
|   |   `-- run_receipt.tmpl.json
|   `-- playbooks/
|       `-- oci_referrers_playbook.(md|pdf)
|-- src/
|   |-- pipelines/
|   |   |-- fetchers/
|   |   `-- validators/
|   `-- models/
|       `-- run_receipt.py
|-- web/
|   `-- src/components/ReceiptViewer/
|       |-- ReceiptViewer.tsx
|       |-- ReceiptViewer.css
|       `-- README.md
`-- .github/workflows/
    |-- kfm__provenance-policy-gate.yml
    |-- pr-verify-receipts.yml
    |-- kfm__watchers.yml
    `-- kfm__ops-acceptance.yml
```

This layout mirrors the source doc’s repo wiring suggestions for watchers and promotion gates, including schemas, policy modules, and GitHub Actions workflows.

1.2 What you get (integration artifacts)

Artifact	Type	Purpose	Primary gate
schemas/run_receipt.v1.schema.json	JSON Schema	Contract for per-run receipts (inputs, outputs, checks, timestamps).	Conftest schema + policy
schemas/watcher.v1.schema.json	JSON Schema	Contract for Watchers Registry entries (signed allow-list).	CI validate + signature verify
policy/opa/receipt.rego	Rego policy	Deny promotion if required receipt fields missing or provider unapproved.	Conftest deny-by-default
.github/workflows/kfm__provenance-policy-gate.yml	CI workflow	Run watcher → emit receipt → policy gate → update catalogs → draft PR.	Required status checks
web/src/components/ReceiptViewer	UI component	Safe, read-only receipt viewer: schema validate → signature verify → render.	Unit tests + policy assertions
ops/playbooks/oci_referrers_playbook	Ops playbook	Publish datasets as OCI artifacts and attach PROV/SBOM/receipts as referrers.	Promotion lane policy

Source anchors: run_receipt + ORAS attach + Cosign attest examples (pp. 55-57); Rego policy for receipts (pp. 208-210); watchers registry contract + repo wiring (pp. 221-226); ReceiptViewer design + guardrails (pp. 183-187).

2. Canonical Contracts (Schemas)

Contracts are the foundation of fail-closed governance: if data does not match the schema, it cannot be trusted, signed, promoted, or shown with “green” affordances. This pack defines three contracts:

- **Run Receipt** — one per ingestion/transform run (typed, hashable).
- **Run Manifest** — promotion-oriented rollup (digests, Rekor UUID, rights).
- **Watchers Registry** — signed allow-list that enables automation safely.

2.1 Run Receipt (v1) — minimal required fields

The source doc includes a minimal run_receipt JSON shape with the key idea that **spec_hash** is computed from canonicalized inputs and stored alongside checks and timestamps (pp. 55-56).

```
{
  "$schema": "https://kfm.dev/schemas/run_receipt_v1.json",
  "kfm_run_id": "string",
  "dataset": "string",
  "subject": "sha256:<digest>",
  "spec_hash": "sha256:<64-hex>",
  "pipeline": "raw→processed→catalog→prov",
  "runner": "gha://workflows/provenance-guard@v3",
  "inputs": ["uri://..."],
  "outputs": ["stac items", "prov bundle"],
  "checks": { "stac": "ok|fail", "prov": "ok|fail", "policy": "ok|fail" },
  "timestamps": { "start": "RFC3339", "end": "RFC3339" }
}
```

Required invariants

- spec_hash MUST be computed from canonicalized content (JCS/RFC 8785-style canonicalization is used in examples).
- subject MUST be a digest (avoid tags) to make verification unambiguous.
- checks MUST be present and interpreted fail-closed (any “fail” blocks promotion).
- timestamps MUST be RFC3339 and must bound the run (start <= end).

2.2 Run Manifest (v1) — promotion-focused rollup

The source doc distinguishes a per-run receipt from a promotion-focused run_manifest and shows PR gates that require invariants like rekor_uuid, spec_hash, and rights (pp. 99-102).

```
{
  "dataset": "soils.ssurgo",
  "run_id": "2026-02-02T12:00Z-abc123",
  "spec_hash": "sha256:...",
  "input_sha256": "sha256:...",
  "rekor_uuid": "c7f2...e1",
  "rights": { "license": "...", "attribution": "...", "redistribution": "..." },
  "produced_artifacts": ["oci://...@sha256:<digest>"],
  "materiality_delta": 0.03,
  "publish_candidate": true,
  "signatures": ["..."] ,
  "attestations": ["oci://..."]
}
```

2.3 Watchers Registry (v1) — signed allow-list

A Watchers Registry is a typed list of automated fetchers with strict fields so CI and Focus Mode can make deterministic allow/deny decisions (pp. 221-226).

```
{
  "$id": "https://kfm.org/schemas/watcher.v1.json",
  "$schema": "https://json-schema.org/draft/2020-12/schema",
  "title": "KFM Watcher",
  "type": "object",
  "additionalProperties": false,
  "required": ["watcher_id","canonical_id","endpoint","poll","policy","outputs","schema_url","versi
  "properties": {
    "watcher_id": { "type": "string", "pattern": "^[a-z0-9:_-]+$" },
    "endpoint": { "type": "string", "format": "uri" },
    "poll": { "type":"object", "required":["interval_seconds"], "properties": {
      "interval_seconds": { "type":"integer", "minimum": 60 },
      "mode": { "type":"string", "enum":["poll","webhook","hybrid"], "default":"poll" }
    }},
    "policy": { "type":"object", "properties": {
      "staleness_s": { "type":"integer", "minimum": 60 },
      "spec_change_pct": { "type":"number", "minimum": 0, "maximum": 1 },
      "geom_shift_m": { "type":"number", "minimum": 0 }
    }},
    "signature_ref": { "type":"string" }
  }
}
```

Why signatures are part of the schema

CI and Focus Mode must verify signatures and spec_hash before enabling watchers. This prevents silent automation drift and makes the allow-list portable.

Source anchors: run_receipt schema + ORAS referrers attach (pp. 55-57); PR verify gate requiring rekor_uuid/spec_hash/rights (pp. 99-102); Watchers Registry schema and wiring (pp. 221-226).

3. Policy Packs (OPA/Rego + Conftest)

Policies operationalize governance. The repeated pattern is deny-by-default, and only allow promotion when all required evidence is present and verified.

3.1 Receipt policy (deny missing spec_hash; approved providers)

```
package kfm.receipt
default allow = false

approved_providers := {"http", "usgs", "noaa", "kansas-data-portal"}

missing(field) { not input[field] }

deny[msg] {
  missing("spec_hash")
  msg := "spec_hash is required"
}

deny[msg] {
  not input.provider
  msg := "provider must be set"
}

deny[msg] {
  input.provider != ""
  not input.provider == approved_providers[_]
  msg := sprintf("unapproved provider: %v", [input.provider])
}

allow { count(deny) == 0 }
```

Run in CI:

```
conftest test run_receipt.json -p policy/opa
```

3.2 Ops acceptance gate (station telemetry + artifact integrity)

A compact acceptance gate blocks promotion when station health degrades or remote artifacts change materially (pp. 227-230).

```
# policy/opa/gates/station.rego
package kfm.gates

deny[msg] {
  input.station.missing_pct > 0.25
  msg := "station_missingness_exceeds_25%"
}

deny[msg] {
  input.station.longest_gap_seconds > 2 * input.station.sample_interval_seconds
  msg := "station_gap_gt_2x_sample_interval"
}

deny[msg] {
  input.station.uptime_24h < 0.90
  msg := "station_uptime_below_90_percent_24h"
}

# policy/opa/gates/artifact.rego
package kfm.gates
```

```
deny[msg] {
  some k
  input.artifact.current[k] != input.artifact.recorded[k]
  k == "etag"  # or "last_modified" or "sha256" or "content_md5"
  msg := sprintf("artifact_%s_changed", [k])
}
```

3.3 PR gate: require run_manifest + attestations + kill-switch

A merge-blocking PR gate fails closed if kill-switch is active, if run_manifest is missing, if cosign verification fails, or if critical fields are absent (pp. 100-102).

```
# policy/opa/receipts_pr_gate.rego
package kfm.receipts

default allow = false

allow {
  rm := input.run_manifest
  rm.signatures != null
  rm.attestations != null
  rm.rights != null
  rm.rekor_uuid != null
  rm.spec_hash != null
}
```

Policy development notes

- Keep policies small and composable: one file per concern.
- Add _test.rego unit tests (golden pass/fail fixtures) to prevent silent policy drift.
- Version your policy pack and reference it from receipts/manifests (so audits can reconstruct the rule set used).

Source anchors: receipt.rego and Conftest invocation (pp. 208-210); ops acceptance gate policies (pp. 227-230); PR verify receipts policy (pp. 100-102).

4. CI/CD Workflows (Fail-Closed by Default)

The goal is automated evidence production: CI should answer what changed, why it changed, and whether it is allowed to ship.

4.1 Reference workflow: provenance + policy gate → draft PR

```
Event (schedule/webhook/push)
  → Run watcher/fetcher (ETag/Last-Modified aware)
  → Canonicalize inputs + compute spec_hash
  → Emit typed run_receipt/run_manifest
  → Policy checks (Conftest/OPA) + schema validation (JSON Schema, STAC/DCAT/PROV)
  → (Optional) Package outputs as OCI artifact by digest
  → Attach receipts/provenance as OCI referrers + create cosign attestation
  → Open a draft PR that embeds receipts + digests
  → Branch protection: merge allowed only if all checks stay green
```

4.2 OCI referrers attach + cosign attestation (digest-addressed)

Attach run_receipt and PROV bundle to the dataset digest using OCI referrers (ORAS attach) and create an in-toto attestation with cosign (pp. 55-56).

```
# Attach artifacts to the subject digest (NOT a tag)
oras attach   --artifact-type application/vnd.kfm.run-receipt.v1   --annotation "kfm.spec_hash=$spe

cosign attest   --predicate run_receipt.json   --type https://in-toto.io/Statement/v1   registry.ex

cosign verify-attestation   --type https://in-toto.io/Statement/v1   registry.example.com/kfm/my-da
```

4.3 Merge blocking: kill-switch + invariant checks

A repo kill-switch can force checks to fail; invariant checks require receipts, attestations, and rights before merge (pp. 3-6, 100-102).

```
# Fail closed if kill-switch enabled
if [ -f .github/KILL_SWITCH ] || [ "${DEPLOY_KILL_SWITCH}" = "1" ]; then
  echo "KILL_SWITCH_ACTIVE"
  exit 1
fi
```

4.4 Validation harness (metadata + provenance)

```
# STAC
stac-validator data/catalog/stac/<dataset>/collection.json

# PROV (example; depends on your PROV tooling)
provconvert -infile data/catalog/prov/<dataset>.provn -validator

# Policy checks
conftest test run_receipt.json --policy policy/opa
```

Source anchors: ORAS attach + cosign attest/verify (pp. 55-56); watcher/aggregator/kill-switch/rollback flow (pp. 3-6); promotion lane rules (pp. 99-101).

5. OCI Distribution Playbook (Datasets + Referrers)

Publish datasets as immutable OCI artifacts so provenance, SBOMs, and attestations live beside the content as discoverable referrers (pp. 164-165).

5.1 Quick mental model

- **Subject**: the dataset blob pushed to a registry and addressed by digest.
- **Referrers**: PROV bundles, run receipts, SBOMs, QA logs, signatures-linked to the subject digest.
- **Discovery**: `oras discover` or the registry's referrers API lists attached artifacts.

5.2 Pre-flight checklist (once per registry)

- Registry supports OCI Distribution v1.1 referrers and has retention configured.
- Cosign key strategy chosen and verification roots published.
- Policy definition: fail-closed means artifacts without valid attestations are not promoted.

5.3 Reference commands (push → capture digest → attach referrers)

```
# Push dataset as OCI artifact (example)
oras push ghcr.io/yourorg/kfm/datasets/kansas/soil:latest   --artifact-type application/vnd.kfm.dat

# Capture digest (pin @sha256:<digest> in catalogs)
DIGEST=$(oras manifest fetch --descriptor ghcr.io/yourorg/kfm/datasets/kansas/soil:latest | jq -r .

# Attach provenance / receipts (referrers)
cosign attest --predicate provenance/prov.json --type slsaprovenance   ghcr.io/yourorg/kfm/datasets

oras attach --artifact-type application/vnd.kfm.run-receipt.v1   --annotation "kfm.spec_hash=${SPEC
```

5.4 Consumption rule (critical)

Downstream systems must pin digests, not tags: **@sha256:**.

Source anchors: OCI referrers playbook framing + steps (pp. 164-165); ORAS attach (pp. 55-57).

6. Deterministic Packaging Playbooks (STAC 1.1.0 + COG + Zarr)

Deterministic outputs are required for trustworthy hashing and reproducible diffs. STAC 1.1.0 is treated as baseline; concrete build commands for COG and Zarr are provided (pp. 33-34).

6.1 Determinism checklist (must enforce)

- Pin tool versions (GDAL/PROJ, xarray, zarr, numcodecs, PySTAC).
- Fix tiling/chunks and compressor settings (no “auto” chunking).
- Canonicalize STAC JSON (sorted keys) before hashing/signing.
- Avoid post-writes to COGs; updating in place can break layout.
- Normalize numeric precision and geometry before hashing (snap-to-grid, force2D, stable rounding).

6.2 Build recipes (excerpt)

```
# Zarr (deterministic chunks + Zstd)
python - <<'PY'
import xarray as xr
from numcodecs import Blosc
ds = xr.open_rasterio('input.tif')
compressor = Blosc(cname='zstd', clevel=5, shuffle=Blosc.SHUFFLE)
encoding = {v: {'compressor': compressor, 'chunks': (1,512,512), 'dtype': str(ds[v].dtype)} for v i
ds.to_zarr('out.zarr', mode='w', encoding=encoding, consolidated=True)
PY

# COG (lock tiling + compression + overviews)
gdal_translate -of COG   -co COMPRESS=ZSTD -co ZLEVEL=9   -co BLOCKXSIZE=512 -co BLOCKYSIZE=512   -
gdaladdo -r average output_cog.tif 2 4 8 16 32
```

6.3 STAC metadata locking

Use STAC 1.1.0 core plus Raster/Projection/Datacube extensions as applicable; set stac_version explicitly; hash only canonicalized JSON.

Source anchors: deterministic packaging guidance + build commands (pp. 33-34).

7. UI Components (Evidence Surfaces)

ReceiptViewer is a compact, safe, read-only UI component for receipts. Trust chain: validate schema → verify signature → render; if verification fails, fail closed (pp. 183-187).

7.1 ReceiptViewer: minimal API and behavior

```
type ReceiptViewerProps = {
  schema: object;        // JSON Schema (AJV-compatible)
  data: unknown;         // raw receipt/manifest
  onAction?: (type: string, payload?: unknown) => void;
  compact?: boolean;     // for PR previews
};
```

Render blocks (recommended)

- Status pill: Valid/Invalid · Verified/Unverified
- Inputs table: key/value with type-aware renderers
- Signatures panel: thumbprint, signer, issued/expires, verify result
- Links panel: safe external links only (noopener + noreferrer)
- JSON viewer: lightweight, collapsed by default

7.2 Security & UX guardrails (non-negotiable)

- Validate first (AJV) before computing derived views; do not trust unvalidated inputs.
- Verify signatures (WebCrypto / JOSE / TweetNaCl). If verification fails: fail-closed.
- No dangerouslySetInnerHTML; if trusted HTML is unavoidable, sanitize (e.g., DOMPurify).
- External links: always target="_blank" and rel="noopener noreferrer".
- Accessibility: headings/landmarks, table semantics, focus order, sufficient contrast.

7.3 Suggested repo placement + tests

Place under **web/src/components/ReceiptViewer/** with a small README and Storybook story. Add CI gates for schema compilation and signature verification unit tests.

Source anchors: ReceiptViewer design and guardrails (pp. 183-187).

8. Domain Playbooks (Ready-to-Run Patterns)

Domain recipes share invariants: canonicalize inputs, compute stable identifiers/hashes, classify materiality, and emit receipts so governance can block/allow promotion.

8.1 GTFS (Schedule + Realtime): canonicalization and stable IDs

Make schedule CSVs deterministic and treat material events differently from noise (pp. 155-156).

- Canonicalize schedule CSVs: deterministic sort (e.g., stops by stop_id; stop_times by (trip_id, stop_sequence)).
- Normalize time/tz: keep GTFS HH:MM:SS semantics (including >24h wrap) but normalize generated timestamps to UTC.
- Trim whitespace and represent nulls consistently.
- Use content-addressed IDs for trips and alerts by hashing canonical rows/payloads.
- For frequencies.txt, diff headway statistics by (route_id, direction_id, time_window) rather than enumerating trips.

8.2 Station telemetry: operational acceptance gate

Use station health policies to block PR creation or promotion when telemetry quality degrades (pp. 227-230).

8.3 Soils (SSURGO/gNATSGO): spec_hash + geom_hash diffing

Compute two hashes (spec_hash + geom_hash) to classify attribute vs geometry changes (pp. 141-142).

```
SELECT
  m.mukey,
  m.muname,
  a.comppct_r,
  md5(CONCAT_WS('|',
    m.mukey::text,
    trim(m.muname),
    COALESCE(ROUND(a.comppct_r::numeric, 6)::text,'')
  )) AS spec_hash,
  md5(ST_AsBinary(ST_Force2D(ST_SnapToGrid(m.geom, 0.0000001)))) AS geom_hash,
  m.geom
FROM mapunit AS m
JOIN muaggatt AS a ON m.mukey = a.mukey;
```

8.4 Materiality in practice

Materiality is a typed decision (publish_candidate: true/false) derived from domain-appropriate diffs.

Source anchors: GTFS playbook (pp. 155-156); ops acceptance gate (pp. 227-230); soils recipe (pp. 141-142).

9. Risks, Tradeoffs, and Verification Steps

The main risks are operational: toolchain drift, key management, noisy diffs, and sensitive-data leakage.

9.1 Risk register (starter)

Risk	Why it matters	Mitigation (minimum)
Toolchain drift (cosign/conftest/STAC tools)	Breaking changes or security issues can invalidate CI gates.	Pin versions; add policy/unit regression suite; periodic gov
Signature verification UX confusion	Users may interpret ‘green’ as trust without understanding scope.	Fail-closed; expose ‘why’ in UI; link to evidence resolver an
False positives in materiality	Too many PRs or blocked releases due to noisy diffs.	Domain-specific normalization; thresholds in policy; staged
Sensitive data leakage through receipts/logs	Receipts/logs include URLs/fields that should not be public.	Classify receipts; redact PII; log structural events only; poli

9.2 Verification steps (minimum work to confirm unknowns)

- Pick one pilot dataset and implement run_receipt + schema validation end-to-end.
- Add Conftest policy gate and make it a required status check on main.
- Publish one OCI artifact by digest and attach one receipt + one PROV bundle as referrers; verify discovery.
- Add ReceiptViewer rendering of that receipt in the UI and ensure invalid/unsigned receipts render as untrusted.
- Measure: PR volume, gate pass rate, time-to-draft-PR, and rollback frequency; tune materiality thresholds.

10. Phased Adoption Plan + Definition of Done

Adopt this in thin slices. Each phase is reviewable, reversible, and merge-safe.

Phase 0 — Contracts first (schemas)

- Add schemas: run_receipt, run_manifest, watcher registry.
- Add fixture examples (valid + invalid) and validate in CI.
- Document fields and invariants in a short README next to schemas.

DoD: CI fails on invalid examples; schema compilation is a required status check.

Phase 1 — Fail-closed policy gates

- Add OPA policies for receipt invariants + PR gate.
- Wire Conftest in PR CI; deny-by-default with readable failure output.
- Add kill-switch support (repo file and/or secret) that forces checks to fail.

DoD: Branch protection blocks merge if any policy denies; kill-switch blocks merge in <1 minute.

Phase 2 — Evidence distribution (OCI referrers)

- Publish one dataset as an OCI artifact by digest (not tag).
- Attach run_receipt + PROV bundle as referrers; add cosign attestation.
- Add a ‘discover attached evidence’ script for auditors and UI backends.

DoD: Fresh environment can verify attestations and discover referrers without prior state.

Phase 3 — Evidence UX (ReceiptViewer + trust badges)

- Add ReceiptViewer component and wire it to the evidence resolver payloads.
- Fail-closed UI states: invalid/unsigned receipts are explicitly untrusted.
- Telemetry: log structural events only (no PII) and link to PROV run IDs.

DoD: Users can click dataset → receipt → evidence and understand why trust is granted/denied.

Phase 4 — Scale out (watchers + aggregator)

- Introduce signed Watchers Registry as the single allow-list for automation.
- Add aggregator rollup check (Checks API) that summarizes all watcher receipts + materiality deltas.
- Add rollback action that reverts last publish and emits a rollback attestation.

DoD: Multi-watcher system produces one rollup and can rollback with attestable evidence.

End of pack. If you want, I can also generate a matching repository scaffold (YAML + schemas + policy bundle + ReceiptViewer skeleton) as commit-ready files.

KFM Integration Kit

Cultivated, integration-ready modules derived from *KFM New Ideas 2-8-26* (Feb 2026).

Version: 0.1 (generated 2026-02-15)

Purpose. Convert the source document's concepts into drop-in modules—contracts, policy gates, metadata profiles, and thin-slice build steps—so engineering teams can implement with minimal ambiguity and strong governance.

Governance posture. Evidence-first, fail-closed. Frontend and external clients never access databases directly; all access is via the governed API + policy boundary. Sensitive fields must be generalized/redacted and surfaced with CARE flags.

Primary source: KFM New Ideas 2-8-26 (internal PDF, 234 pages). Page references in this kit point back into that document.

Table of Contents

Placeholder for table of contents **0**

1. Executive Summary (What to Implement First)

The source document repeatedly converges on a small set of implementation primitives: deterministic identity (**spec_hash**), typed provenance (**prov:run_receipt**), fail-closed policy gates (OPA/Rego + Conftest), standards-first cataloging (STAC/DCAT/PROV linkability), and user-facing traceability (provenance timeline + one-click citations).

This kit organizes those primitives into integration modules. Each module includes: required artifacts, interfaces, CI gates, and a definition of done (DoD).

Thin-slice orde	Module	Why it's first
1	Receipt + Catalog Contract	Everything else depends on stable IDs, digests, and linkable metadata
2	Policy Gate Pack (fail-closed)	Blocks unsafe promotion; encodes licensing/CARE rules into CI.
3	Provenance Timeline (Graph + API + UI)	Turns provenance into a visible product feature.
4	Focus Mode Citation Handshake	Enforces "cite or abstain" and surfaces CARE redactions.
5	Story Node + Evidence Hooks	Demonstrates end-to-end: ETL → catalogs → graph → narrative.

2. Platform Contracts and Drop-in Directory Layout

A recurring pattern is to make governance and provenance physical: predictable folders, machine-validated schemas, and immutable linking between datasets, run receipts, and catalogs.

Recommended directory skeleton (example from the source doc):

```
data/
  raw/
    ///...
  processed/
    /.
  catalog/
    dcat/.json
    stac/collections//collection.json
  prov/
    run_receipts//.json
    datasets//bundle.json
```

Key linkage: each raw/processed artifact must point to exactly one run_receipt; each run_receipt must be discoverable from DCAT and/or STAC.

Source: examples show the structure and PROV bundle linkages. (KFM New Ideas 2-8-26 p.166–167)

3. Module: Receipt + Catalog Contract (spec_hash + prov:run_receipt)

The run receipt is the smallest “evidence object” that makes datasets auditable. The source doc shows a receipt containing HTTP freshness hints (ETag/Last-Modified), deterministic spec_hash, artifact digests, and optional cosign references.

Field	Purpose	Notes
@type	Typed receipt marker	Use prov:run_receipt
spec_hash	Deterministic identity for the run spec	sha256 of canonical JSON (jq -S) or manifest
fetched_at	Clock time of fetch	UTC ISO-8601
accessURL	Source URL	Reject insecure http:// in policy
etag / last_modified	Conditional request keys	Enable cheap “no-change” runs
artifact_digest	sha256 digest of fetched payload	Binds receipt to bytes
fallback_used	True if used cached artifact	Fail-closed if no cache and no payload
cosign.signature_ref / attestation_ref	Supply chain attachment pointers	Optional but recommended

Source: receipt JSON example and storage guidance. (KFM New Ideas 2-8-26 p.166, p.208–209)

Minimal JSON Schema (starter):

```
{
  "$schema": "https://json-schema.org/draft/2020-12/schema",
  "title": "prov:run_receipt (KFM minimal)",
  "type": "object",
  "required": ["@type","spec_hash","fetched_at","accessURL","artifact_digest"],
  "properties": {
    "@type": {"const":"prov:run_receipt"},
    "spec_hash": {"type":"string","pattern":"^sha256[:\-][0-9a-fA-F]{8,}$"},
    "fetched_at": {"type":"string","format":"date-time"},
    "accessURL": {"type":"string","format":"uri"},
    "etag": {"type":"string"},
    "last_modified": {"type":"string"},
    "artifact_digest": {"type":"string","pattern":"^sha256[:\-][0-9a-fA-F]{8,}$"},
    "fallback_used": {"type":"boolean"},
    "cosign": {
      "type":"object",
      "properties":{
        "signature_ref":{"type":"string"},
        "attestation_ref":{"type":"string"}
      }
    }
  }
}
```

DoD:

■ Acceptance criteria
A receipt is emitted for every ingestion run and stored under data/prov/run_receipts/<date>/…
DCAT and/or STAC entries link to the receipt (direct URL or bundle pointer).
Policy pack rejects missing artifact_digest, insecure accessURL, or missing license/CARE fields.

4. Module: Policy Gate Pack (OPA/Rego + Conftest) — “Fail Closed”

Policy gates are used at multiple layers: ingestion scripts run Conftest locally; CI blocks merges unless catalogs validate and receipts/rights rules pass; runtime APIs should remain deny-by-default.

A receipt gate pattern appears in the watcher script: run Conftest over the receipt and abort on failure. (KFM New Ideas 2-8-26 p.208–210)

Starter Rego rules (example):

```
package kfm.guard

deny[msg] {
  input["@type"] != "prov:run_receipt"
  msg := "receipt must be prov:run_receipt"
}

deny[msg] {
  not input.artifact_digest
  msg := "artifact_digest required"
}

deny[msg] {
  startswith(input.accessURL, "http://")
  msg := "insecure accessURL"
}
```

Gate	What it blocks	Where to enforce
Schema validate	Malformed DCAT/STAC/PROV	CI (pre-merge)
Receipt present	Untraceable raw/processed changes	CI + pipeline
Rights/CARE gate	Missing license; restricted data leaking	CI + API policy
Determinism gate	Unstable IDs/spec_hash mismatch	CI + pipeline

Source: CI checklist and rights gate pattern. (KFM New Ideas 2-8-26 p.167–168, p.98–100)

Reusable CI workflow shape (callable) is proposed to reuse gates across repos. (KFM New Ideas 2-8-26 p.50–51)

5. Module: Supply-Chain Evidence (SBOM + Attestations + Verification)

The document provides a GitHub Actions “fail-closed” skeleton: build, generate SBOM, attach/attest with cosign, verify signatures, verify attestation, then conftest policy-gate the results. (KFM New Ideas 2-8-26 p.134–135)

```
# CI/CD “fail-closed” skeleton (abridged)
- name: SBOM + provenance + attach/attest
  run: |
    syft $IMAGE -o cyclonedx-json > sbom.cdx.json
    cosign attach sbom --sbom sbom.cdx.json $IMAGE
    cosign attest --predicate slsa-provenance.json --type https://slsa.dev/provenance $I
MAGE
- name: Verify signatures
  run: cosign verify $IMAGE
- name: Policy gate (fail-closed)
  run: |
    conftest test slsa.att.json
    conftest test sbom.cdx.json
```

Kubernetes runtime enforcement is sketched using Kyverno to deny workloads unless images are cosign-signed and carry an attestation that includes required provenance fields like spec_hash and STAC pointers. (KFM New Ideas 2-8-26 p.136)

■■ Tooling risk note: the doc calls out a Cosign verification issue (CVE-2026-22703) and recommends pinning versions and adding regression tests. (KFM New Ideas 2-8-26 p.142–143)

6. Module: OCI Evidence Bundles (ORAS + Digest-Pinned Distribution)

Push datasets + catalogs + receipts + attestations as an OCI artifact and treat its digest as the canonical address. Clients resolve by digest and selectively fetch children descriptors. (KFM New Ideas 2-8-26 p.141–142)

Artifact piece	Suggested media type
Bundle config	application/vnd.kfm.stac-bundle+json
STAC Catalog/Collection/Item	application/vnd.stac.catalog+json / application/vnd.stac.collection+json / application/geo+js
SBOM	application/spdx+json (or CycloneDX JSON)
Run receipt	application/vnd.kfm.run-receipt+json
Attestation	application/vnd.in-toto.attestation
GeoParquet	application/x-parquet
COG	image/tiff

Drop-in tasks include helper scripts, STAC generator support for oci://…@sha256:… asset hrefs, a CI gate for bundle publish + signing, and an API resolver endpoint GET /bundles/{digest}.

7. Module: Provenance Timeline (Graph → API → UI)

The provenance timeline is a user-facing trace: agent → activity → ingestion → dataset, with visible policy bundles and linkable evidence (STAC/DCAT/PROV IDs + commit SHA). A Cypher query and API shapes are provided, along with suggested UI component paths. (KFM New Ideas 2-8-26 p.37–38)

Cypher query (timeline slice):

```
MATCH (d:Dataset {slug:$dataset_slug})<-[:VERSIONS]-(i:Ingestion)
OPTIONAL MATCH (i)<-[:GENERATED]-(act:Activity)<-[:PERFORMED]-(agent:Agent)
WHERE i.at >= datetime($from) AND i.at < datetime($to)
RETURN i.ingestion_id AS id, i.at AS at, i.status AS status, i.git_sha AS git_sha,
       i.policy_gate AS policy_gate, i.stac_id AS stac_id,
       act.kind AS activity_kind, agent.name AS who
ORDER BY at ASC;
```

REST API (proposal consistent with doc):

```
GET /provenance/timeline?dataset_slug=&from;=&to;= (cached; deny-by-default; cite graph)
```

Component	Responsibility
ProvenanceTimeline	Filter by dataset/date; render ingestion points; link to STAC/DCAT/receipt
Detail drawer	Show who/when/why; policy bundle; lineage parents; commit SHA

8. Module: Focus Mode Citation Handshake (Cite or Abstain)

When Focus Mode answers using items from a collection, the document requires: include dataset DOI + provider key, link to DCAT landing page and STAC collection, and surface CARE flags if sensitive fields were generalized. (KFM New Ideas 2-8-26 p.167)

Requirement	UI/Output behavior
Include dataset DOI + provider keys	Citation block contains DOI + source keys (e.g., datasetKey).
DCAT + STAC links	Citation block includes resolvable links to DCAT landing and STAC collection/item.
CARE/sensitivity surfaced	If coordinates generalized (e.g., 1 km), show care.privacy_level + permission badge

Engineering hook: validate Focus Mode outputs against a schema that requires at least one citation object with dataset identifiers and resolvable evidence links. Treat missing citations as policy-deny.

9. Module: Story Nodes as “Proof of Integration” (Micro-story + Evidence Hooks)

A compact Story Node arc is provided to showcase provenance-first stitching: NHD flowlines + 3DEP elevation change + USDA soils → a 2-minute watershed narrative. It includes front matter, scenes, map layer hints, ‘Cite’ buttons, example Graph/API queries, ETL notes, and acceptance criteria. (KFM New Ideas 2-8-26 p.38–41)

Front matter (example):

```
title: "Where Water Learns the Land"
slug: "watershed-two-minutes"
duration_sec: 120
dataset_refs:
  - usgs:nhd:flowlines
  - usgs:3dep:elevation-change
  - usda:ssurgo:soils
governance:
  posture: evidence-first
  policies: [FAIR, CARE, OPA-fail-closed]
provenance:
  requires_citations: true
stac_collections:
  - kfm:nhd:flowlines:v1
  - kfm:3dep:elvchange:v1
  - kfm:usda:ssurgo:v1
```

Evidence hook pattern:

Each annotation renders a **Cite** button that opens a side panel listing STAC item, license, lineage hash, and the API query used. Tooltips auto-append source_ref IDs like stac://… (KFM New Ideas 2-8-26 p.40).

DoD (micro-story):

■ Acceptance criteria
Story plays in ≤2 minutes with timeline scrubber.
Each on-screen claim has one-click provenance (STAC/DCAT/receipt).
Map layers load in <2 seconds at default zoom on dev hardware.
Story Node lints against schema and passes OPA policy-pack:story.

10. Module: Watchers for Authority Drift (Canonicalize → Digest → Diff → PR)

Watcher flow: seed inputs → normalize (URDNA2015 preferred; JCS fallback) → sha256 digests per stable ID → classify diffs into typed deltas → policy & privacy gate → signed run_receipt → human review PR → deterministic apply to PostGIS/Neo4j/Story Node revisions. (KFM New Ideas 2-8-26 p.224–229)

Folder layout (example):

```
/watchers/authorities/
  fetch/        (lcnaf.py, viaf_dump.py, wikidata_stream.py, kansas_oai_pmh.py)
  normalize/    (rdf_c14n.py URDNA2015, json_c14n.py JCS)
  diff/         (engine.py, detectors.py)
  policy/       (rego/, conftest/)
  emit/         (delta_writer.py, pr_report.py)
  data/
    prov/runs/YYYY/MM/DD/...
    deltas/YYYY/MM/DD/authority/...
    cache/etags/...
```

Privacy guard (CARE defaults):

If dateOfBirth suggests a living person, collapse to decade/year; drop exact street addresses; store generalized geo (county or bbox); mark care:livingPerson=true and kfm:restricted=true; deny story-node mutation unless policy allows. (KFM New Ideas 2-8-26 p.227)

11. Module: Governed Entity Packages (Example: ‘Bleeding Kansas’ Cluster)

Commit-ready entity cluster scaffolding: authority map for people (LCNAF/VIAF), site exemplars (KHRI/NRHP), Story Node metadata JSON, PROV receipts, pipeline hooks, and ops checklist. (KFM New Ideas 2-8-26 p.168–172)

Entity schema intent (abridged):

```
required: [canonical_id, class, provenance]
canonical_id: "kfm:taxon:itis:40462" | "kfm:person:lc:n84024412"
class: BiologicalEntity | HistoricalPerson
provenance: { run_receipt, spec_hash, source_url, fetched_at }
```

Operational guidance: resolve canonical IDs first; treat KHRI/NRHP geometry as authoritative; respect restricted coordinates and sensitivity flags; capture rights_holder/license/contact; publish Story Nodes only after receipts + policy gates. (KFM New Ideas 2-8-26 p.168–172)

12. Proposed 2–4 Sprint Integration Plan (Safe Defaults)

The source doc provides many “drop-in” components. This section turns them into a small, reversible delivery plan. The critical idea is to land contracts + gates before scaling ingestion.

Sprint	Deliverables	Acceptance / verification
S1: Contracts	Receipt schema + validator; DCAT/STAC link rules; Dirac trip skeleton	Dirac trip → receipt → catalogs validate
S2: Gates	Reusable CI workflow; conftest policy pack; rights/OARE checks	Merge blocked on missing receipt/license; policy tests green
S3: Provenance UX	Neo4j prov model + /provenance/timeline; UI timeline shows policy	UI shows policy bundle + commit SHA + STAC links
S4: Proof Story	2-minute watershed Story Node + Cite panel; Focus Mode citations schema	Each claim has evidence; Focus outputs cite or abstain

Verification steps (minimum work): (1) run schema validation on a known-good DCAT/STAC sample; (2) intentionally break a receipt (remove artifact_digest) and confirm conftest denies; (3) run the provenance timeline query for a known dataset slug and confirm policy_gate + stac_id show up; (4) run a Focus Mode output fixture through the output schema and confirm missing citations fail.

Appendix A. Copy/Paste Checklists

A1) PR checklist for dataset promotion

Checklist item
DCAT identifier[] includes DOI + provider keys
STAC collection.id stable; properties.ids.canonical[] present
License + rights holder mapped; care.* and sensitivity captured
Bulk ingestion used (no ad-hoc scrapes where prohibited)
prov:run_receipt written with digests + (optional) cosign
CI passes: schema, receipts, rights gates
Focus Mode citations show dataset DOI + provider keys

Source: PR checklist bullets. (KFM New Ideas 2-8-26 p.168)

A2) Watcher success checklist (authority drift)

Checklist item
ETag/Last-Modified respected for HTTP sources
URDNA2015/JCS canonicalization locked and versioned
sha256 per ID stored; diff classified into typed deltas
CARE/privacy defaults on; living-person fields generalized
OPA checks pass; failures block with clear reason
prov:run_receipt signed and stored
Human-review PR includes machine-readable changelog
Deterministic apply to PostGIS/Neo4j + Story Node revision

Source: watcher checklist bullets. (KFM New Ideas 2-8-26 p.228–229)

Appendix B. Notes on What Is Not Confirmed Here

This kit is derived from a single internal document. Where specific endpoint names, schema URLs, or repo folder names differ in your repository, treat them as placeholders and verify by: (1) searching the repo for existing OpenAPI/GraphQL contracts, (2) checking Story Node v3 schema definitions, and (3) aligning policy bundle names/IDs with your policy-as-code repo.

Kansas Frontier Matrix (KFM)

Integration Idea Pack

Cultivated from internal pattern compendium:
Latest ideas 1-4-26-1.pdf

Generated: 2026-02-15 | **Doc version:** 1.0

Purpose

This document reorganizes and expands the most integration-ready patterns from the source compendium into a set of implementation cards that can be pasted into KFM planning, tickets, and governed documentation. Each card includes: intent, interfaces/contracts, suggested repo paths, CI/policy gates, observability hooks, and a crisp Definition of Done.

Governance note

Several patterns include security, provenance, and policy enforcement mechanisms (Cosign/SLSA, OPA/Rego gates, FAIR+CARE metadata). Treat these as governed system behaviors: review with KFM governance owners before enabling in production.

Table of Contents

Purpose . . . 1

Governance note . . . 1

Table of Contents **2**

Executive Summary **6**

Integration Module Registry . . . 6

UI-AUTOMATION-STATUS: UI Automation Status Badges (MapLibre + Cesium) **7**

Intent & UX behavior . . . 7

Contracts & endpoints . . . 7

Placement in KFM (clean layers) . . . 7

CI, tests, and telemetry . . . 8

Definition of Done . . . 8

GRAPH-QA-PROMOTION-GATE: Graph QA + Promotion Gate (Neo4j Cypher) **9**

Intent . . . 9

Graph invariants (example Cypher) . . . 9

CI integration . . . 9

Security & governance notes . . . 10

STAC-CARE-EXT: KFM CARE STAC Extension (policy-aware metadata) **11**

Intent . . . 11

Schema & decision logic . . . 11

Implementation steps . . . 12

DISCOVERY-ENVELOPE: Upstream Discovery Envelope (event-first, poll-fallback) **13**

Intent . . . 13

Discovery contract . . . 13

Implementation patterns (event-first, poll-fallback) . . . 13

Where discoveries go next . . . 14

DISCOVERY-TO-DRAFT-PR: Discovery → Draft PR (STAC/DCAT + PROV, deterministic IDs) 15

Intent . . . 15

Deterministic IDs & minimal provenance keys . . . 15

Repo layout (recommended excerpt) . . . 15

Generated artifacts (examples) . . . 16

Policy gates . . . 16

PR-PREVIEWS-ROLLBACK-FIRST: PR-Centric Orchestration (Preview → Attest → Promote, rollback-first) 17

Intent . . . 17

Core mechanics . . . 17

Preview run sketch . . . 17

Governance & compliance . . . 17

Definition of Done . . . 17

SUPPLY-CHAIN-HARDENING: Supply-chain hardening (SBOMs + SLSA/Cosign attestations) 18

Intent . . . 18

CI build pattern (high-level) . . . 18

Local attestation/verification sketch . . . 18

LICENSE-SPDX-GUARD: SPDX License Guard (CI gate for code + data catalogs) 19

Intent . . . 19

What it checks . . . 19

Drop-in workflow (sketch) . . . 19

OBSERVABILITY-OTEL-MAPPINGS: Telemetry Pattern (OTel → Prometheus/Mimir → Tempo, mapped to catalogs) 20

Intent . . . 20

Canonical keys . . . 20

Catalog linkage . . . 20

SOILS-ARTIFACT-RUNBOOK: Soil Data Production Artifacts (GeoParquet/PMTiles + COG) runbook 21

Intent . . . 21

Deterministic naming . . . 21

Governance & reproducibility checklist . . . 21

Outputs (suggested) . . . 21

WZDX-PIPELINE: WZDx v4.x → GeoParquet/PMTiles pipeline (CI-friendly) 22

Intent . . . 22

Drop-in directory layout . . . 22

Field mapping & determinism . . . 22

Freshness watch (CI efficiency) . . . 22

AIR-QUALITY-INGEST: Air Quality Ingest Flow (vendors → Fleet Registry → SensorThings → STAC/PROV) 23

Intent . . . 23

Minimal metadata fieldset (registry + STAC/PROV) . . . 23

Directory layout (pasteable) . . . 23

Operational checklist . . . 23

ARROW-NATIVE-GEO: Arrow-native Geospatial Pipeline Pattern (GeoArrow → GeoParquet 1.1 → IPC/Flight) 25

Intent . . . 25

Reference layout (drop-in) . . . 25

Data contract (columns & metadata) . . . 25

QA & CI checks . . . 25

Example flow & placement . . . 25

KG-SUMMARY-NODES: Knowledge-Graph Summary Nodes (rollups to speed Focus Mode) 27

Intent . . . 27

Placement & contract . . . 27

Prototype plan (thin slice) 27

Governance & observability 27

Definition of Done 27

Integration Roadmap (thin slices) **28**

Governance Review Checklist **29**

Glossary (selected) **30**

Source Reference **31**

Executive Summary

The source document contains many drop-in patterns. This pack prioritizes the ones that most directly strengthen KFM's trust membrane: deterministic IDs, catalog/provenance completeness, promotion gates, supply-chain attestations, and policy-aware metadata. The recommended integration order is:

- **Foundation (P1):** Discovery envelope → Draft PR automation → PR-preview orchestration + rollback-first → Supply-chain + provenance gates.
- **Governance (P1):** Graph QA promotion gate + CARE metadata extension (consent/constraints) to enforce safe publication.
- **UX (P1):** On-map automation status badges to make lineage visible to users without exposing secrets.
- **Domain pipelines (P2+):** Soil artifacts, WZDx ingestion, Air-quality SensorThings pipeline, Arrow-native geo performance pattern.
- **Performance (P3):** Summary nodes to accelerate Focus Mode reads with staleness checks and redaction.

Integration Module Registry

A quick index of the cultivated modules and where they map into KFM:

Priority	Module ID	Name	Domain	Source pages
P1	UI-AUTOMATION-STATUS	UI Automation Status Badges (MapLibre + Cesium)	UI / Provenance UX	pp. 2-6
P1	GRAPH-QA-PROMOTION-GATE	Graph QA + Promotion Gate (Neo4j Cypher)	Governance / QA	pp. 10-14
P1	STAC-CARE-EXT	KFM CARE STAC Extension (policy-aware metadata)	Governance / Metadata	pp. 48-50
P1	DISCOVERY-ENVELOPE	Upstream Discovery Envelope (event-first, poll-fallback)	Ingestion / Orchestration	pp. 202-209
P1	DISCOVERY-TO-DRAFT-PR	Discovery → Draft PR (STAC/DCAT + PROV, deterministic IDs)	Ingestion / Catalog Promotion	pp. 209-211
P1	PR-PREVIEWS-ROLLBACK-FIRST	PR-Centric Orchestration (Preview → Attest → Promote, rollback-first)	CI/Ops / Safety	pp. 92-96
P1	SUPPLY-CHAIN-HARDENING	Supply-chain hardening (SBOMs + SLSA/Cosign attestations)	Security / CI	pp. 111-118, 122
P2	LICENSE-SPDX-GUARD	SPDX License Guard (CI gate for code + data catalogs)	Compliance	pp. 160-166
P2	OBSERVABILITY-OTEL-MAPPINGS	Telemetry Pattern (OTel → Prometheus/Mimir → Observability)	Observability (mapped to catalogs)	pp. 50-55
P2	SOILS-ARTIFACT-RUNBOOK	Soil Data Production Artifacts (GeoParquet/PMTiles specs)	Pipelines / Geospatial	pp. 71-80, 82-84
P2	WZDX-PIPELINE	WZDx v4.x → GeoParquet/PMTiles pipeline (CI-friendly)	Transportation	pp. 166-169
P2	AIR-QUALITY-INGEST	Air Quality Ingest Flow (vendors → Fleet Registry → SensorThings → STAC/PROV)	Air Quality / Time-series	pp. 22-24
P3	ARROW-NATIVE-GEO	Arrow-native Geospatial Pipeline Pattern (GeoArrow/GeoParquet → IPC/Flight)	Performance / Geospatial	pp. 96-99
P3	KG-SUMMARY-NODES	Knowledge-Graph Summary Nodes (rollups to speed Focus Mode)	Knowledge Graph / Focus Mode	pp. 159-160

UI-AUTOMATION-STATUS: UI Automation Status Badges (MapLibre + Cesium)

Priority: P1 **Domain:** UI / Provenance UX **Source:** pp. 2-6

Why this exists: Show per-feature data freshness and pipeline health directly on the map, with drill-down to attestations, SBOMs, manifests, and logs.

Intent & UX behavior

Render small health badges on top of map features (2D MapLibre and 3D Cesium) that update in near-real time. Clicking a badge opens a drawer with run metadata and verified provenance links (attestation/SBOM/manifest/logs) through a server proxy.

- Join key: **featureId** in status events ↔ **id** (or configured property) in map features.
- Update model: streaming first (SSE/WebSocket), fallback polling every 15s if stream unavailable.
- Degrade gracefully: if attestation missing, show status but disable provenance link with explanation.
- Accessibility: keyboard focus, ARIA labels, and safe color semantics.

Contracts & endpoints

Define a small, stable JSON schema for badge events (validated in CI). Expose two API surfaces: (1) a streaming endpoint for live updates, and (2) a polling endpoint for degraded mode.

```
// Example (TypeScript) - AutomationBadge event
export type AutomationBadge = {
  featureId: string;
  status: "healthy" | "degraded" | "failing" | "running" | "unknown";
  last_run_ts: string; // RFC3339
  lineage: {
    run_id: string;
    backend: "openlineage" | "prov";
    links: {
      attestation?: string;
      sbom?: string;
      manifest?: string;
      logs?: string;
    };
  };
  extras?: Record<string, unknown>; // e.g., SLO, error rates
};
```

- GET **/api/lineage/stream** (SSE or WebSocket) → pushes AutomationBadge events.
- GET **/api/lineage/status?since=<ts>** → returns list of events since timestamp.
- Server must verify Cosign/Sigstore attestations before exposing links; browser must never fetch untrusted attestations directly.

Placement in KFM (clean layers)

Keep UI rendering logic in the frontend (React/MapLibre/Cesium), but keep provenance verification and policy decisions behind the governed API gateway (trust membrane).

- **Frontend:** overlay renderer + drawer component; consumes only API responses (no direct DB/object-store access).
- **Backend service:** lineage status service translates OpenLineage/PROV into AutomationBadge events.
- **Integration adapters:** MapLibre adapter and Cesium adapter attach overlays without coupling to provenance source.
- **Infrastructure:** optional event bus + cache (for polling) + signature verification tooling.

Suggested repo paths (adapt to your layout):

- web/packages/ui-automation-status/ (shared badge logic + schemas)
- web/apps/map/ (MapLibre integration)
- web/apps/globe/ (Cesium integration)
- api/services/lineage-status/ (SSE + polling endpoints; attestation verification proxy)
- schemas/ui/automation_badge.schema.json (AJV/JSONSchema)

CI, tests, and telemetry

- **Unit:** adapter tests (mock map APIs), datasource transform tests, AJV schema validation on fixtures.
- **E2E:** Playwright camera paths snapshots to verify overlay placement across zoom/tilt.
- **OTel events/metrics:** emit badge-rendered and status-changed events; measure stream latency and overlay render time.

Definition of Done

- Badges render and update within **< 250ms** after event receipt (target from source).
- Works in both MapLibre and Cesium demo apps with graceful fallback polling.
- Schema validation enforced in CI; provenance links are verified server-side before exposure.

GRAPH-QA-PROMOTION-GATE: Graph QA + Promotion Gate (Neo4j Cypher)

Priority: P1 **Domain:** Governance / QA **Source:** pp. 10-14

Why this exists: Hard-fail missing attestations, duplicate asset href collisions, and digest mismatches before any dataset promotion.

Intent

Use the knowledge graph (Neo4j) as the last-line quality gate before promotion. Treat provenance completeness as graph invariants and run repeatable Cypher checks in CI.

- Detect missing attestations for promoted artifacts.
- Detect duplicate asset href collisions (two STAC items pointing at same href unintentionally).
- Detect digest mismatches between STAC metadata and artifact digests.
- Fail the pipeline on any anomaly; otherwise, mark the release as promotable.

Graph invariants (example Cypher)

Compact sketches of the three checks (tune labels/props to your KFM graph model):

```
// Example (Cypher) - anomaly pattern sketches
// 1) Missing attestation for generated item
MATCH (i:StacItem)<-[:GENERATED]-(r:Run)
WHERE NOT (r)-[:HAS_ATTESTATION]->(:Attestation {verified:true})
RETURN "MISSING_ATTESTATION" AS kind, i.id AS item_id, r.run_id AS run_id;

// 2) Duplicate asset href across items
MATCH (a:Asset)
WITH a.href AS href, collect(a) AS assets
WHERE size(assets) > 1
RETURN "DUPLICATE_ASSET_HREF" AS kind, href, size(assets) AS count;

// 3) Digest mismatch
MATCH (a:Asset)
WHERE a.digest_sha256 IS NOT NULL AND a.digest_sha256 <> a.stac_digest_sha256
RETURN "DIGEST_MISMATCH" AS kind, a.href AS href;
```

CI integration

- Run these queries as a **pre-promotion** job in CI.
- Export anomaly report (CSV/JSON) as build artifact and attach to PR comment.
- Fail if any anomaly exists; on success require **promotable=true** to allow release tags.

Suggested repo paths (from source)

- tools/validation/graph/qa_stac_attestations.cypher
- .github/workflows/graph-qa.yml
- docs/quality/graph-qa/neo4j_stac_attestation_checks.md

Security & governance notes

- Verification must be Sigstore/Cosign (or approved key mode) and recorded as graph entities.
- Preserve immutable history: new content ⇒ new version ⇒ new digests/attestations.

STAC-CARE-EXT: KFM CARE STAC Extension (policy-aware metadata)

Priority: P1 **Domain:** Governance / Metadata **Source:** pp. 48-50

Why this exists: Encode consent, authority-to-control, usage constraints, and ethics risk mitigations into STAC so policy can be enforced consistently.

Intent

Add a small STAC extension block that encodes CARE-like governance signals (consent status, stewardship review requirement, usage constraints, and ethics sensitivity). These fields become indexable facets for policy-aware search and can be enforced in CI and at the API boundary.

Schema & decision logic

Keep the schema opinionated and small. Put long narrative policy text in governance docs and link via a basis field. Below is a trimmed example showing the fields referenced by the source document.

```
// Illustrative JSON Schema excerpt (trimmed)
{
  "$id": "https://kfm.example.org/stac/extensions/kfm-care/schema.json",
  "type": "object",
  "properties": {
    "properties": {
      "type": "object",
      "properties": {
        "kfm:care": {
          "type": "object",
          "properties": {
            "consent": {
              "type": "object",
              "properties": {
                "status": { "enum": ["granted","conditional","withheld"] },
                "basis": { "type": "string" }
              }
            },
            "steward_review": {
              "type": "object",
              "properties": { "required": { "type": "boolean" } }
            },
            "usage_constraints": {
              "type": "object",
              "properties": {
                "prohibited": { "type": "array", "items": { "type": "string" } }
              }
            },
            "ethics": {
              "type": "object",
              "properties": { "sensitive": { "type": "boolean" } }
            }
          }
        }
      }
    }
  }
}
```

Minimal decision table (policy gate)

- **consent.status = withheld** → deny access to assets; expose only a metadata stub.
- **consent.status = conditional** and **steward_review.required = true** → queue steward approval before distributing assets.
- **ethics.sensitive = true** → require risk mitigations before distribution.
- **intended use** ∈ **usage_constraints.prohibited** (e.g., law-enforcement-surveillance) → block and record an audit note.

Implementation steps

- Declare the extension URL in **stac_extensions** for Collections/Items.
- Add CI validation via pystac/stac-check with extension schema enabled.
- Default policy gate at Collection; allow Item-level overrides when consent/mitigations differ.
- Index kfm:care.* fields in catalog/search to enable policy-aware filtering.
- Reuse the same did: agent IDs in PROV (prov:wasAssociatedWith) so policy ↔ pipeline actions are traceable.

Suggested validation command

```
pystac validate --extensions schemas/stac/extensions/kfm-care/schema.json \
  data/stac/collections/your_collection.json
stac-check data/stac/items/**/*.json
```

DISCOVERY-ENVELOPE: Upstream Discovery Envelope (event-first, poll-fallback)

Priority: P1 **Domain:** Ingestion / Orchestration **Source:** pp. 202-209

Why this exists: Normalize upstream change detection into one idempotent message, enabling deterministic pipelines and PR-bot automation.

Intent

Standardize something changed upstream into a single discovery message that is idempotent, auditable, and safe to process repeatedly. This becomes the trigger for ETL runs, catalog updates, and PR automation.

- Use object-store notifications when available; fall back to polite HTTP HEAD polling with ETag/Last-Modified.
- Compute an **idempotency_key** from stable provenance fields so duplicate triggers collapse naturally.
- Include provenance fields (URI, ETag, size, modified time) and log every discovery as JSONL for audits.

Discovery contract

The discovery envelope should be schema-validated and stable. Keep it small; it is a trigger + provenance stub, not the payload itself.

```
// Example discovery envelope (JSON)
{
  "schema": "io.kfm.discovery.v1",
  "idempotency_key": "sha256:…",
  "occurred_at": "2026-01-05T12:34:56Z",
  "kind": "object|feed_entry",
  "source": { "provider": "cloudflare:r2|aws:s3|http", "bucket": "...", "key": "...", "url": "..." },
  "provenance": { "uri": "r2://bucket/key", "etag": ""…"", "size_bytes": 12345, "last_modified": "…" }
  "hints": { "expected_followup": "GET", "format": "application/json" },
  "governance": { "license_hint": "…", "sensitivity_hint": "public|restricted" }
}
```

Implementation patterns (event-first, poll-fallback)

Pattern 1: Object-store event → webhook (example: Cloudflare R2).

```
// Cloudflare Worker sketch (trimmed)
export default {
  async fetch(req, env) {
    const ev = await req.json();
    const { bucket, object } = ev;
    const base = `${bucket}|${object.key}|${object.etag||""}|${object.size||""}`;
    const idempotency_key = "sha256:" + await sha256(`cloudflare:r2|${base}`);
    const discovery = {
      schema: "io.kfm.discovery.v1",
      idempotency_key,
      occurred_at: new Date().toISOString(),
      kind: "object",
      source: { provider: "cloudflare:r2", bucket, key: object.key },
      provenance: { uri: `r2://${bucket}/${object.key}`, etag: object.etag, size_bytes: object.size }
    };
```

```
    await fetch(env.OUT_WEBHOOK, { method:"POST", body: JSON.stringify(discovery) });
    return new Response("ok");
  }
};
```

Pattern 2: HTTP poller with ETag/If-None-Match fallback (only emits when changed).

```
// HTTP HEAD poller sketch (trimmed)
const headers = { "If-None-Match": cp.etag ?? "", "If-Modified-Since": cp.last_modified ?? "" };
const res = await fetch(url, { method: "HEAD", headers });
if (res.status === 304) return { next: cp };
const etag = res.headers.get("etag");
const last_modified = res.headers.get("last-modified");
const idempotency_key = sha256(`http|${url}|${etag||""}|${last_modified||""}`);
```

- Persist polling checkpoints in a small KV (Dynamo/R2 KV/Redis/etc.).
- Schedule every 1–5 minutes; throttle politely and respect robots.txt/ToS for HTTP sources.
- Do not inline secrets; use workload identity / IAM roles and include emitter identity for audit trails.

Where discoveries go next

- Send to an event bus/queue (with DLQ) or directly to an orchestrator webhook.
- Optionally send to a GitHub App PR-bot that produces catalog updates (see next module).

DISCOVERY-TO-DRAFT-PR: Discovery → Draft PR (STAC/DCAT + PROV, deterministic IDs)

Priority: P1 **Domain:** GitOps / Catalog Promotion **Source:** pp. 209-211

Why this exists: Turn dataset discovery into a draft PR that updates catalogs and blocks merge until provenance checks pass.

Intent

When a dataset is discovered, automatically open a draft PR that adds/updates STAC Items/Collections, DCAT datasets, and PROV run bundles with deterministic IDs. Block merge until validation and attestation policy gates pass.

- PR becomes the single source of truth for discovery metadata and promotion intent.
- Catalog updates are generated deterministically so diffs are reviewable.
- OPA/Rego checks enforce minimal STAC/DCAT/PROV keys plus cryptographic provenance.

Deterministic IDs & minimal provenance keys

- **Item ID:** <domain>-<source>-<yyyymmdd>-<geohash5>-<sha8(fingerprint)>
- **Collection ID:** <domain>-<theme>-<version>
- **Dataset fingerprint:** stable hash over canonicalized fields (URIs, extents, bbox, checksum of primary artifact).

Minimal required provenance keys (KFM core)

- prov:wasGeneratedBy
- prov:used
- kfm:run_id
- kfm:artifact_digest
- kfm:attestation_uri
- kfm:source_license
- kfm:data_sensitivity

Repo layout (recommended excerpt)

```
data/
  <domain>/
    stac/
      items/
      collections/
    dcat/
      datasets/
    prov/
      runs/
```

```
tools/
  validation/
    rego/
.github/
  workflows/
    discovery-draft-pr.yml
```

Generated artifacts (examples)

The PR-bot should generate small, valid JSON documents (STAC + DCAT + PROV). Keep optional fields out of the critical path; add them over time.

```
// Minimal STAC Item sketch (trimmed)
{
  "type": "Feature",
  "stac_version": "1.0.0",
  "id": "hydrology-usgs-20251214-dr5-6f12a9cd",
  "collection": "hydrology-streamflow-v1",
  "properties": {
    "datetime": "2025-12-14T00:00:00Z",
    "kfm:run_id": "RUN-2025-12-14T01-03-22Z-0af1c9",
    "kfm:artifact_digest": "sha256:…",
    "kfm:attestation_uri": "oci://…@sha256:…",
    "kfm:source_license": "…",
    "kfm:data_sensitivity": "public"
  },
  "links": [
    { "rel": "prov", "href": "./prov/runs/<run_id>.jsonld", "type": "application/ld+json" }
  ]
}
```

Policy gates

- Run Conftest/OPA against STAC/DCAT/PROV JSON to enforce required keys and allowed values.
- Verify Cosign/SLSA attestations for referenced artifacts before allowing merge.
- Block merge until validations pass; exceptions must reference a governance ticket.

PR-PREVIEWS-ROLLBACK-FIRST: PR-Centric Orchestration (Preview → Attest → Promote, rollback-first)

Priority: P1 **Domain:** GitOps / Safety **Source:** pp. 92-96

Why this exists: Use lakeFS preview branches + idempotent ETL + policy gates; rollback quickly if post-merge checks fail.

Intent

Make every pull request produce an isolated preview of data and artifacts, with full provenance and an explicit rollback path. The PR preview environment is reproducible via an idempotency key.

Core mechanics

- **Preview branch (lakeFS):** create preview-$PR_NUMBER from prod; copy-on-write data changes.
- **Idempotency:** derive key from input URI + model digest + config digest; reruns overwrite the same preview outputs.
- **Digest-addressed outputs:** write artifacts to preview paths using digests to avoid ambiguity.
- **Lineage:** emit OpenLineage events and PROV bundles and link them in STAC/DCAT.

Preview run sketch

```
# Example preview run (trimmed)
./run_etl.sh --out "s3://kfm-staging/preview-$PR_NUMBER/" \
  --idempotency-key "$IDEMP_KEY" --seed 424242
lakefs commit -m "local preview $PR_NUMBER"
emit_openlineage --run_id "local-$PR_NUMBER"
create_stac_item --properties "kfm:branch=preview-$PR_NUMBER,..."
```

Governance & compliance

- All runs must produce STAC/DCAT updates, OpenLineage events, Cosign attestations, and PROV activities.
- Merges require passing OPA/Rego policy + signature verification; exceptions must reference a governance ticket.

Definition of Done

- Preview artifacts reproducible via idempotency key.
- STAC/DCAT entries updated in draft PR.
- OpenLineage + PROV present and linked.
- Cosign verification + OPA checks green.
- Rollback plan validated by automated tests.

SUPPLY-CHAIN-HARDENING: Supply-chain hardening (SBOMs + SLSA/Cosign attestations)

Priority: P1 **Domain:** Security / CI **Source:** pp. 111-118, 122

Why this exists: Bind artifacts to builds and runs; verify signatures and provenance before promotion.

Intent

Make artifact trust verifiable: every important artifact (COG, PMTiles, GeoParquet, containers) should be bound to a build identity and config hash, with an SBOM and a cryptographic attestation that CI verifies before promotion.

CI build pattern (high-level)

- Use CI with least privilege: contents:read; id-token:write only when keyless signing is needed.
- Build and push container images by digest; generate an SBOM (e.g., CycloneDX).
- Write a PROV JSON-LD build activity artifact with commit SHA, workflow, SBOM reference, and image tag/digest.
- Sign/attest artifacts with Cosign (DSSE) and verify attestations in pre-promotion jobs.

Local attestation/verification sketch

```
# Minimal make targets (trimmed)
digest:
  sha256sum results/artifact.bin > results/artifact.bin.sha256
  sha256sum params/params.yaml > results/config.sha256

attest: digest
  cosign attest --yes --type slsaprovenance \
    --predicate predicate.json --bundle results/attest.jsonl \
    results/artifact.bin

verify:
  cosign verify-attestation --type slsaprovenance \
    --bundle results/attest.jsonl results/artifact.bin
```

- Canonicalize configs before hashing (stable key ordering, minified JSON/YAML) to avoid false changes.
- Pin container image digests in CI to make predicates reproducible.
- Store bundle URIs in STAC/DCAT/PROV so downstream users can find attestations reliably.
- Gate publish/promote: verify-attestation must pass.

LICENSE-SPDX-GUARD: SPDX License Guard (CI gate for code + data catalogs)

Priority: P2 **Domain:** Compliance **Source:** pp. 160-166

Why this exists: Fail PRs when licensing signals are missing or invalid across repo, STAC, and DCAT.

Intent

Enforce that every PR contains SPDX-accurate licensing signals for both code and data artifacts, including catalog metadata (STAC/DCAT license fields). Fail fast when licensing is missing or invalid.

What it checks

- Package manifests (package.json, pyproject.toml, etc.).
- LICENSE files and SPDX-License-Identifier headers in source files.
- STAC/DCAT license fields for datasets/distributions.
- Outputs actionable hints when violations occur.

Drop-in workflow (sketch)

```
# .github/workflows/license-spdx-guard.yml (conceptual)
name: license-spdx-guard
on: [pull_request]
jobs:
  spdx:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v4
      - run: node tools/license/check_spdx.js  # or python equivalent
```

Implementation note: use a real SPDX parser and map results into a short report. Treat missing catalog license fields the same as missing code license metadata.

OBSERVABILITY-OTEL-MAPPINGS: Telemetry Pattern (OTel → Prometheus/Mimir → Tempo, mapped to catalogs)

Priority: P2 **Domain:** Observability **Source:** pp. 50-55

Why this exists: Standardize run identifiers and status in metrics/traces and reuse them in STAC/DCAT/PROV properties.

Intent

Make every pipeline run traceable across metrics, traces, and catalogs by standardizing on canonical attribute keys (OTel) and label names (Prometheus/Mimir). Reuse the same identifiers in STAC/DCAT/PROV to link runtime telemetry to dataset artifacts.

Canonical keys

Semantic	OTel attribute	Prom label	Notes
Run ID	kfm.job.run_id	job_run_id	UUID/ULID
Commit	kfm.git.commit_sha	commit_sha	7-40 hex
Status	kfm.job.status	status	scheduled|running|succeeded|failed|canceled
Started	kfm.job.started_at (RFC3339)	started_at	also emit unix seconds
Ended	kfm.job.ended_at (RFC3339)	ended_at	also emit unix seconds
Dataset	kfm.dataset.id	dataset_id	STAC/DCAT id
Pipeline	kfm.pipeline.name	pipeline	stable slug
Env	kfm.env	env	dev|staging|prod

Emit timestamps twice: RFC3339 strings for traces/catalog metadata, and unix seconds for Prom math. Use these keys in exporters and relabel consistently.

Catalog linkage

- Write run identifiers into STAC Item properties (kfm:run_id, commit sha, status).
- In PROV JSON-LD, model the run as prov:Activity and attach times and associated agents.
- Use dataset_id to join telemetry dashboards to catalog pages in Focus Mode.

SOILS-ARTIFACT-RUNBOOK: Soil Data Production Artifacts (GeoParquet/PMTiles + COG) runbook

Priority: P2 **Domain:** Pipelines / Geospatial **Source:** pp. 71-80, 82-84

Why this exists: A reproducible recipe for turning soil vectors/rasters into web and analytics artifacts with deterministic naming and sidecar metadata.

Intent

Provide a reproducible geospatial artifact pipeline that turns soil datasets into both analytics-friendly formats (GeoParquet) and map-ready formats (PMTiles vector tiles, COG rasters), with deterministic naming and STAC/PROV sidecars for governance.

Deterministic naming

```
# Deterministic run metadata (example)
DATASET_SLUG="soil-harmonized"
SRC_HASH="$(find data/work/soil/vectors -type f -print0 | sort -z | xargs -0 shasum | \
  shasum | awk '{print $1}')"
DATE_TAG="$(date -u +%Y%m%d)"
RUN_ID="${DATASET_SLUG}-${DATE_TAG}-${SRC_HASH:0:8}"
echo "RUN_ID=${RUN_ID}"
```

- Vectors: normalize → GeoParquet → tippecanoe → PMTiles.
- Rasters: translate to COG with tiling/overviews and cloud-optimized checks.
- Emit STAC + PROV sidecars per artifact; validate in CI.

Governance & reproducibility checklist

- Inputs are immutable (content-addressable or pinned commit/DOI).
- CRS is explicit and consistent (often EPSG:4326 unless required otherwise).
- Filenames include RUN_ID / RASTER_RUN_ID.
- STAC/PROV sidecars exist per artifact and validate against schemas.
- COGs pass gdalinfo cloud-optimized checks; PMTiles pass pmtiles inspect.
- CI logs attach RUN_ID and artifact digests (sha256).
- Licenses and attribution for sources included in STAC/DCAT.

Outputs (suggested)

```
data/processed/soil/
  parquet/
  tiles/   # PMTiles
  cog/     # Cloud-Optimized GeoTIFF
data/stac/soil/
data/prov/soil/
```

WZDX-PIPELINE: WZDx v4.x → GeoParquet/PMTiles pipeline (CI-friendly)

Priority: P2 **Domain:** Transportation **Source:** pp. 166-169

Why this exists: Validate GeoJSON against official schema, normalize, emit deterministic GeoParquet and optional PMTiles bundle.

Intent

Turn WZDx v4.x road event feeds into deterministic GeoParquet (analytics) and optional PMTiles (maps), with schema validation, freshness detection, and STAC packaging for KFM promotion.

Drop-in directory layout

```
tools/wzdx/
  README.md
  requirements.txt
  wzdx_validate_transform.py
  wzdx_freshness_watch.py
  schemas/
    road_event_feed.json   # pinned schema
  outputs/
    wzdx.parquet
    wzdx.pmtiles           # optional
  stac/
    wzdx.item.json         # minimal STAC template
```

Field mapping & determinism

Keep a core subset of fields with spec-accurate typing, and preserve the raw payload as an optional asset. Determinism rules: sort columns, sort rows by id, and cast types stably.

- features[].id → id (stable unique feature id)
- road_event_type → road_event_type
- start_date/end_date → dt_start/dt_end (UTC)
- lanes[].status → lane_statuses
- feed_info → kfm_source (map into STAC properties)
- geometry → geometry (cast to EPSG:4326 and validate)

Freshness watch (CI efficiency)

- Use conditional GET (ETag / Last-Modified) so CI only recomputes when the feed changes.
- Record each fetch attempt, status (304 vs full), and computed digests as PROV entities/activities.
- Emit STAC fields to track age and validation status if useful.

AIR-QUALITY-INGEST: Air Quality Ingest Flow (vendors → Fleet Registry → SensorThings → STAC/PROV)

Priority: P2 **Domain:** Air Quality / Time-series **Source:** pp. 22-24

Why this exists: A governed pipeline blueprint with deterministic device identity, QC hooks, and promotion metadata.

Intent

Create a governed air-quality ingestion pipeline that consolidates multiple vendors into a canonical fleet registry and SensorThings-compatible time-series, with deterministic deduplication and STAC/PROV capture at every promotion step.

Minimal metadata fieldset (registry + STAC/PROV)

- hardware_id (stable KFM UID for physical device; never re-used)
- sensorFirmware (vendor version string)
- calibration_reference (DOI/URL if possible)
- method_version (KFM method hash for transforms & QC)
- provenance_edges (activity/agent/input_ref/output_ref/time for PROV)
- install_lat/install_lon/install_alt (WGS84 + datum notes)
- owner_org, license, data_rights (align with FAIR+CARE)
- sensor_model/sensor_make/measurement_caps (PM2.5/PM10/NO2/etc.)
- site_context (indoor/outdoor, roadside/urban/background)
- valid_from/valid_to (lifecycle intervals)

Directory layout (pasteable)

```
data/air-quality/
  raw/
  work/
  processed/
  stac/
  prov/
  catalog/dcat/
src/pipelines/air_quality/
  ingest_vendor_pull.py
  normalize_frames.py
  dedupe_qc_promote.py
  sensorthings_export.py
schemas/air/
  observations.schema.json
  fleet_registry.schema.json
```

Operational checklist

- Acquire raw feeds and store verbatim by vendor/date path.
- Normalize to common schema (timestamp, ids, lat/lon, observed properties, units).
- Join/update Fleet Registry (vendor_device_id → canonical hardware_id).
- QC hooks (range, rate-of-change, correction, drift detection); never drop data silently.
- Deterministic dedupe with tolerance windows.
- SensorThings promotion (Thing/Sensor/Datastream upserts; Observations export).
- Create STAC Item per batch (raw bundle, normalized parquet, QC report, export).
- Emit PROV activity chain and attach attestations & telemetry; block promotion on failures.

ARROW-NATIVE-GEO: Arrow-native Geospatial Pipeline Pattern (GeoArrow → GeoParquet 1.1 → IPC/Flight)

Priority: P3 **Domain:** Performance / Interop **Source:** pp. 96-99

Why this exists: Use Arrow as the backbone for fast, zero-copy-ish geospatial processing, with QA and perf guardrails.

Intent

Use Apache Arrow as the backbone for geospatial pipelines: GeoArrow arrays in memory → GeoParquet 1.1 on disk → Arrow IPC or Arrow Flight for interchange. This minimizes serialization and keeps geometry columnar.

Reference layout (drop-in)

```
src/pipelines/geo/arrow_native/
  ingest_flight.py
  to_geoarrow.py
  write_geoparquet.py
  read_arrow_ipc.py
  h3_enrich.py
  tests/
```

Data contract (columns & metadata)

- geometry: GeoArrow (buffers + CRS via WKT/EPSG).
- attrs: scalar Arrow columns (int/float/utf8/decimal).
- Parquet key metadata: geo (GeoParquet JSON), creator, optional h3 fields.

QA & CI checks

- Schema pin: Arrow/Parquet fields + GeoParquet metadata match spec.
- Round-trip: GeoArrow → GeoParquet → GeoArrow invariants.
- Perf guardrail: throughput and open-time thresholds.
- Provenance: PROV activity for each conversion; STAC/DCAT links for datasets.

Example flow & placement

- Load features → cast/validate → optional enrich → write GeoParquet → publish IPC/Flight → register catalogs.

  Suggested placement

- ETL: src/pipelines/geo/arrow_native/*
- Artifacts: data/processed/*.geoparquet, data/ipc/*.arrow
- Catalogs: catalogs/stac/, catalogs/dcat/, prov/

- CI: schema round-trip, perf thresholds, PROV/STAC checks

KG-SUMMARY-NODES: Knowledge-Graph Summary Nodes (rollups to speed Focus Mode)

Priority: P3 **Domain:** Knowledge Graph / Focus Mode **Source:** pp. 159-160

Why this exists: Precompute rollups (e.g., watersheds) into summary nodes to reduce query cost and improve Focus Mode latency with governance controls.

Intent

Speed up Focus Mode by writing precomputed summary nodes in the knowledge graph (e.g., hydrology rollups). Focus Mode should read from summaries first and fall back to raw traversal if missing. Summaries must be freshness-checked and governed (redaction/aggregation where required).

Placement & contract

```
src/pipelines/
  kg/
  hydrology/
  summary-nodes/
docs/patterns/kg/hydrology/summary-nodes/
docs/catalogs/prov/
```

- Expose an API contract such as GET /v1/watersheds/:huc8 returning summary fields (schema-validated).
- CI should fail if rollups are stale beyond the class SLA (e.g., >24h for daily summaries).

Prototype plan (thin slice)

- Pick one pilot region (single HUC8).
- Run attach + rollup job; publish GET /v1/watersheds/:huc8.
- Wire Focus Mode story to read summary node first, then raw traversal fallback.
- Measure before/after: P95 latencies, DB hits, render time.

Governance & observability

- If summaries overlap protected areas/territories, apply redaction or aggregation before publishing.
- Log rollups with energy/carbon telemetry; inherit licenses from source datasets.
- Metrics: rollup duration, nodes written, stale node count; traces around Focus Mode summary fetch.

Definition of Done

- Summary node exists for pilot region with fresh rollups.
- API schema + contract tests green.
- Provenance JSON-LD written and linked.
- Focus Mode renders from summary node; P95 improves on pilot.

Integration Roadmap (thin slices)

This roadmap is designed to minimize risk and maximize reviewability: each slice produces a governed artifact or a measurable capability, and each slice is reversible.

Phase	Goal	Outputs	Primary gates
0. Baseline	Define core schemas + policy stack	Discovery schema, AutomationBadge schema, CARE extension schema, OPA/Rego rules skeleton	CARE extension schema + OPA/Rego
1. Catalog automation	Discovery → Draft PR working end-to-end	PR bot generates STAC/DCAT/PROV with deterministic IDs; draft JSON-LD	Contract tests/DB + JSON-LD schema checks
2. Trustable promotion	Attestation + Graph QA gates	Cosign verify job; Neo4j anomaly queries; promotion flag	Signature verification + graph QA pass
3. Preview safety	PR previews with rollback-first	LakeFS preview branches; idempotent ETL; rollback test harness	Dry-run + rollback test
4. UX lineage visibility	Dirty map badges in Focus Mode	SSE stream + polling fallback; MapLibre + Cesium overlays; API proxy	Playwright E2E + proxy verification
5. Domain pipelines	Implement 1-2 exemplar pipelines	Soils artifacts runbook OR WZDx OR Air-quality ingest	End-to-end catalog + attestation + UI badge inte

Governance Review Checklist

Use this checklist before enabling any pattern in production environments.

- Policy review: CARE/consent fields and any data sovereignty flags reviewed by governance owners.
- Trust membrane: frontend never accesses databases/object stores directly; all access through governed API boundary.
- Security: Cosign/Sigstore verification implemented server-side; browser only sees verified, sanitized links.
- Licensing: SPDX signals present for code and data; STAC/DCAT license fields filled.
- Immutability: promoted artifacts are content-addressed; updates create new versions with new digests/attestations.
- Audit: discovery events and promotion actions logged as JSONL/PROV and retained.

Glossary (selected)

- **STAC**: SpatioTemporal Asset Catalog - JSON catalogs describing geospatial assets.
- **PROV-O**: W3C Provenance Ontology for expressing lineage (entities, activities, agents).
- **DCAT**: Data Catalog Vocabulary; used to describe datasets and distributions.
- **OpenLineage**: Standard event model for data lineage (jobs, runs, datasets).
- **Cosign/Sigstore**: Tooling for signing and verifying artifacts, images, and attestations.
- **SLSA**: Supply-chain Levels for Software Artifacts; provides provenance models and levels.
- **SBOM**: Software Bill of Materials (SPDX/CycloneDX).
- **OPA/Rego**: Policy-as-code used for CI gates (Conftest) and runtime checks.
- **GeoParquet**: Geospatial encoding for Apache Parquet; efficient analytics format.
- **PMTiles**: Single-file vector tile container optimized for static hosting/CDNs.
- **COG**: Cloud-Optimized GeoTIFF; tiled + overviews for efficient HTTP range reads.
- **H3**: Hexagonal hierarchical indexing system useful for tiling and diffing.
- **SensorThings**: OGC standard for sensor observation time-series APIs.
- **lakeFS**: Git-like branching for object storage; enables preview branches and rollback.

Source Reference

This idea pack was cultivated from: *Latest ideas 1-4-26-1.pdf*. Source pages referenced in each module header correspond to the locations of the originating pattern in the compendium.

If you want this pack exported as individual KFM Story Nodes or catalog-ready objects (JSON), the module cards in this PDF are already structured to map 1:1 into those formats.

KFM Integration Ideas Pack

Cultivated, integration-ready concepts derived from **Latest Ideas-2.pdf**

Source	Latest Ideas-2.pdf
Scope	Governed CI/CD, provenance, PR-first promotion, deterministic pipelines
Audience	Engineering leads, data stewards, governance reviewers
Output	A drop-in set of patterns, contracts, and thin-slice implementation steps

How to use this pack

Treat each section as an integration module: (1) adopt the non-negotiable invariants, (2) copy the suggested repo layout and CI skeleton, then (3) add one domain lane (for example soils or work zones) and wire it to policy and provenance gates.

Table of Contents

Table of Contents	2
Executive Summary	3
1. Non-negotiable Invariants	3
2. Agent System: Watcher, Planner, Executor (WPE)	3
3. Deterministic IDs, Idempotency, and Lineage	4
4. CI/GitOps Baseline: Reusable Lanes + Policy PR Gate	5
5. Policy Pack: OPA/Rego + Conftest	5
6. Provenance and Observability	6
7. STAC Extension: kfm:watcher-signals	6
8. Domain Lane Example: Soils Pipeline Runbook Pattern	7
9. Domain Lane Example: Work Zones (WZDx) as a Governed Feed	8
10. Thin-slice Implementation Plan (Buildable Increments)	8
11. Risks, Tradeoffs, and Governance Review	9
References	9

Executive Summary

Latest Ideas-2 proposes a governed, PR-first way to evolve KFM datasets, catalogs, and documentation without bypassing policy: a Watcher observes and emits immutable facts, a deterministic Planner turns facts into change plans, and a constrained Executor opens PRs with attached evidence (never merging directly). This pack extracts those ideas into a clean set of integration modules.

Core promise

Every change is reproducible, auditable, and reviewable: the system emits STAC/DCAT records, PROV-O lineage, and telemetry events; it is idempotent (safe to rerun) and fail-closed (no PR if gates fail).

1. Non-negotiable Invariants

Invariant	Why it exists
PR-first publishing:	Executors open or update PRs; they never merge and never push to protected branches.
Idempotency:	Every run has stable keys so reruns do not duplicate lineage or PRs.
Determinism:	Plans derive from pinned inputs plus a commit seed; same inputs rebuild identical artifacts.
Fail-closed gates:	Schema, policy, QA, and reproducibility checks must pass before PR creation.
Kill-switch:	A single central flag can stop Planner/Executor immediately.
Network boundaries:	No direct writes to production stores; only governed repo changes via PRs.

These invariants appear throughout the WPE agent blueprint, CI gates, and the deterministic validation patterns in Latest Ideas-2.

2. Agent System: Watcher, Planner, Executor (WPE)

The WPE split is the document's main architecture move: it reduces risk by separating observation, decision, and action into components with sharply limited permissions.

Component	What it does	What it must never do
Watcher	Observe sources and normalize signals into immutable facts and alerts.	Mutate code/data; write to production stores.

Planner	Convert facts + policies into deterministic change plans (plan.yml + diff.patch).	Make network calls that change outputs nondeterministically.
Executor	Open/update PRs and attach proofs (attestations, telemetry, gate reports).	Merge PRs; push to protected branches; bypass policy gates.

```
Suggested governed layout (drop-in)

docs/specs/agents/
  README.md
  WATCHER_CONTRACT.md
  PLANNER_CONTRACT.md
  EXECUTOR_CONTRACT.md
  examples/

artifacts/
  watcher/<ts>/facts.ndjson
  watcher/<ts>/alerts.json
  plans/<ts>/<subject>/plan.yml
  plans/<ts>/<subject>/diff.patch
  plans/<ts>/<subject>/evidence/

prov/
  watcher/<ts>/bundle.jsonld
  planner/<ts>/bundle.jsonld
  executor/<ts>/bundle.jsonld

telemetry/
  agents/*.events.json
```

3. Deterministic IDs, Idempotency, and Lineage

Latest Ideas-2 provides a concrete recipe for stable run keys and idempotency keys. The goal is twofold: (1) reruns do not create duplicate PRs or duplicate lineage entries, and (2) validation and promotion are content-addressed so they can be audited.

Deterministic key scheme (recommended)

run_id = sha256(repo_head_sha + planner_fingerprint + utc_day).
idempotency_key = sha256(run_id + lane_name + lane_input_digest).
content_digest = SHA256 (or BLAKE3) over canonicalized validated artifacts.

```
Python sketch (from the document's common/idempotency helpers)

def canonical_json(obj) -> bytes:
    return json.dumps(obj, sort_keys=True, separators=(",", ":")).encode("utf-8")

def run_id(repo_sha: str, planner_fp: str, utc_day: str) -> str:
    return sha256_hex(f"{repo_sha}|{planner_fp}|{utc_day}".encode())

def lane_key(run: str, lane: str, content_digest: str) -> str:
    return sha256_hex(f"{run}|{lane}|{content_digest}".encode())
```

4. CI/GitOps Baseline: Reusable Lanes + Policy PR Gate

The document includes a minimal GitHub Actions baseline: (a) a main entry workflow that calls a reusable workflow per lane, (b) a Conftest-based PR gate that enforces policy-as-code, and (c) OIDC permissions for keyless signing/attestation.

Artifact	Purpose	Integrate by
.github/workflows/kfm-ci.yml	Cron/manual entry that calls reusable CI.	Copy and parameterize lane and matrix defaults.
.github/workflows/reusables/kfm-reusable-ci.yml	Shared lane runner with guardrails, concurrency, and artifact packaging.	Adopt as the golden path for all domain lanes.
.github/workflows/kfm-policy-gate.yml	Merge-blocking policy checks on PRs via Conftest.	Point to tools/validation/policy and add unit tests.
.github/actions/setup-conftest	Tiny composite action to install a pinned Conftest version.	Pin tool versions; treat updates as governed changes.

Baseline guardrails

Adopt cancel-in-progress concurrency, environment/repo kill-switch, and short artifact retention by default. The reusable workflow pattern keeps lanes consistent: lanes differ by inputs, not structure.

5. Policy Pack: OPA/Rego + Conftest

A ready-to-drop policy folder structure is described, covering required catalog fields (STAC/DCAT/PROV), SPDX license allowlists, and domain QC bundles. This provides a single, merge-blocking governance layer across domains.

```
policy/
  rego/
    common/helpers.rego
    common/license_allowlist.rego
    catalogs/stac_required.rego
    catalogs/dcat_required.rego
    catalogs/prov_required.rego
    domains/dirt_geometry.rego
    domains/air_nowcast.rego
    domains/logistics_temporal.rego
    bundles.rego
  tests/
    *_test.rego
    samples/
```

Policy design guideline

Keep gates explainable: deny messages should clearly point to the missing field, the violated constraint, and the remediation. Treat policy bundles as versioned, reviewed artifacts.

6. Provenance and Observability

The document uses two complementary provenance streams: OpenLineage events for observability tooling, and PROV-O JSON-LD bundles for standards-based lineage (who/what/when used or generated). It also provides a minimal mapping from GitHub PR activity into PROV so CI can emit merge provenance and enforce invariants.

Stream	Where it helps	Minimum artifact
OpenLineage	Run-level telemetry: job/run IDs, inputs/outputs, execution state.	artifacts/openlineage.json
PROV-O JSON-LD	Auditable lineage: entities, activities, agents, used/generated relationships.	prov/<component>/<ts>/bundle.jsonld
Build provenance + attestations	Supply-chain evidence attached to builds and PRs.	artifacts/build_provenance.json + Sigstore/Cosign attest

Idempotent gate semantics

Rerunning CI is expected: downstream sinks should be keyed by idempotency_key so repeated validations do not duplicate lineage. Promotion only occurs after green validation.

7. STAC Extension: kfm:watcher-signals

To make automated updates explainable, the document proposes a compact STAC extension that records watcher/fetch metadata (what changed, how it was detected, what was validated) without bloating STAC Items.

Design goals

Compact and forward-compatible. Keep all fields optional and additive under the kfm: prefix. Use JSON Schema draft-07 for ecosystem compatibility.

```
Suggested placement

docs/specs/api/schemas/stac/extensions/KFM__WATCHER_SIGNALS.md
docs/specs/api/schemas/stac/extensions/kfm-watcher-signals.schema.json
```

8. Domain Lane Example: Soils Pipeline Runbook Pattern

Latest Ideas-2 includes a concrete, end-to-end soils runbook that illustrates the preferred pattern: snapshot raw sources, transform deterministically, emit modern delivery formats (GeoParquet/COG/PMTiles), then produce STAC/DCAT/PROV artifacts, sign them, and open a promotion PR with a rollback tag.

Step	Output	Gate(s)
1) Snapshot raw sources	RAW + MANIFEST + checksums	Checksum verification; retention policy
2) Geometry fix + CRS	Fixed vector layer	CRS/axis sanity; geometry validity
3) Key normalization	Stable, de-duplicated keys	Uniqueness; referential integrity
4) Attribute QC	QC report	Policy bounds; null/domain constraints
5) Delivery formats	GeoParquet, COG, PMTiles	Format validators (gdalinfo, parquet schema, sample tile)
6) Catalogs + lineage	STAC/DCAT/PROV	Schema validation; required fields; signatures
7) Promotion proposal	Signed PR + rollback tag	PR body checklist complete

Promotion-ready output layout

data/soils/raw/<provider>/<yyyymmdd>/
data/soils/processed/<region>/<job_id>/ (GeoParquet, COG, PMTiles)
data/stac/soils/<region>/<job_id>/
data/catalog/dcat/soils/<region>/<job_id>/
data/prov/runs/<job_id>/

9. Domain Lane Example: Work Zones (WZDx) as a Governed Feed

The document highlights Work Zone Data Exchange (WZDx) feeds as a strong candidate for watcher-based ingestion: a standardized GeoJSON FeatureCollection, schema-first validation, and natural mapping into STAC Items (per RoadEvent) linked to a DCAT Dataset (per feed).

Concept	Practical integration
Conditional fetch	Use ETag / If-None-Match when supported to avoid unnecessary downloads.
STAC mapping	Each RoadEvent Feature becomes a candidate STAC Item with geometry/time and prov:wasDerivedFrom the raw snapshot.
Policy + schema validation	Treat upstream JSON Schemas as inputs to Conftest/OPA gates before any promotion.

10. Thin-slice Implementation Plan (Buildable Increments)

To integrate quickly without over-committing, adopt these increments. Each is reviewable and merge-safe.

Increment	What you add	Definition of Done
Slice A: Repo scaffolding + CI skeleton	- Add reusable CI workflow + Conftest PR gate. - Add kill-switch secret/flag and guardrails step. - Create manifests/ci/ci_manifest.yml as the single source of truth for lanes.	CI runs and blocks merges on failing policy; no domain logic yet.
Slice B: Deterministic detection + lane runner	- Add tools/detection/detect_changes.py and tools/detection/planner.py. - Add tools/validation/common/idempotency.py and lineage.py. - Implement one minimal lane validator (for example lane_stac_catalog.py).	A push or scheduled run deterministically selects lanes and emits OpenLineage events.
Slice C: WPE offline end-to-end demo	- Add the single-file watcher->planner->PR demo with local gates (stac-lint, prov-schema-check, story-node-render). - Ensure it is fully offline and deterministic via a fixed seed.	Engineers can prove the full loop locally in minutes and validate artifacts.

Slice D: First domain lane (Soils or WZDx)	- Implement a watcher (ETag or upstream snapshot) and a lane executor. - Emit STAC/DCAT/PROV and sign bundles. - Open a promotion PR with attached evidence and rollback tag.	A real dataset can be promoted via governed PRs with provenance and policy enforcement.

11. Risks, Tradeoffs, and Governance Review

Known integration risks

- Tool and spec drift: pin versions for Conftest/Cosign and schemas; treat updates as governed changes.
- Nondeterministic dependencies: network calls, unordered JSON, and unpinned containers break reproducibility.
- Sensitive data: watchers must classify/redact before any PR; policies should be deny-by-default for restricted layers.

Governance review checklist (minimum):

- [ ] Policies are merge-blocking and have unit tests.
- [ ] Required catalog/provenance fields validate (STAC/DCAT/PROV).
- [ ] Idempotency keys exist for runs, lanes, and PR creation.
- [ ] Kill-switch is documented and tested (neutral exit code path).
- [ ] Evidence artifacts (gate reports, provenance, signatures) are attached and retained per policy.

References

Latest Ideas-2. Latest Ideas-2.pdf (internal working document; 183 pages). Key sections referenced: agent architecture (pp. 1-6), deterministic IDs and idempotency (pp. 9-10), detection/planner/lanes (pp. 11-19), reusable CI workflows and policy gates (pp. 77-83), WZDx notes and watcher-signals extension (pp. 43-46), soils pipeline runbook (pp. 96-99), offline WPE demo (pp. 99-104), change-trigger QC promotion example (pp. 109-113).

This PDF is an integration synthesis. It intentionally avoids inventing new system facts not present in the source document; when you adopt it, pin exact versions and paths inside your repository as governed decisions.